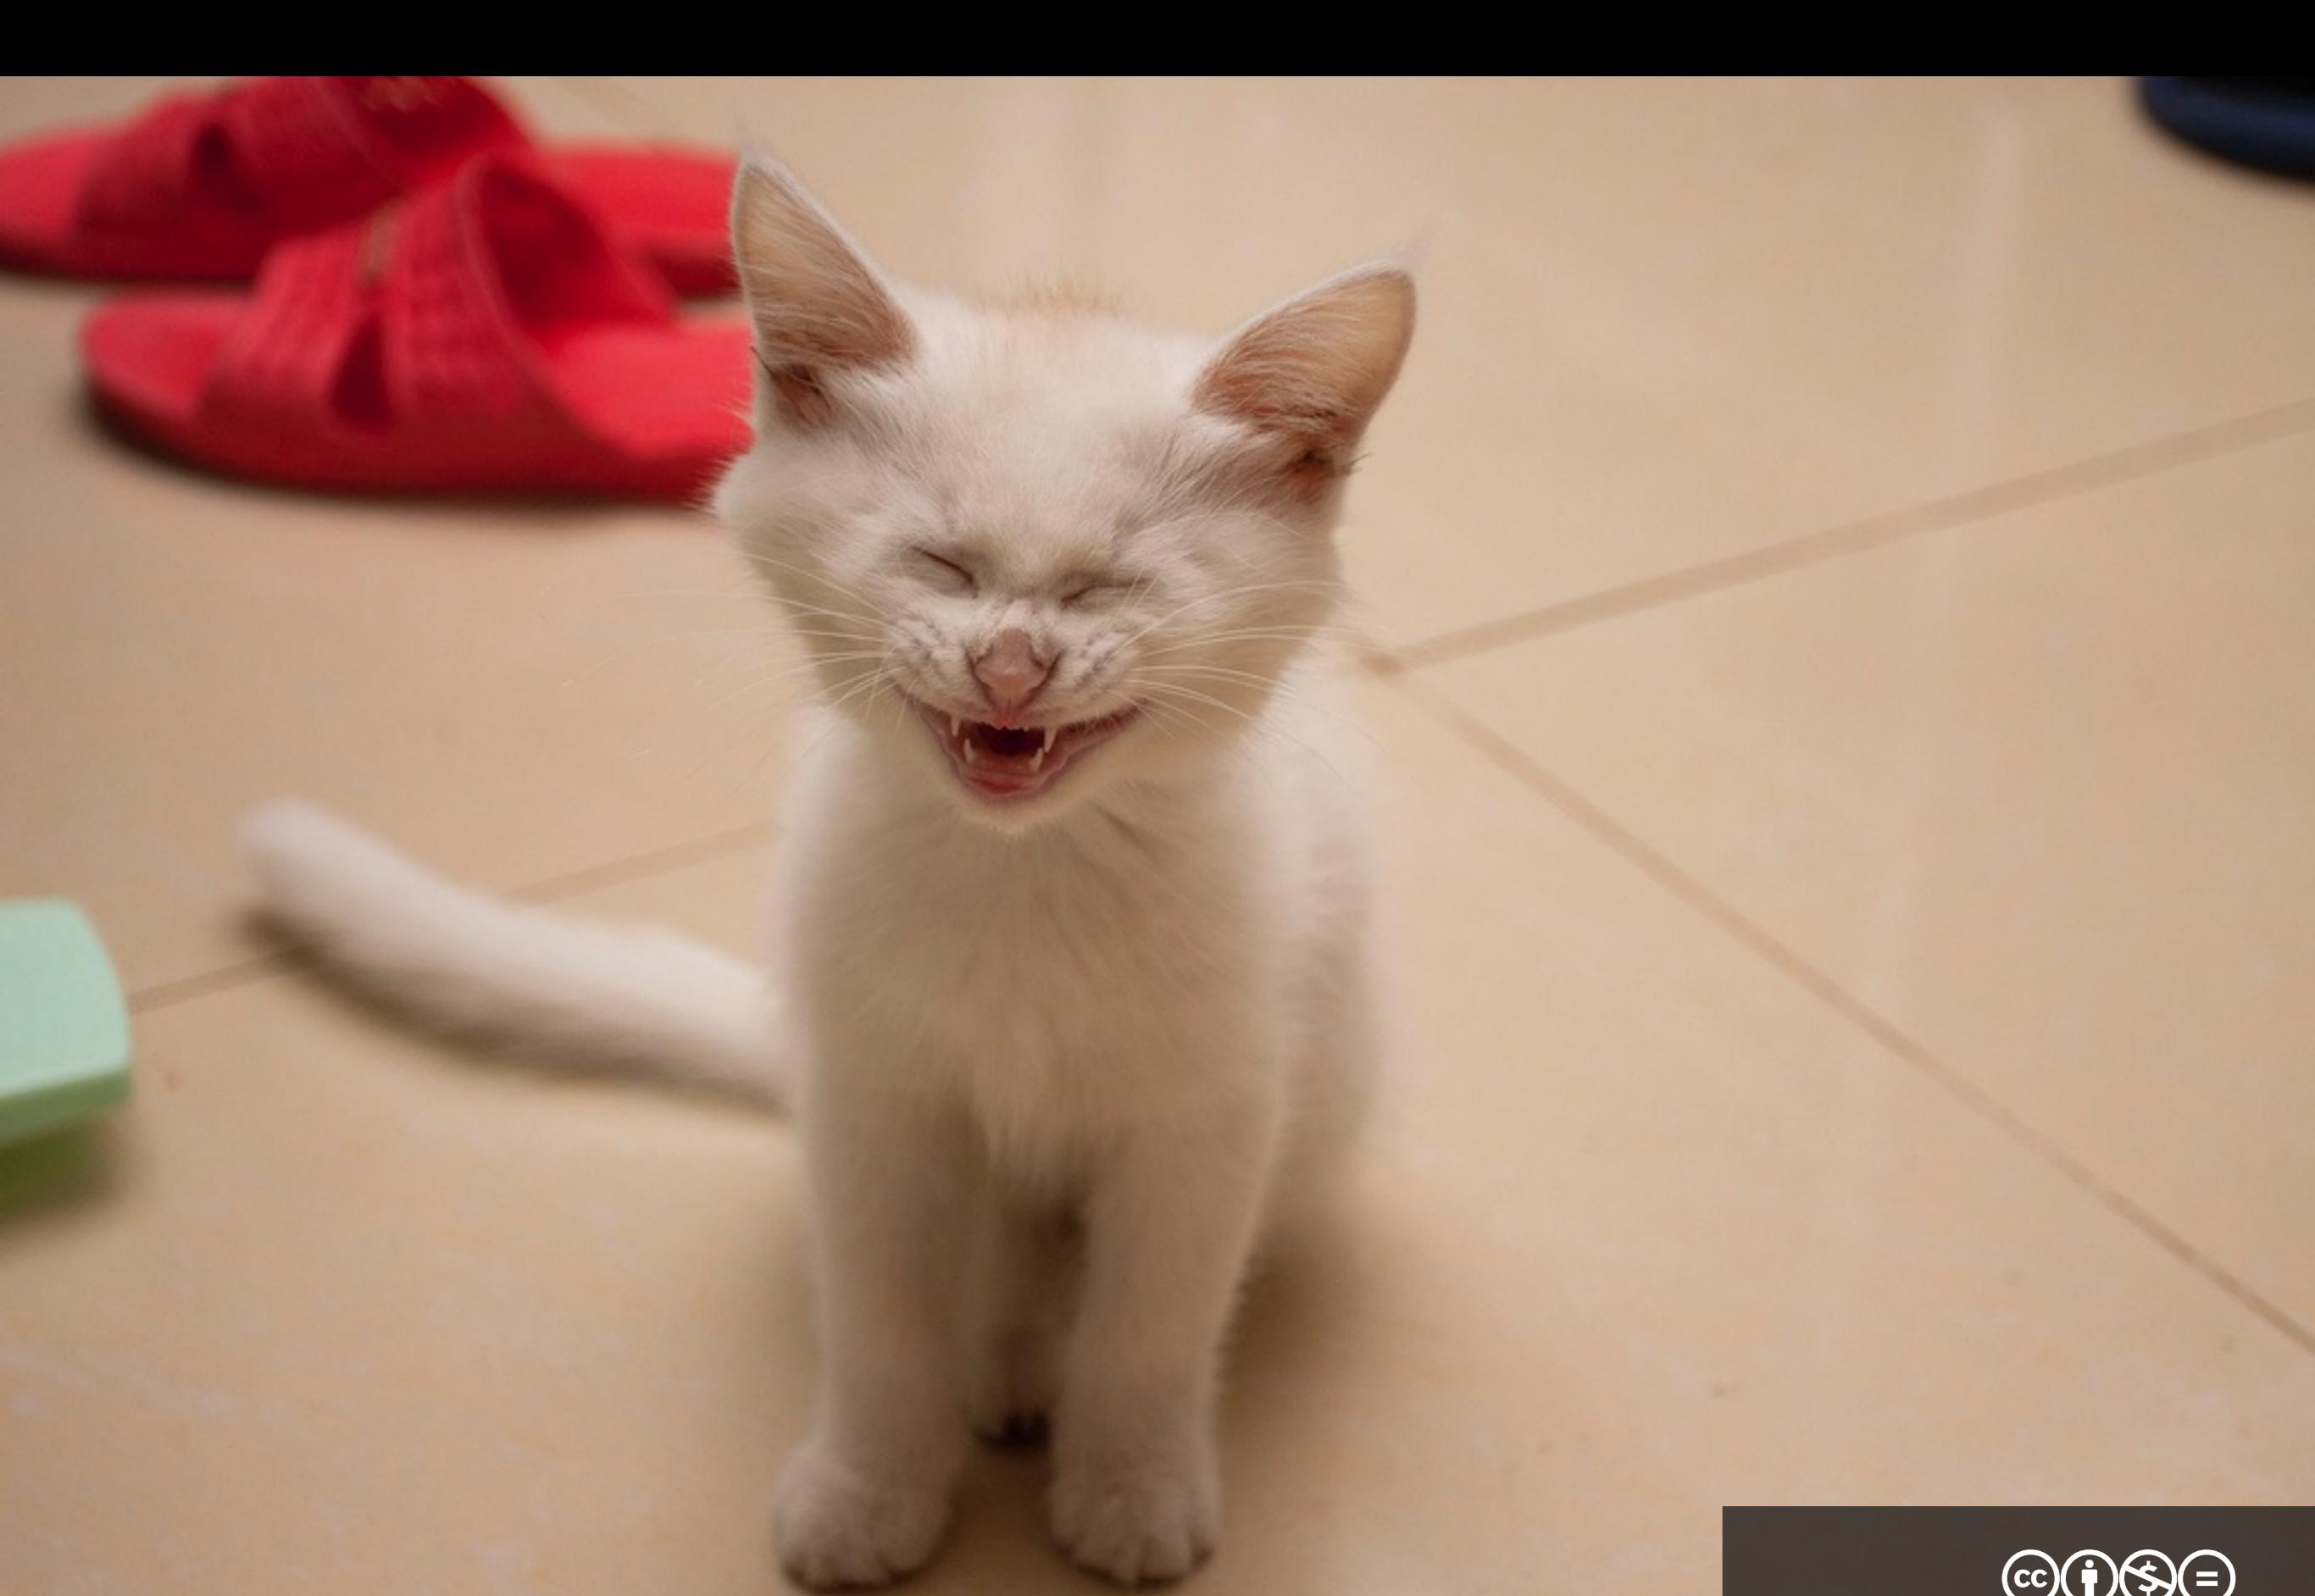


# AN ASTRONOMER'S INTRODUCTION TO GAUSSIAN PROCESSES

Dan Foreman-Mackey

CCPP@NYU // [github.com/dfm](https://github.com/dfm) // [@exoplaneteer](https://twitter.com/exoplaneteer) // [dfm.io](https://dfm.io)



Flickr user **lizphung**

[github.com/dfm/gp](https://github.com/dfm/gp)

---

[gaussianprocess.org/gpml](http://gaussianprocess.org/gpml)

---

Rasmussen & Williams

I write **code** for good & **astrophysics**.

---

I (probably) do **data science**.

---

not data  
science.

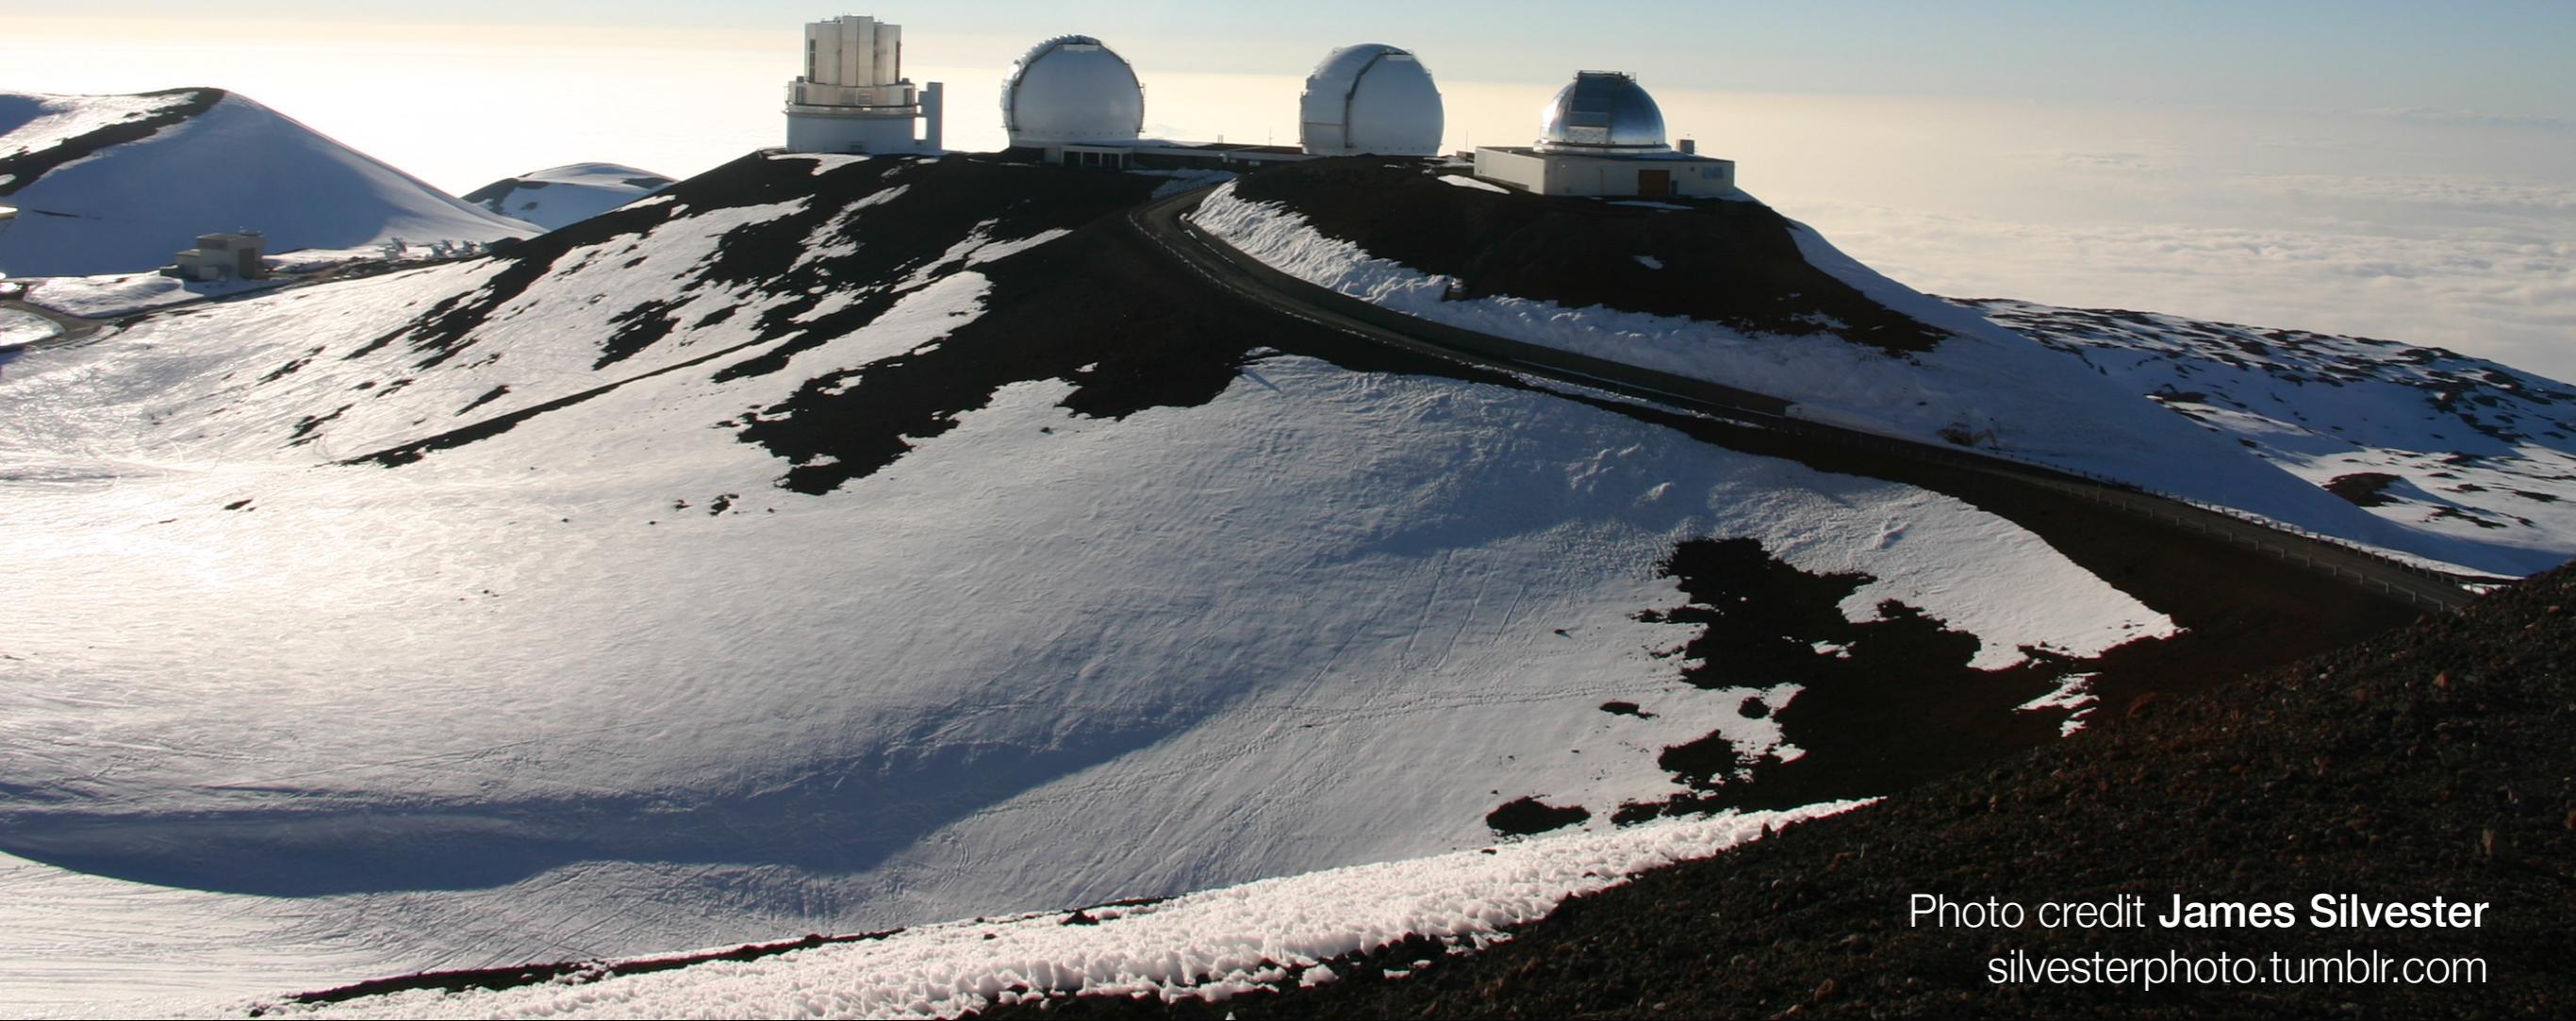


Photo credit **James Silvester**  
[silvesterphoto.tumblr.com](http://silvesterphoto.tumblr.com)

CATERING PREP

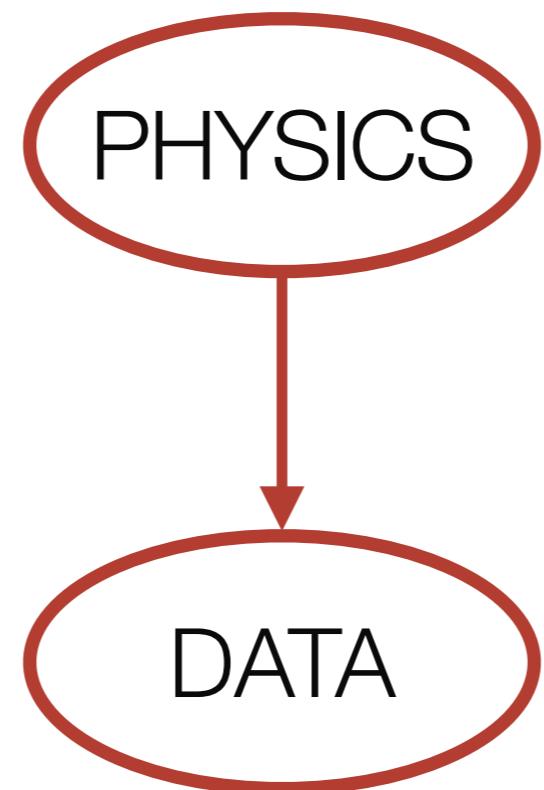
# data science.



Flickr user Marcin Wichary

# Data Science

---



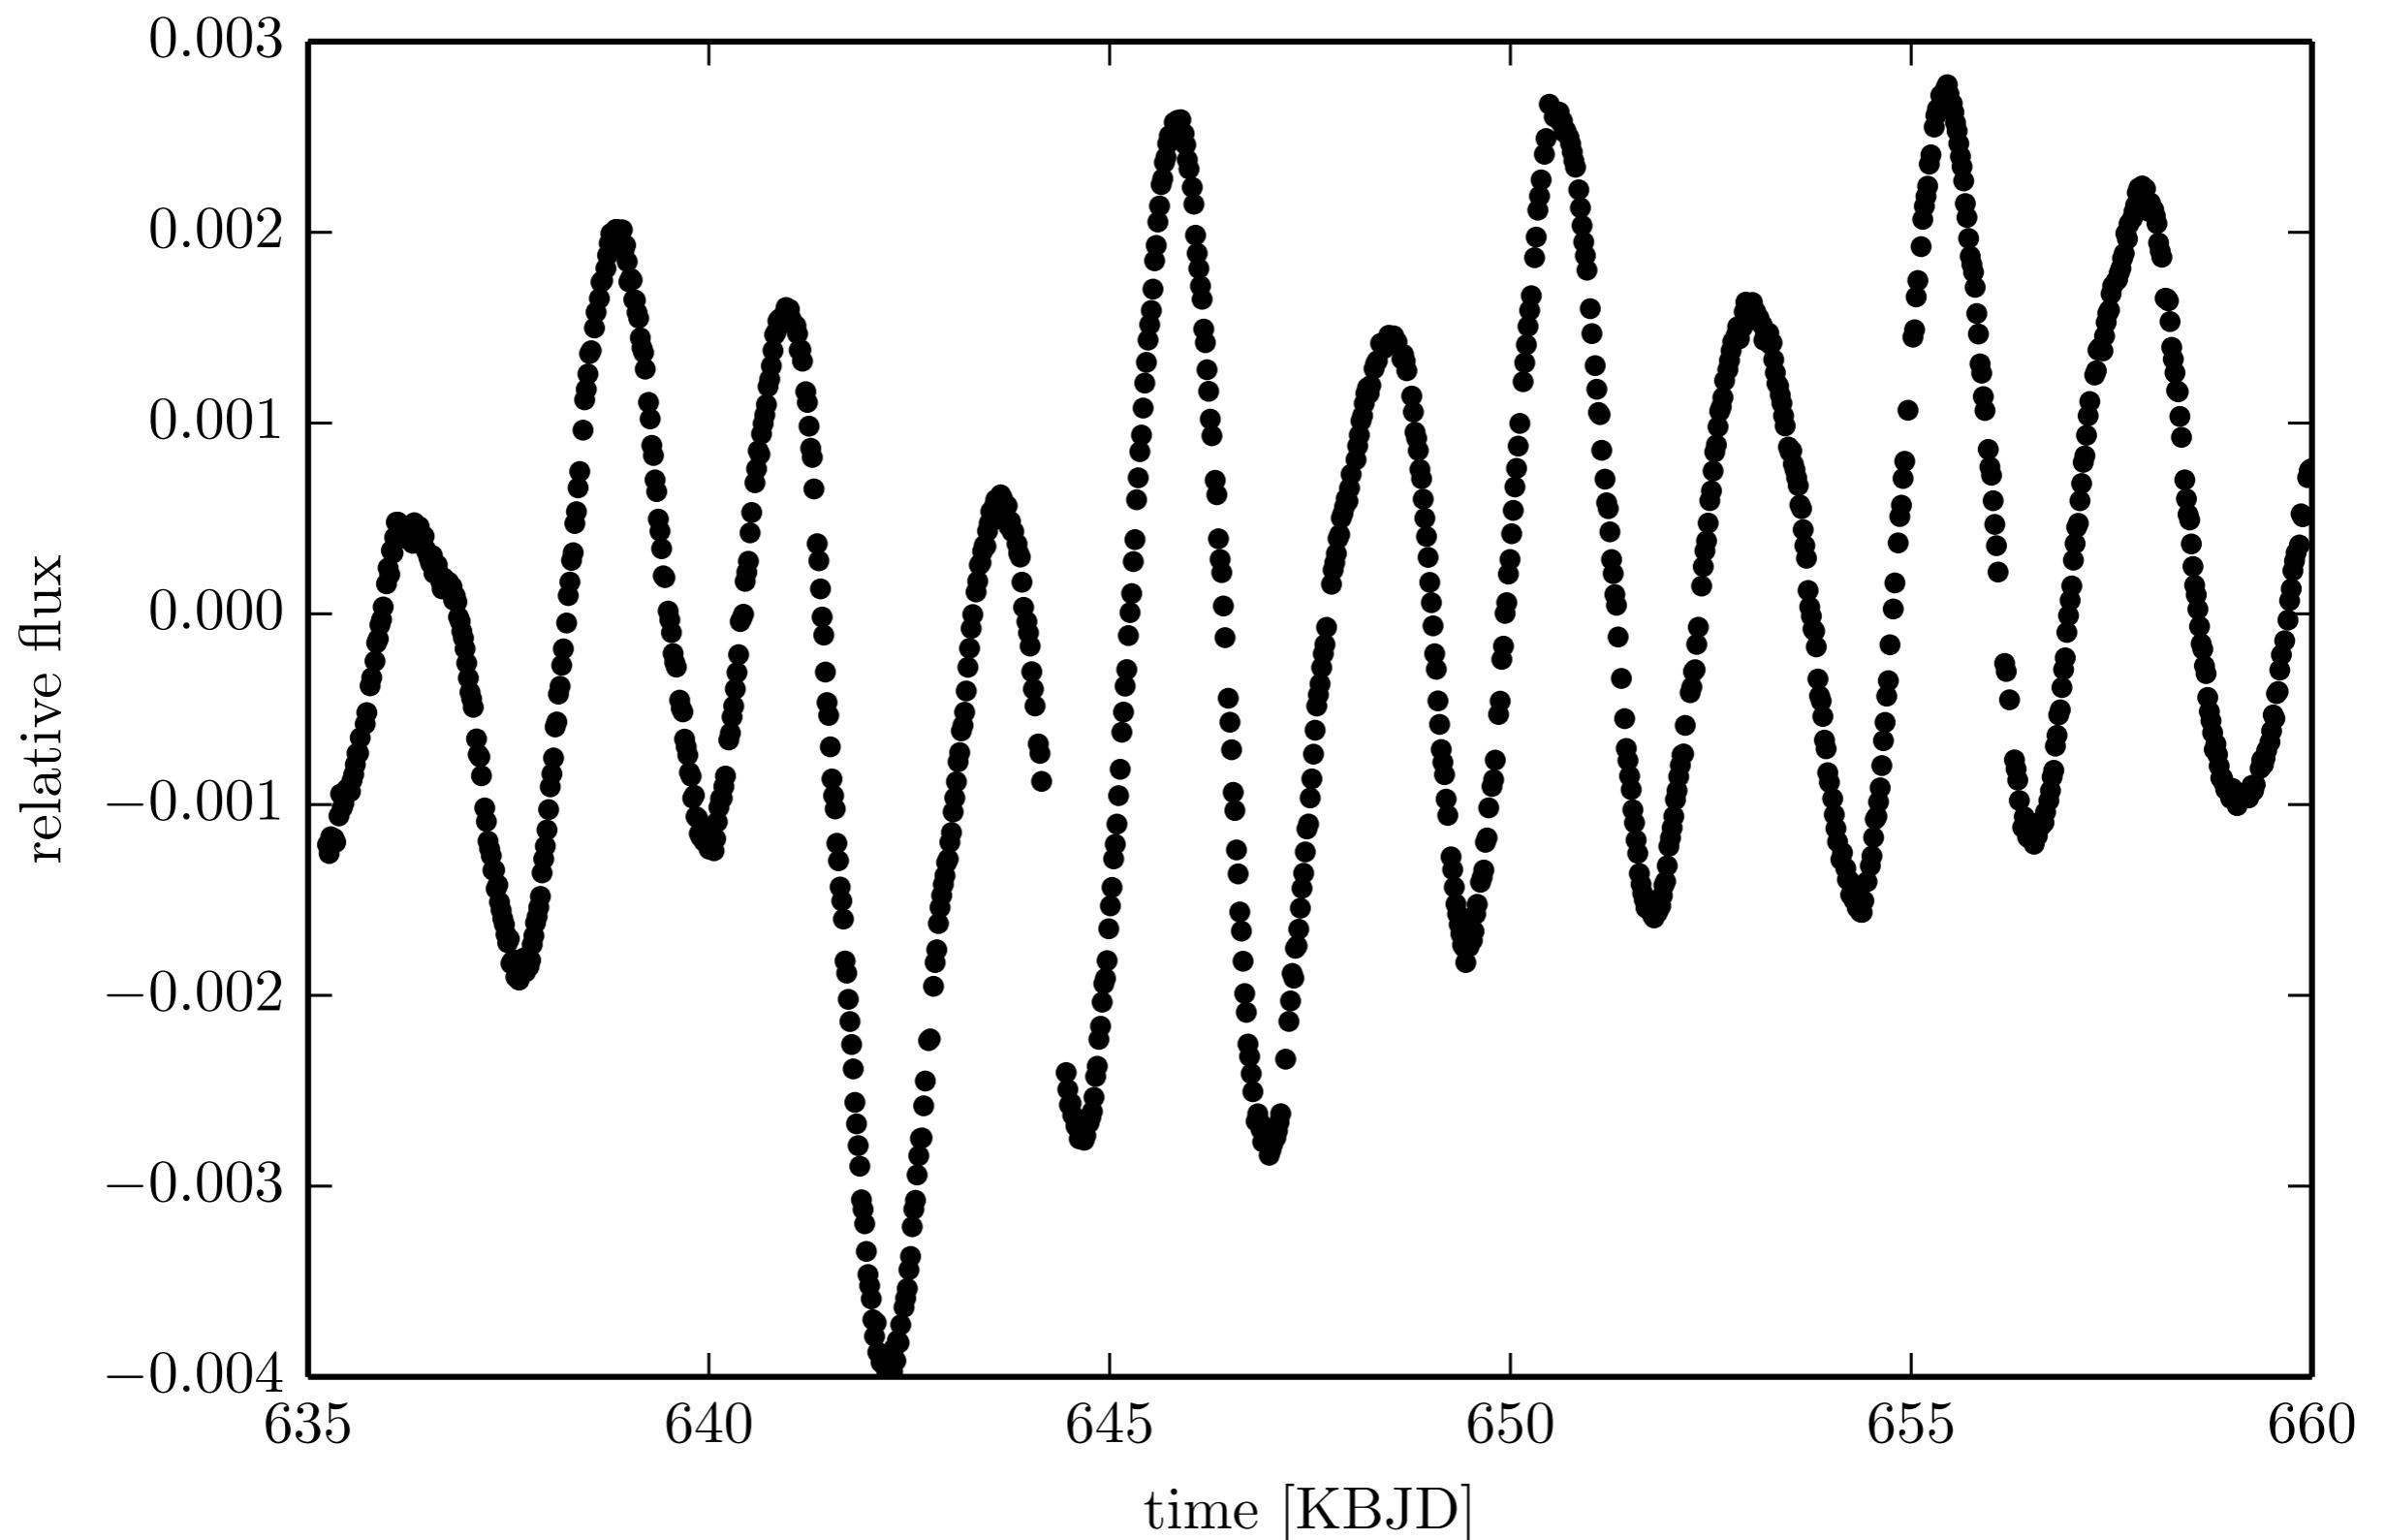
$$p(\text{data} \mid \text{physics})$$

I work with **Kepler** data.

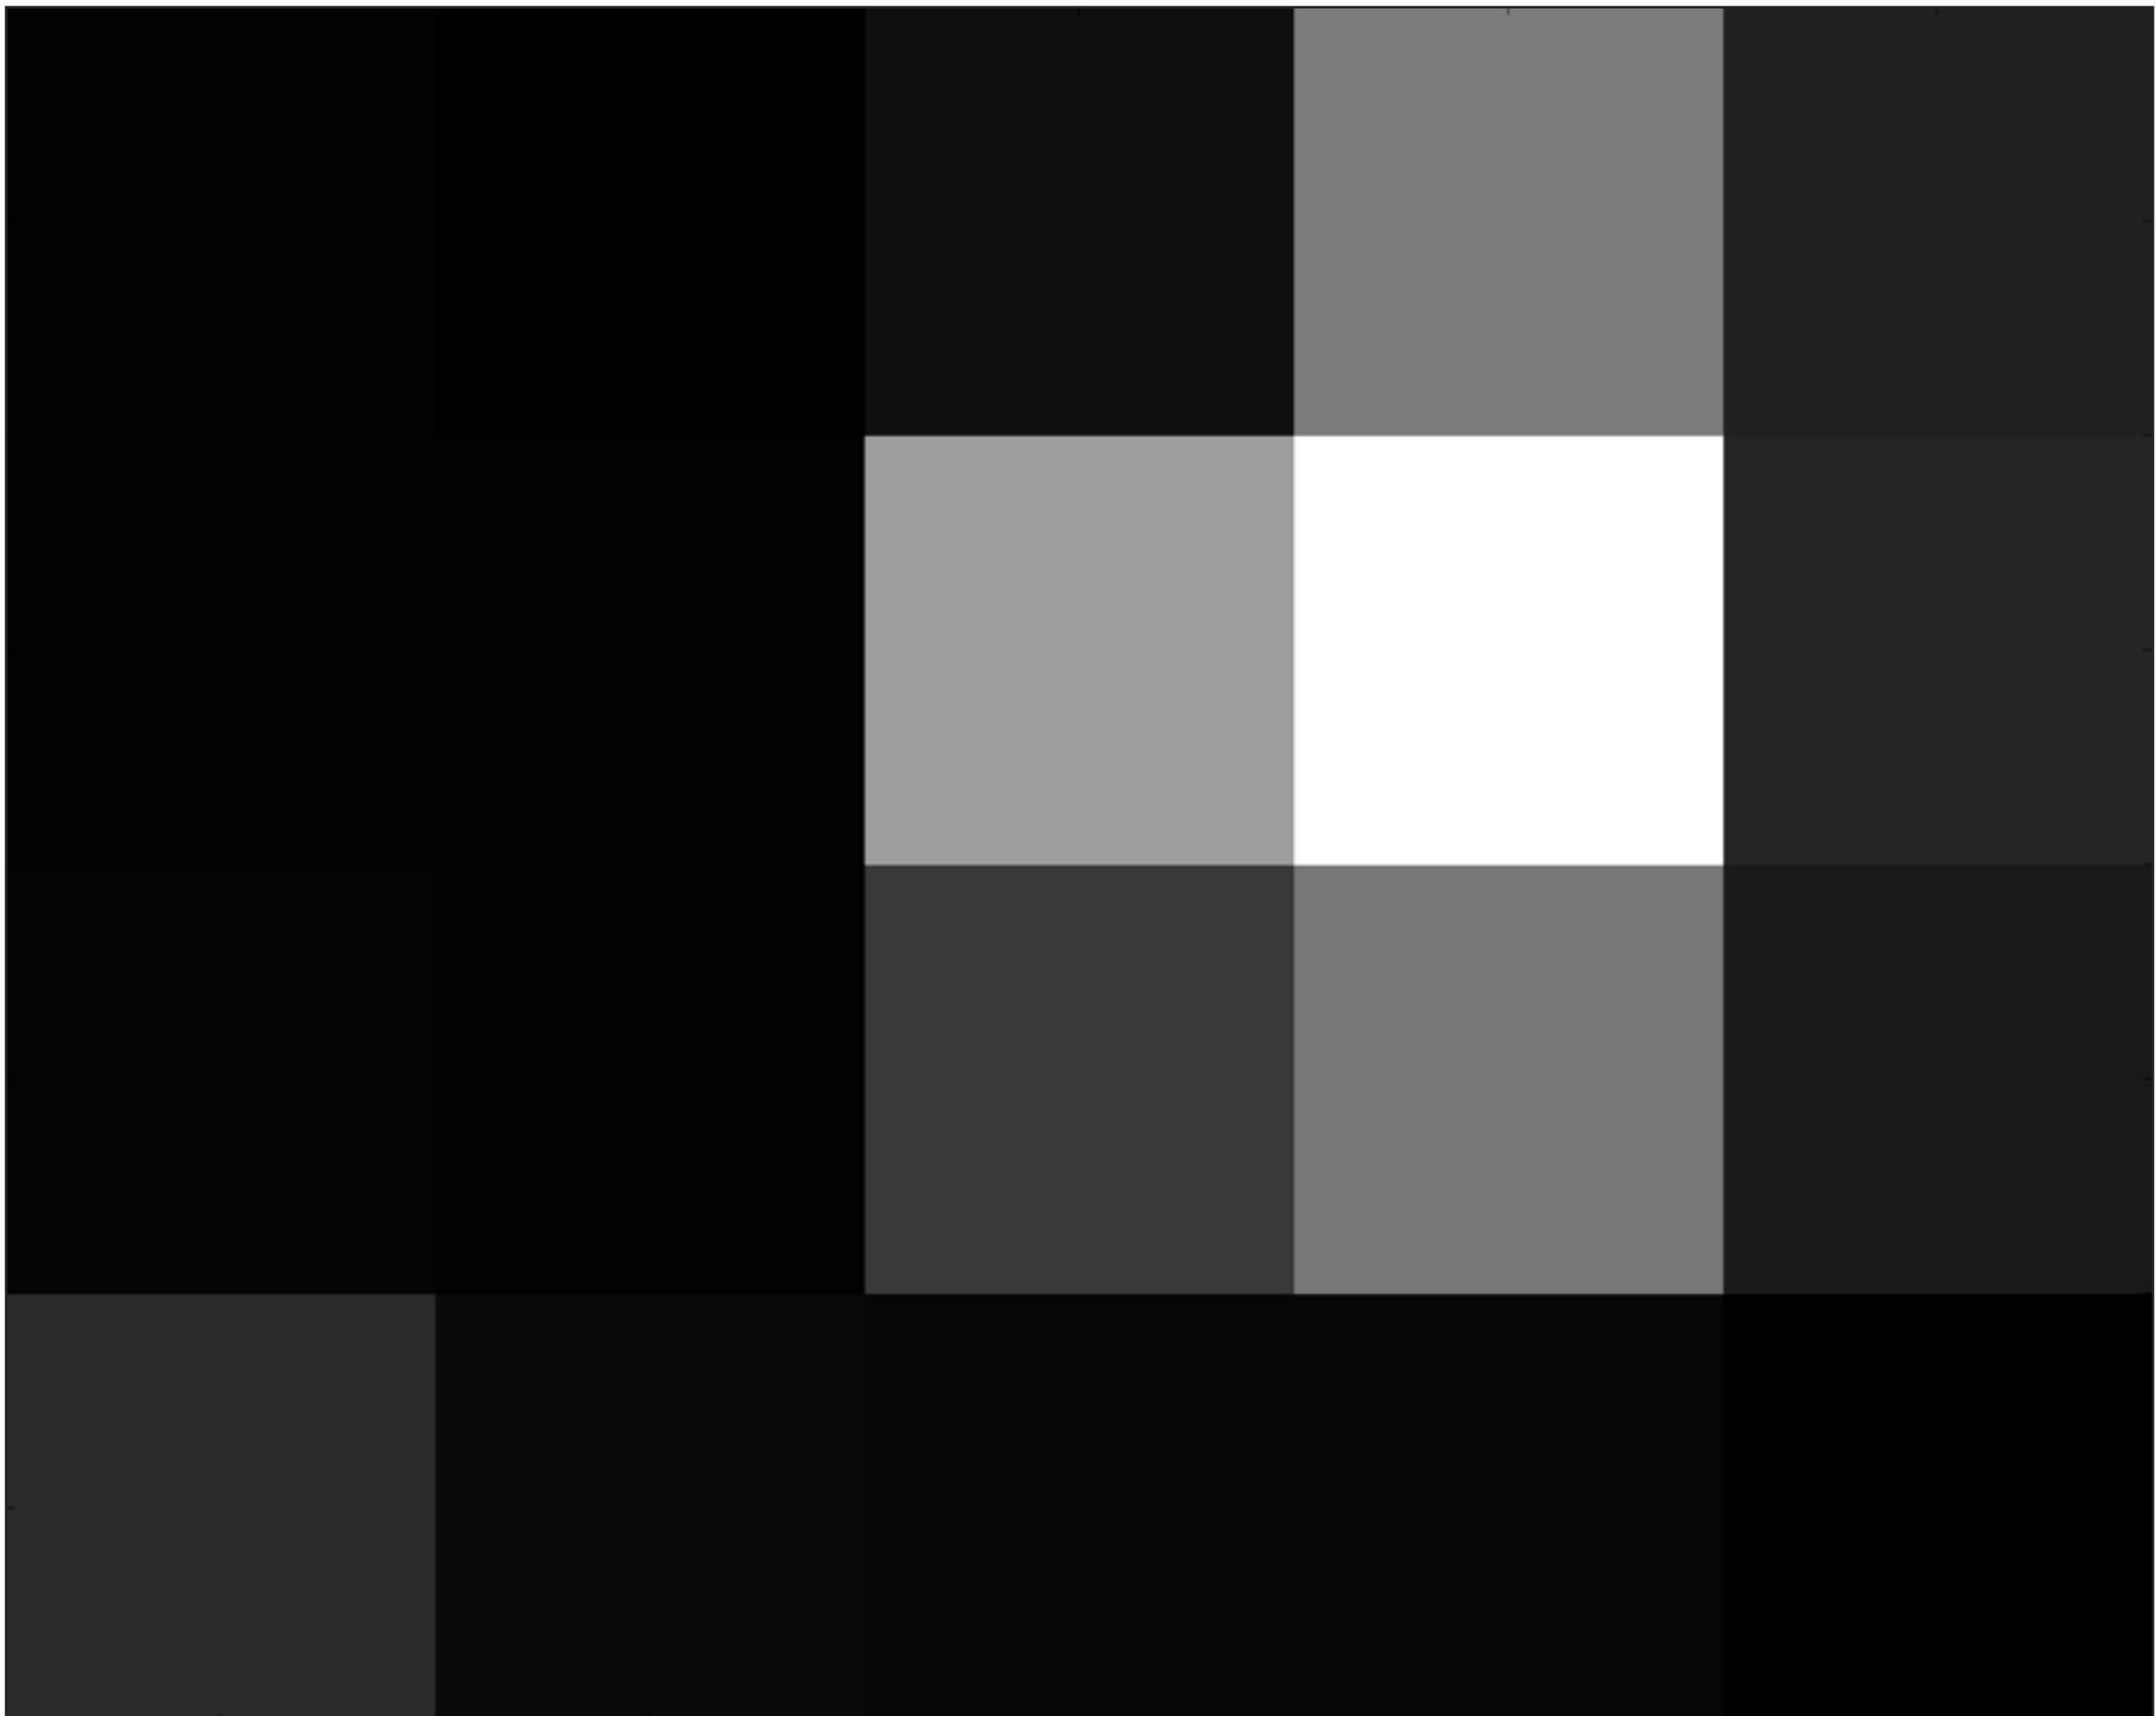
---

I get really passionate about

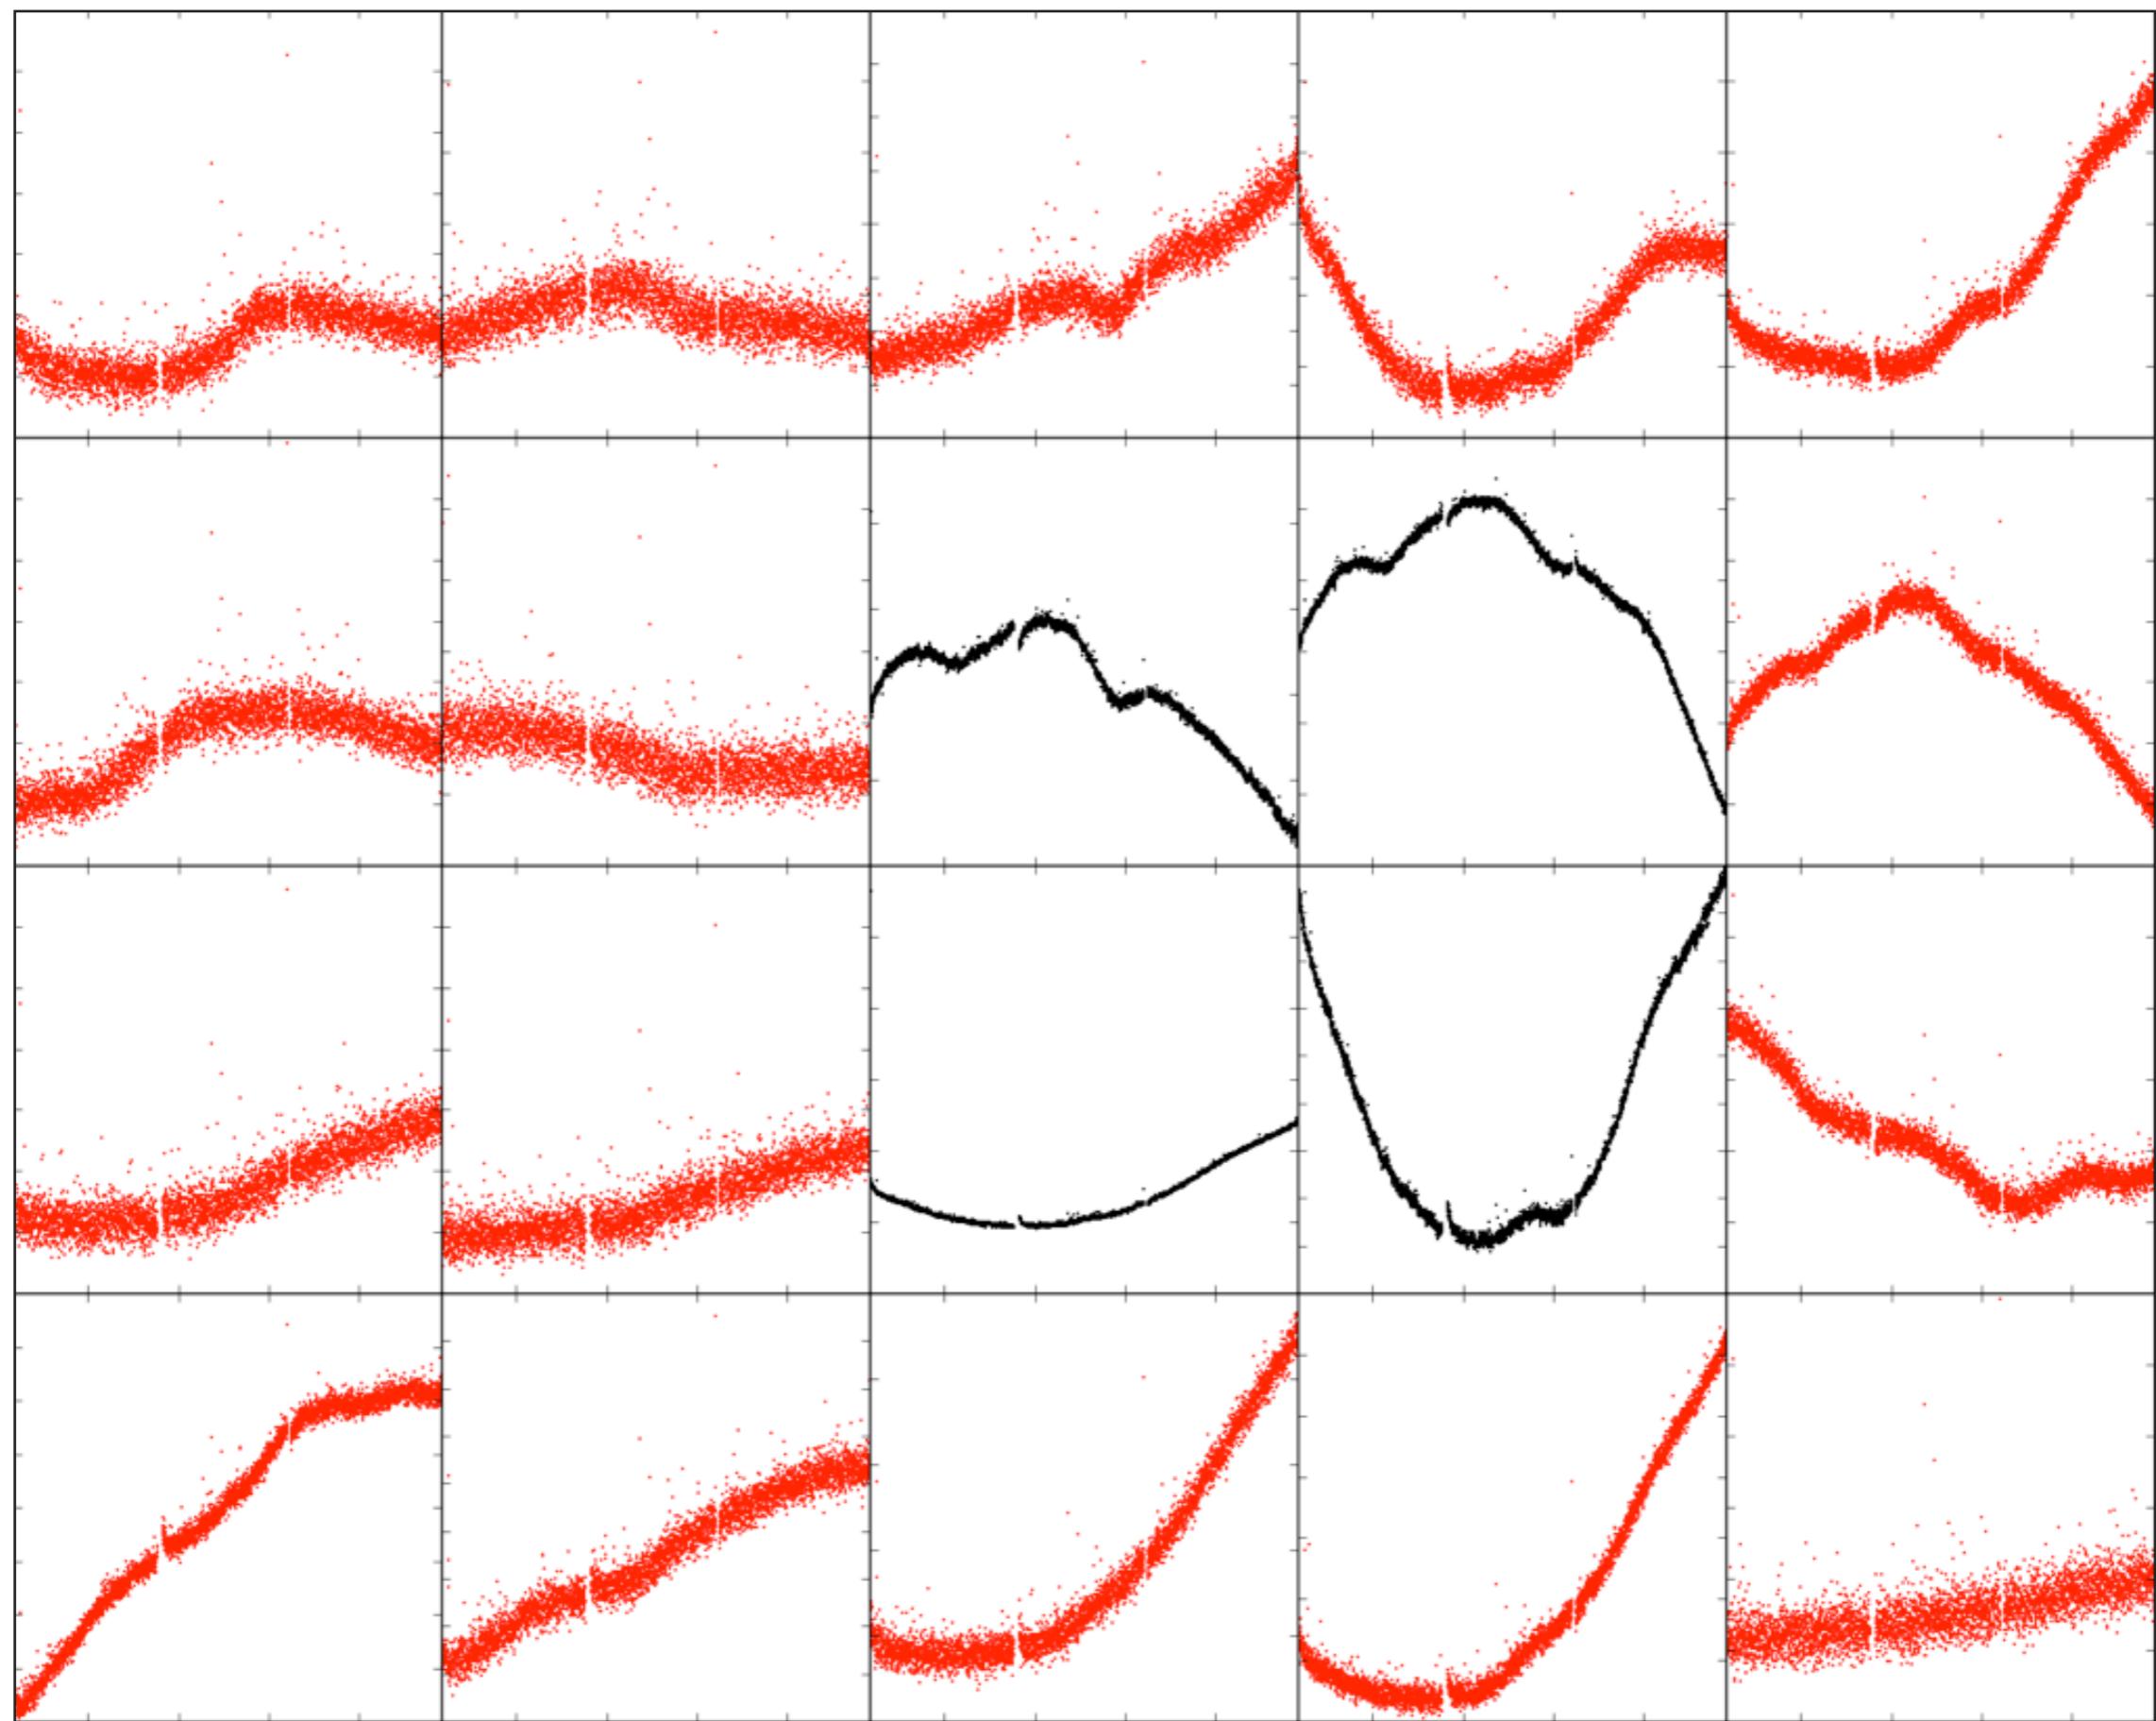
**NOISE**



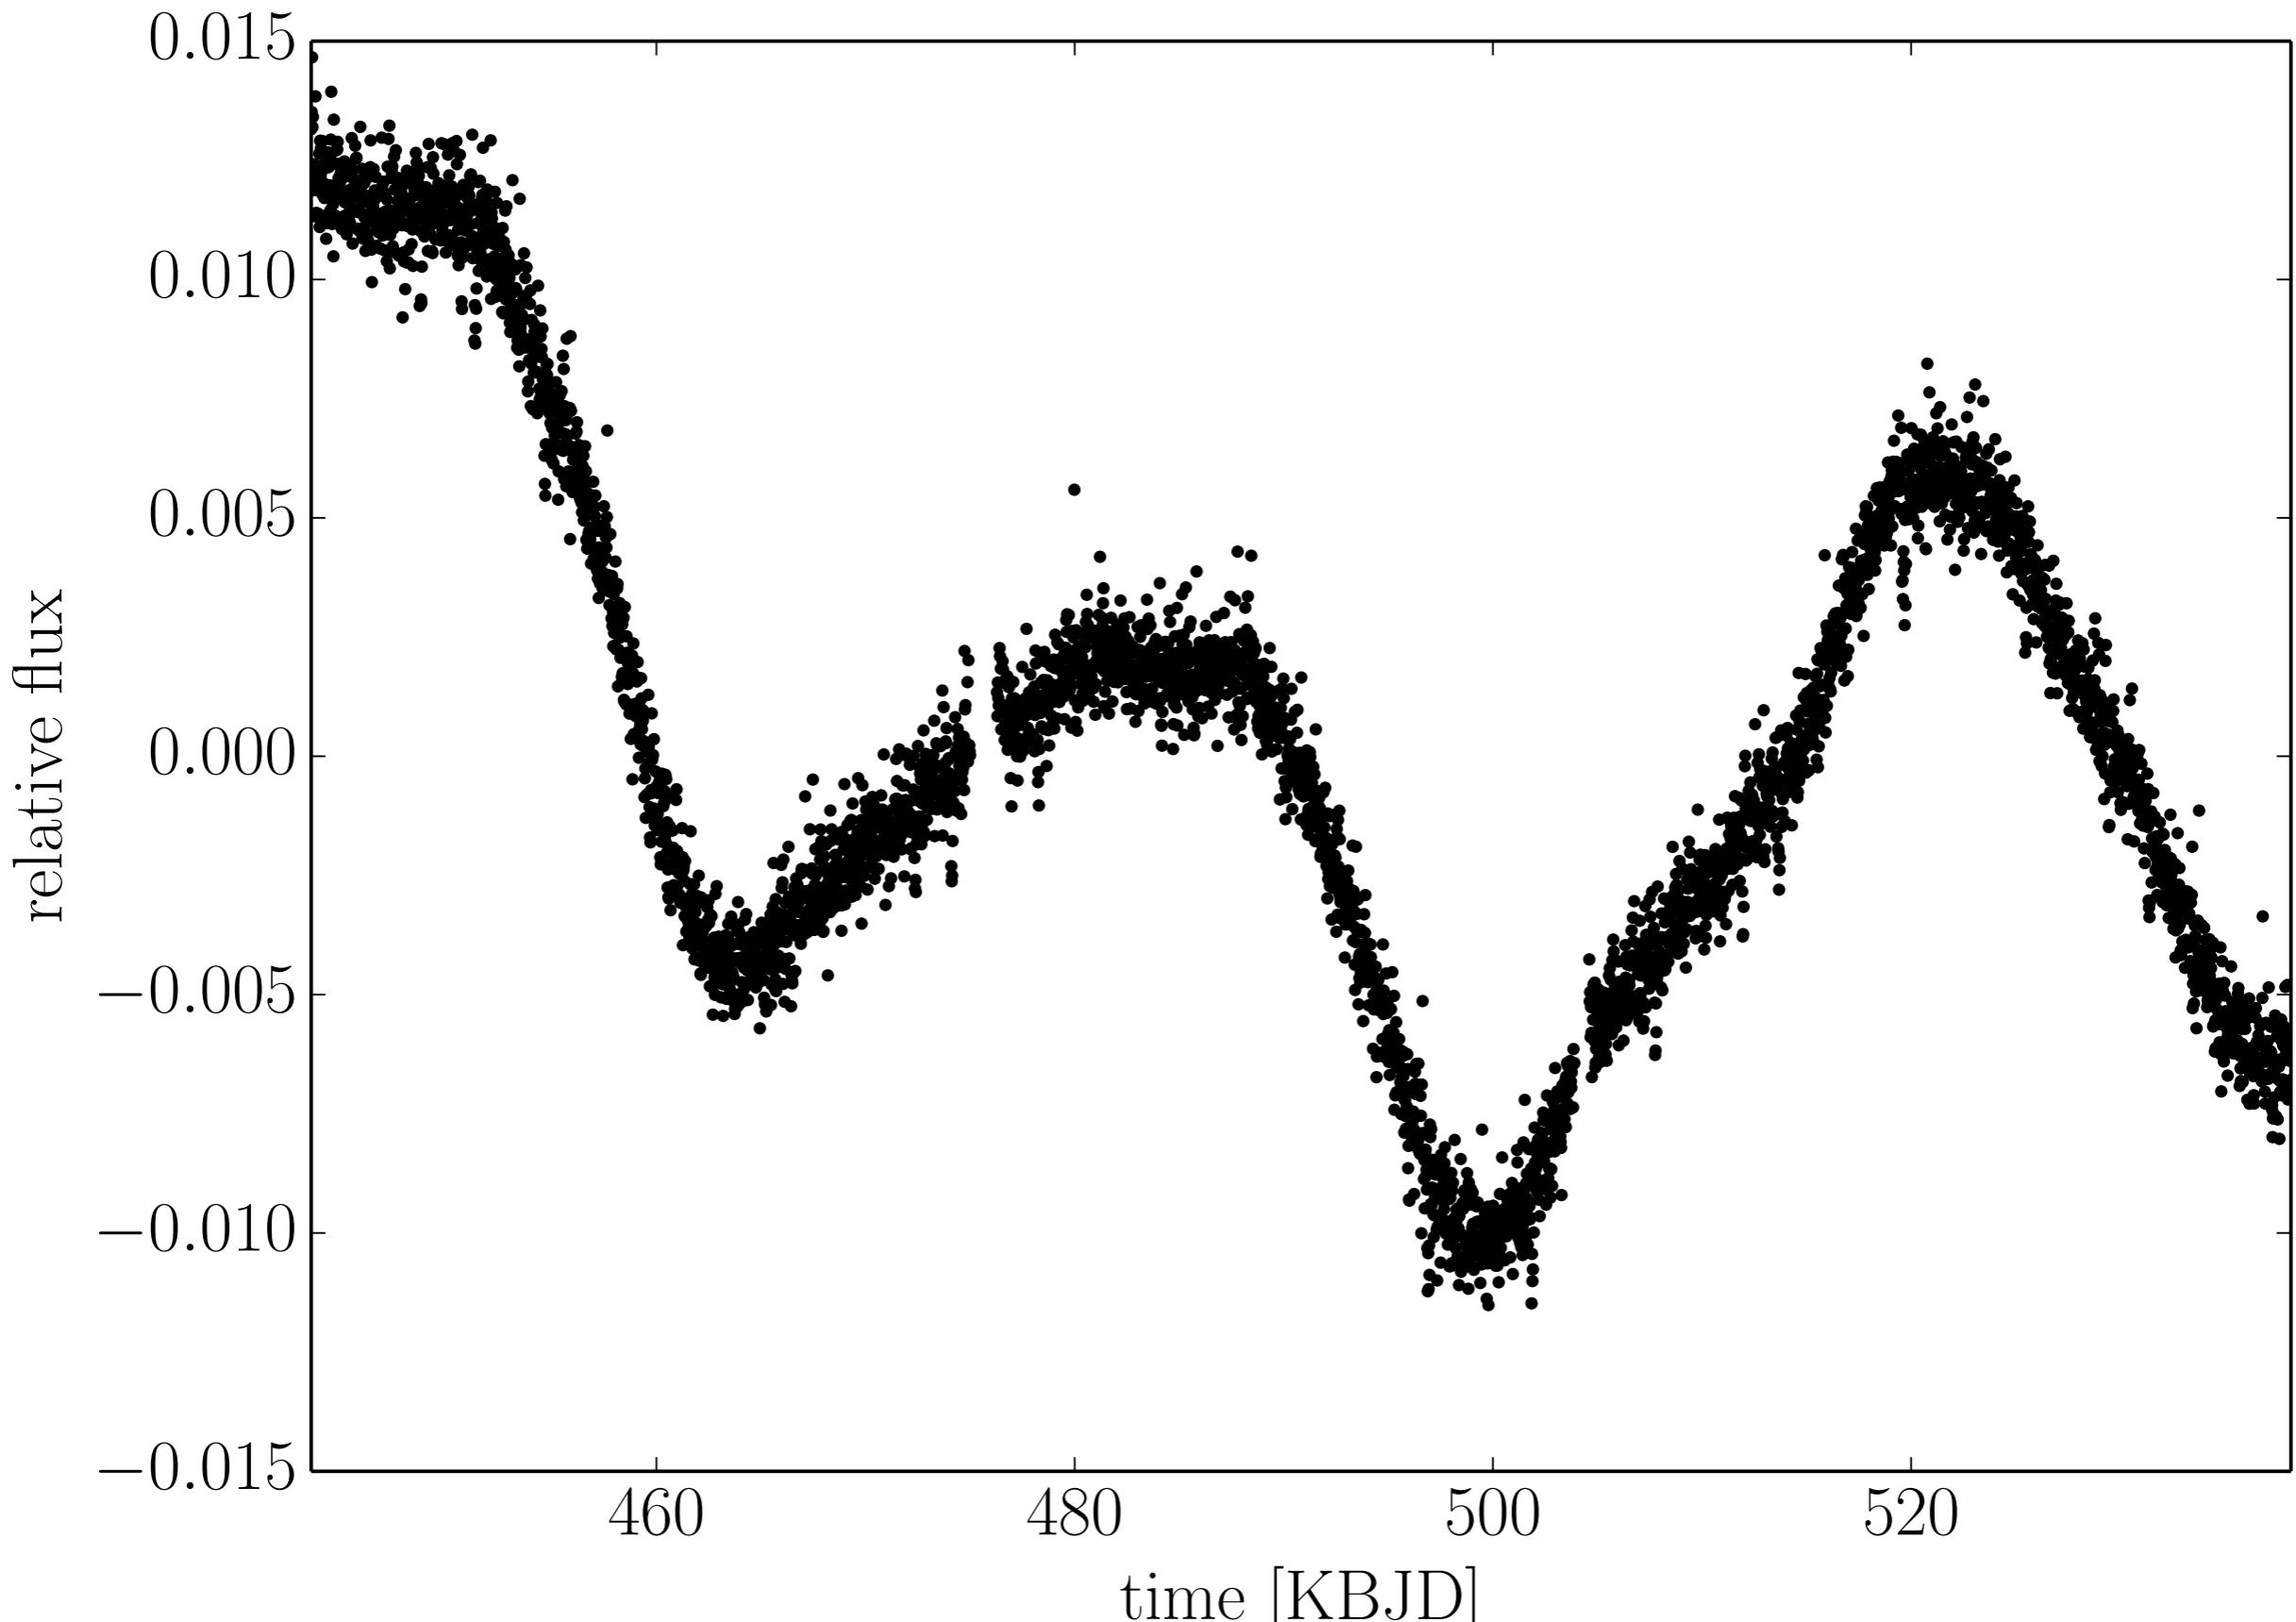
KIC 3223000



Kepler 32



Kepler 32



Kepler 32

# HACKS ALL HACKS

and I'm *righteous!*

Why not model all the things?

---

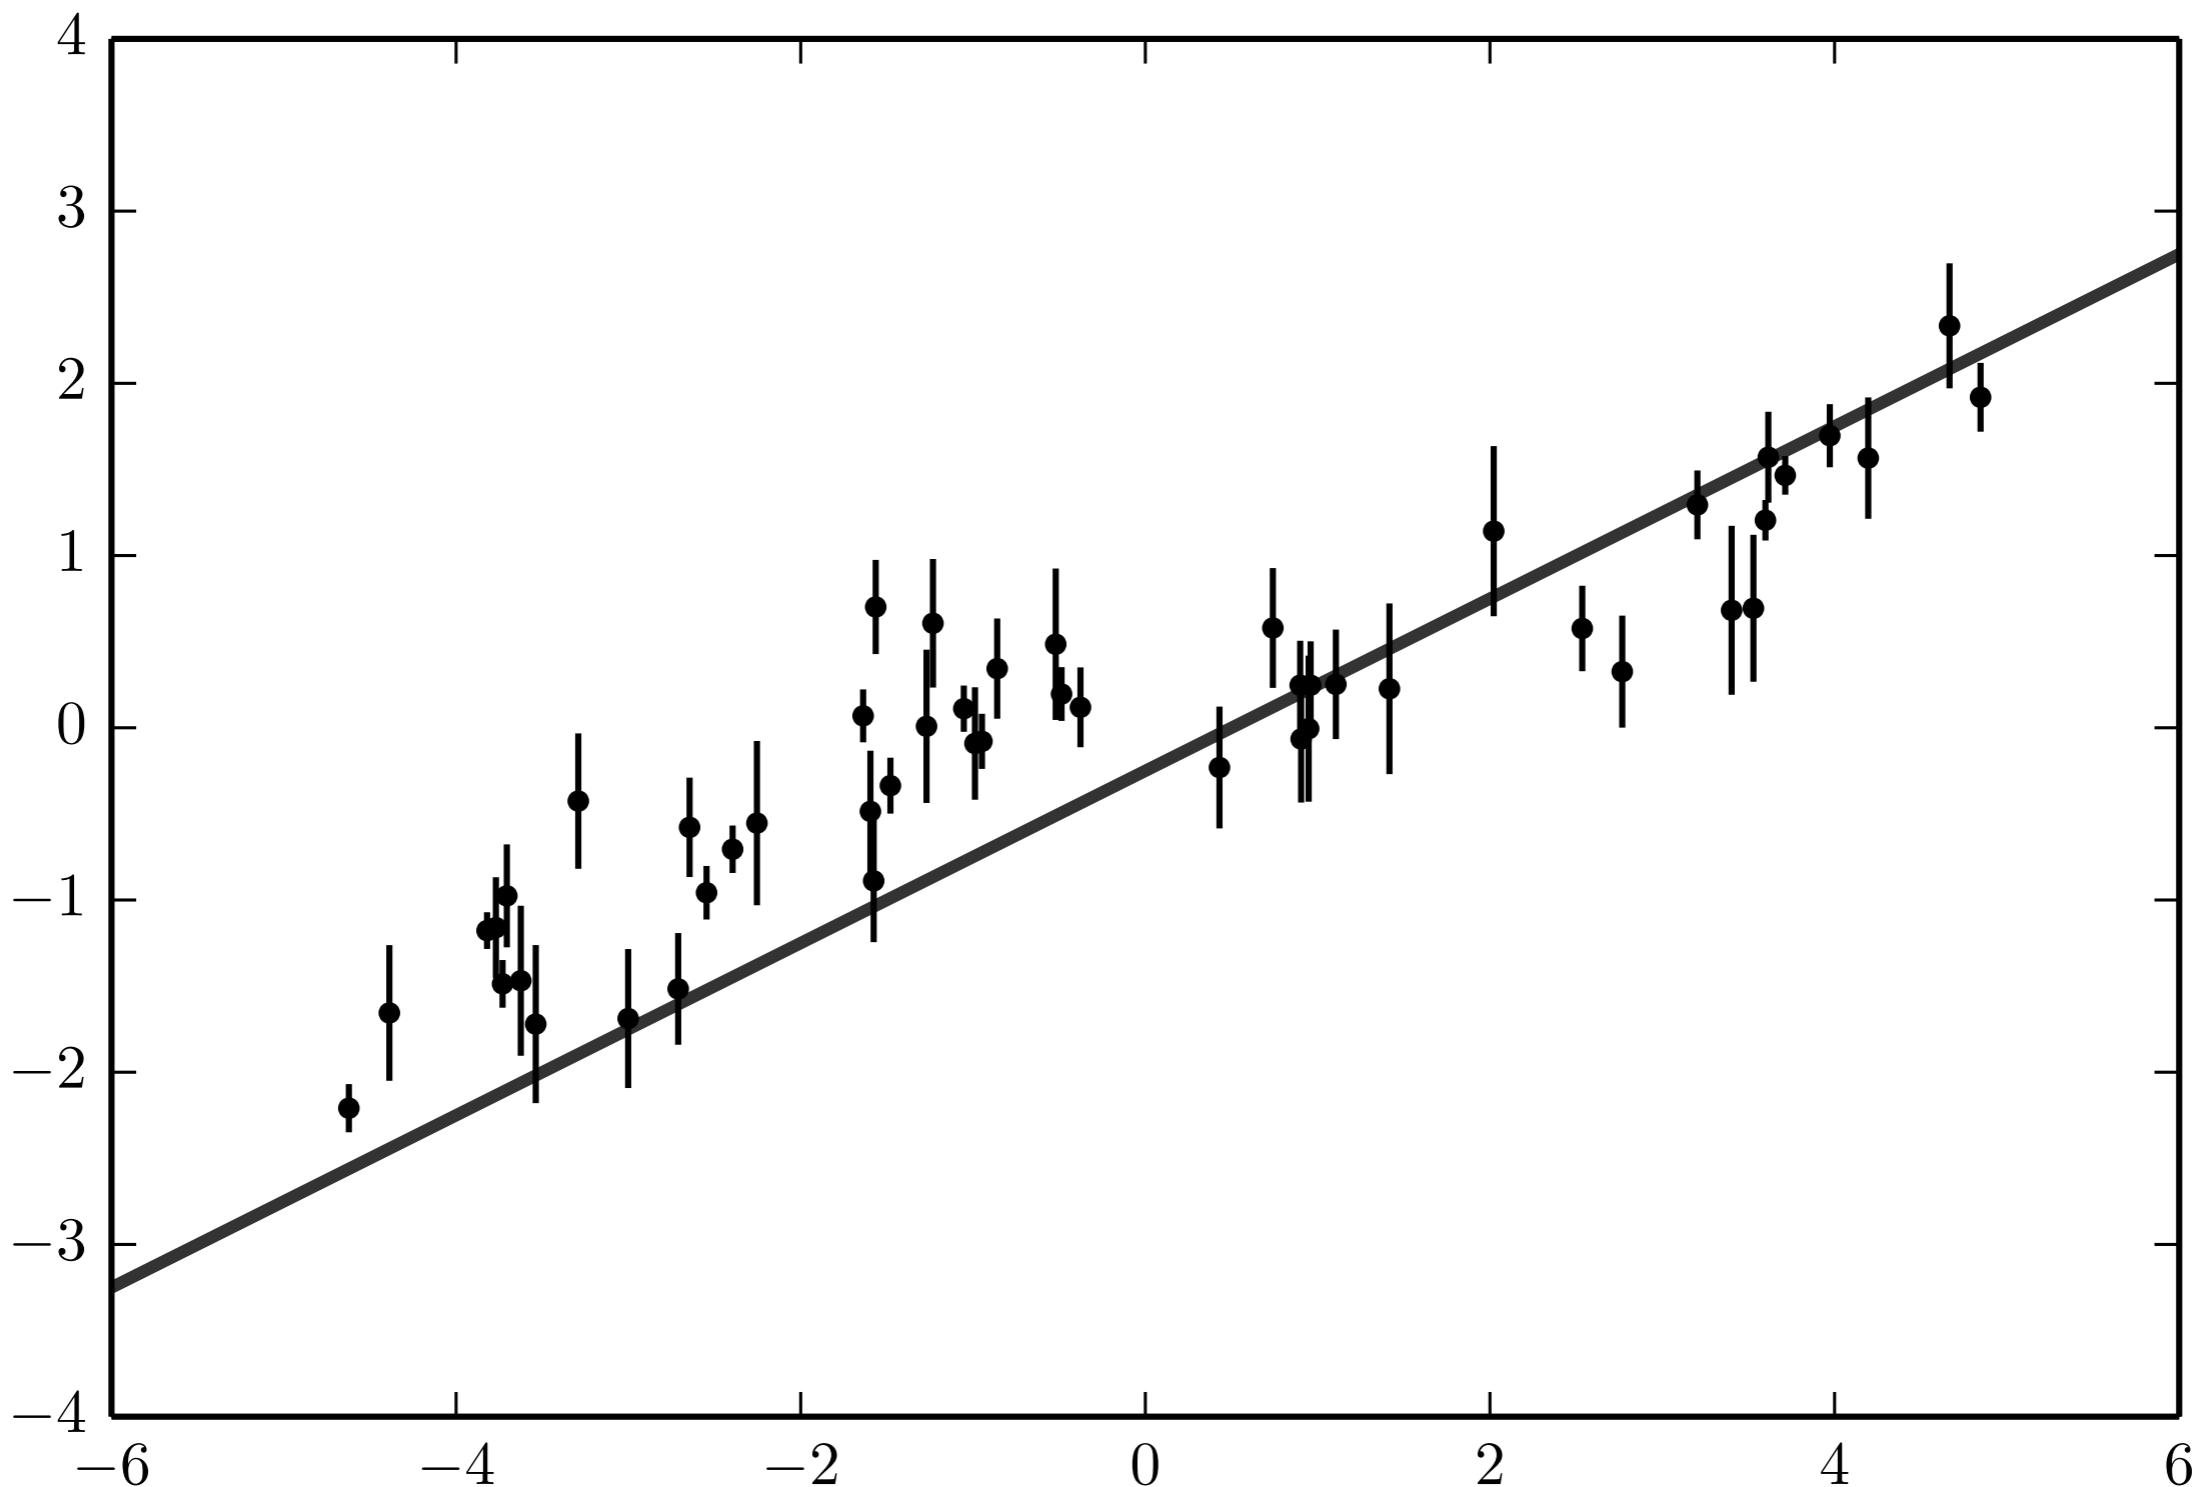
with, for example, a Gaussian Process

1

The power of **correlated noise.**

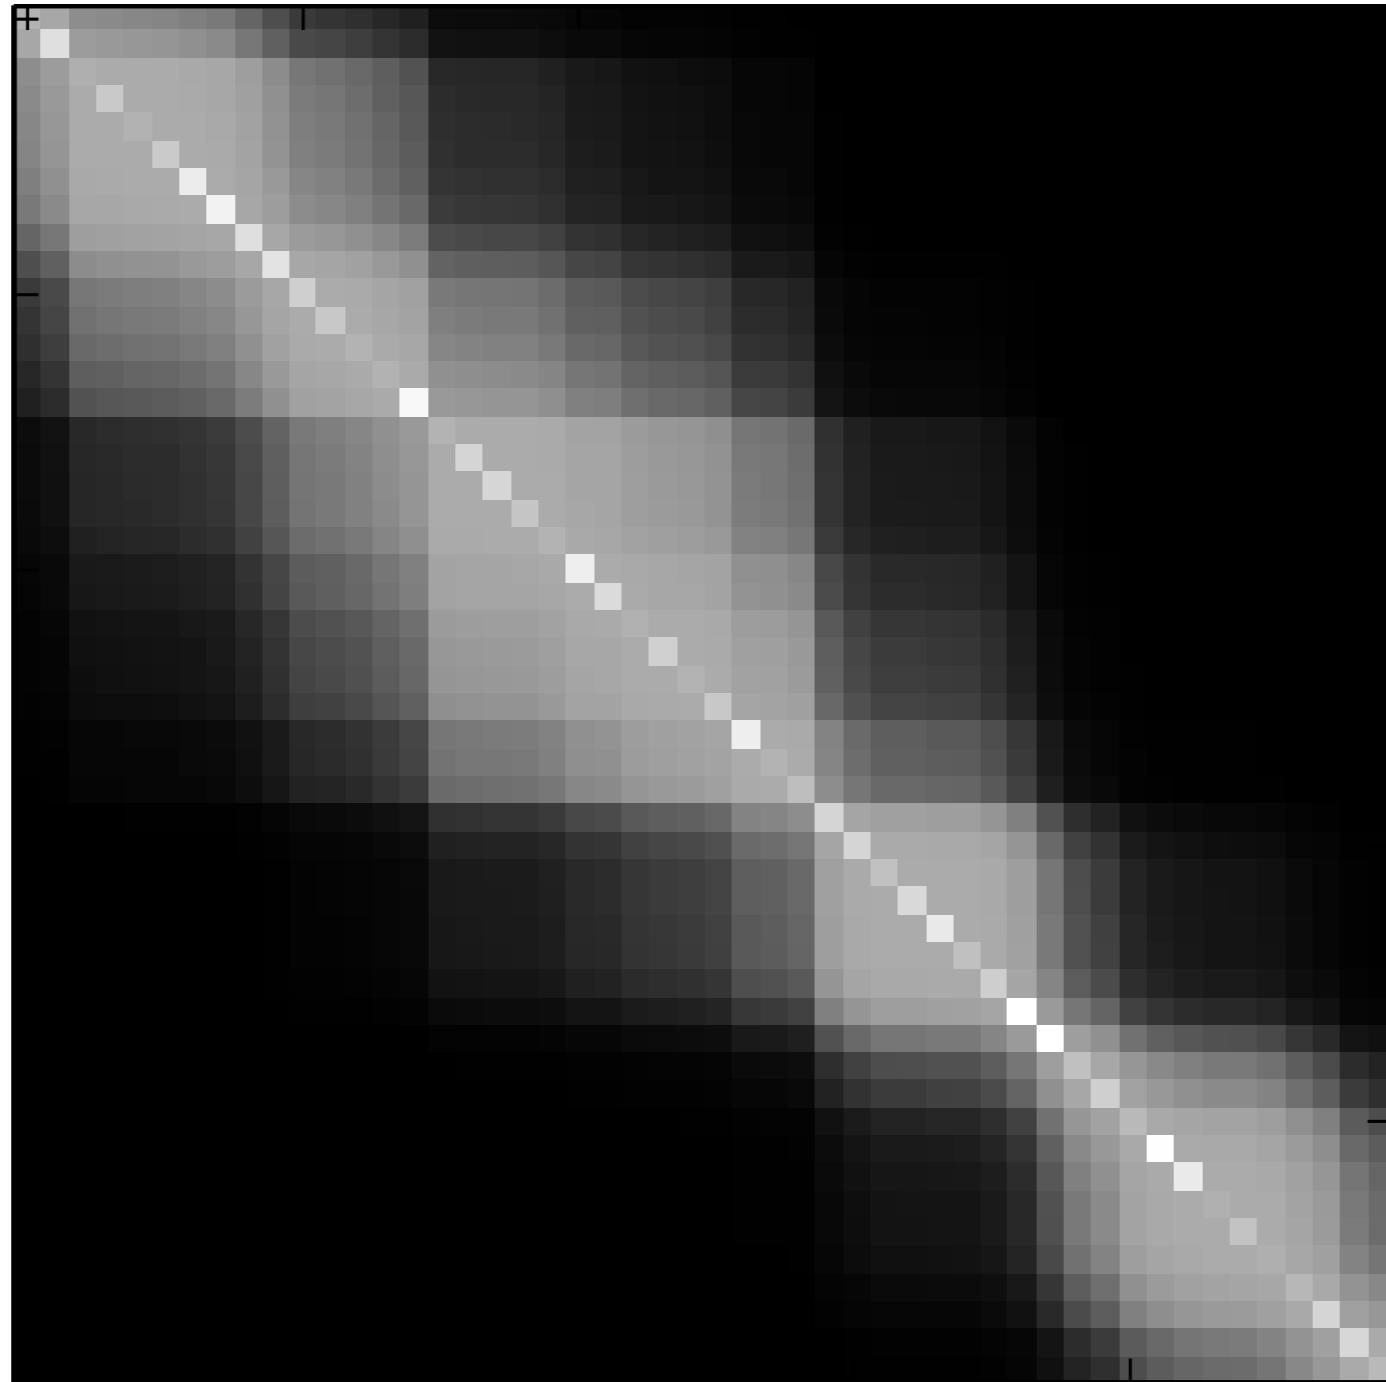
---

$$y = m x + b$$



The **true** covariance of the **observations**.

---



Let's **assume** that the **noise** is...

---

independent



$$\log p(\mathbf{y} | \mathbf{x}, \sigma, \theta) = -\frac{1}{2} \sum_{n=1}^N \left[ \frac{[y_n - f_{\theta}(x_n)]^2}{\sigma_n^2} + \log 2\pi\sigma_n^2 \right]$$



Gaussian with  
known variance

Or equivalently...

---

$$\log p(\mathbf{y} \mid \mathbf{x}, \boldsymbol{\sigma}, \boldsymbol{\theta}) = -\frac{1}{2} \mathbf{r}^T \mathbf{C}^{-1} \mathbf{r} - \frac{1}{2} \log \det \mathbf{C} - \frac{N}{2} \log 2\pi$$

Or equivalently...

---

data covariance



$$\log p(\mathbf{y} \mid \mathbf{x}, \boldsymbol{\sigma}, \boldsymbol{\theta}) = -\frac{1}{2} \mathbf{r}^T \mathbf{C}^{-1} \mathbf{r} - \frac{1}{2} \log \det \mathbf{C} - \frac{N}{2} \log 2\pi$$

Or equivalently...

---

data covariance



$$\log p(\mathbf{y} \mid \mathbf{x}, \boldsymbol{\sigma}, \boldsymbol{\theta}) = -\frac{1}{2} \mathbf{r}^T \mathbf{C}^{-1} \mathbf{r} - \frac{1}{2} \log \det \mathbf{C} - \frac{N}{2} \log 2\pi$$

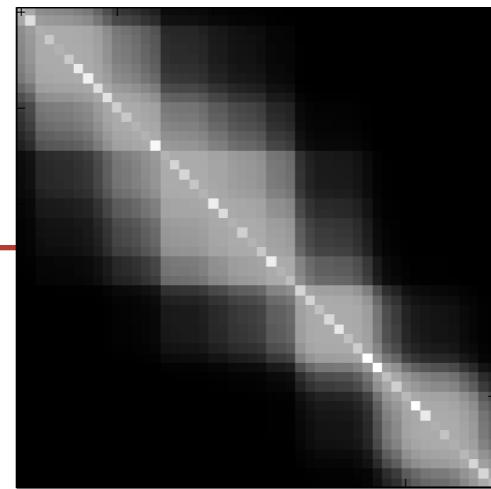
residual vector

$$\mathbf{r} = [ y_1 - f_{\boldsymbol{\theta}}(x_1) \quad y_2 - f_{\boldsymbol{\theta}}(x_2) \quad \cdots \quad y_n - f_{\boldsymbol{\theta}}(x_n) ]^T$$

Or equivalently...

---

data covariance



$$\log p(\mathbf{y} \mid \mathbf{x}, \boldsymbol{\sigma}, \boldsymbol{\theta}) = -\frac{1}{2} \mathbf{r}^T \mathbf{C}^{-1} \mathbf{r} - \frac{1}{2} \log \det \mathbf{C} - \frac{N}{2} \log 2\pi$$



residual vector

$$\mathbf{r} = [ y_1 - f_{\boldsymbol{\theta}}(x_1) \quad y_2 - f_{\boldsymbol{\theta}}(x_2) \quad \cdots \quad y_n - f_{\boldsymbol{\theta}}(x_n) ]^T$$

## Linear least-squares.

---

$$\begin{bmatrix} m \\ b \end{bmatrix} = \mathbf{S} \mathbf{A}^T \mathbf{C}^{-1} \mathbf{y} \quad \mathbf{S} = [\mathbf{A}^T \mathbf{C}^{-1} \mathbf{A}]^{-1}$$

maximum likelihood & posterior covariance  
in this case only mean of posterior

---

$$\mathbf{A} = \begin{bmatrix} x_1 & 1 \\ x_2 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \quad \mathbf{C} = \begin{bmatrix} \sigma_1^2 & 0 & \cdots & 0 \\ 0 & \sigma_2^2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \sigma_n^2 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

## Linear least-squares.

---

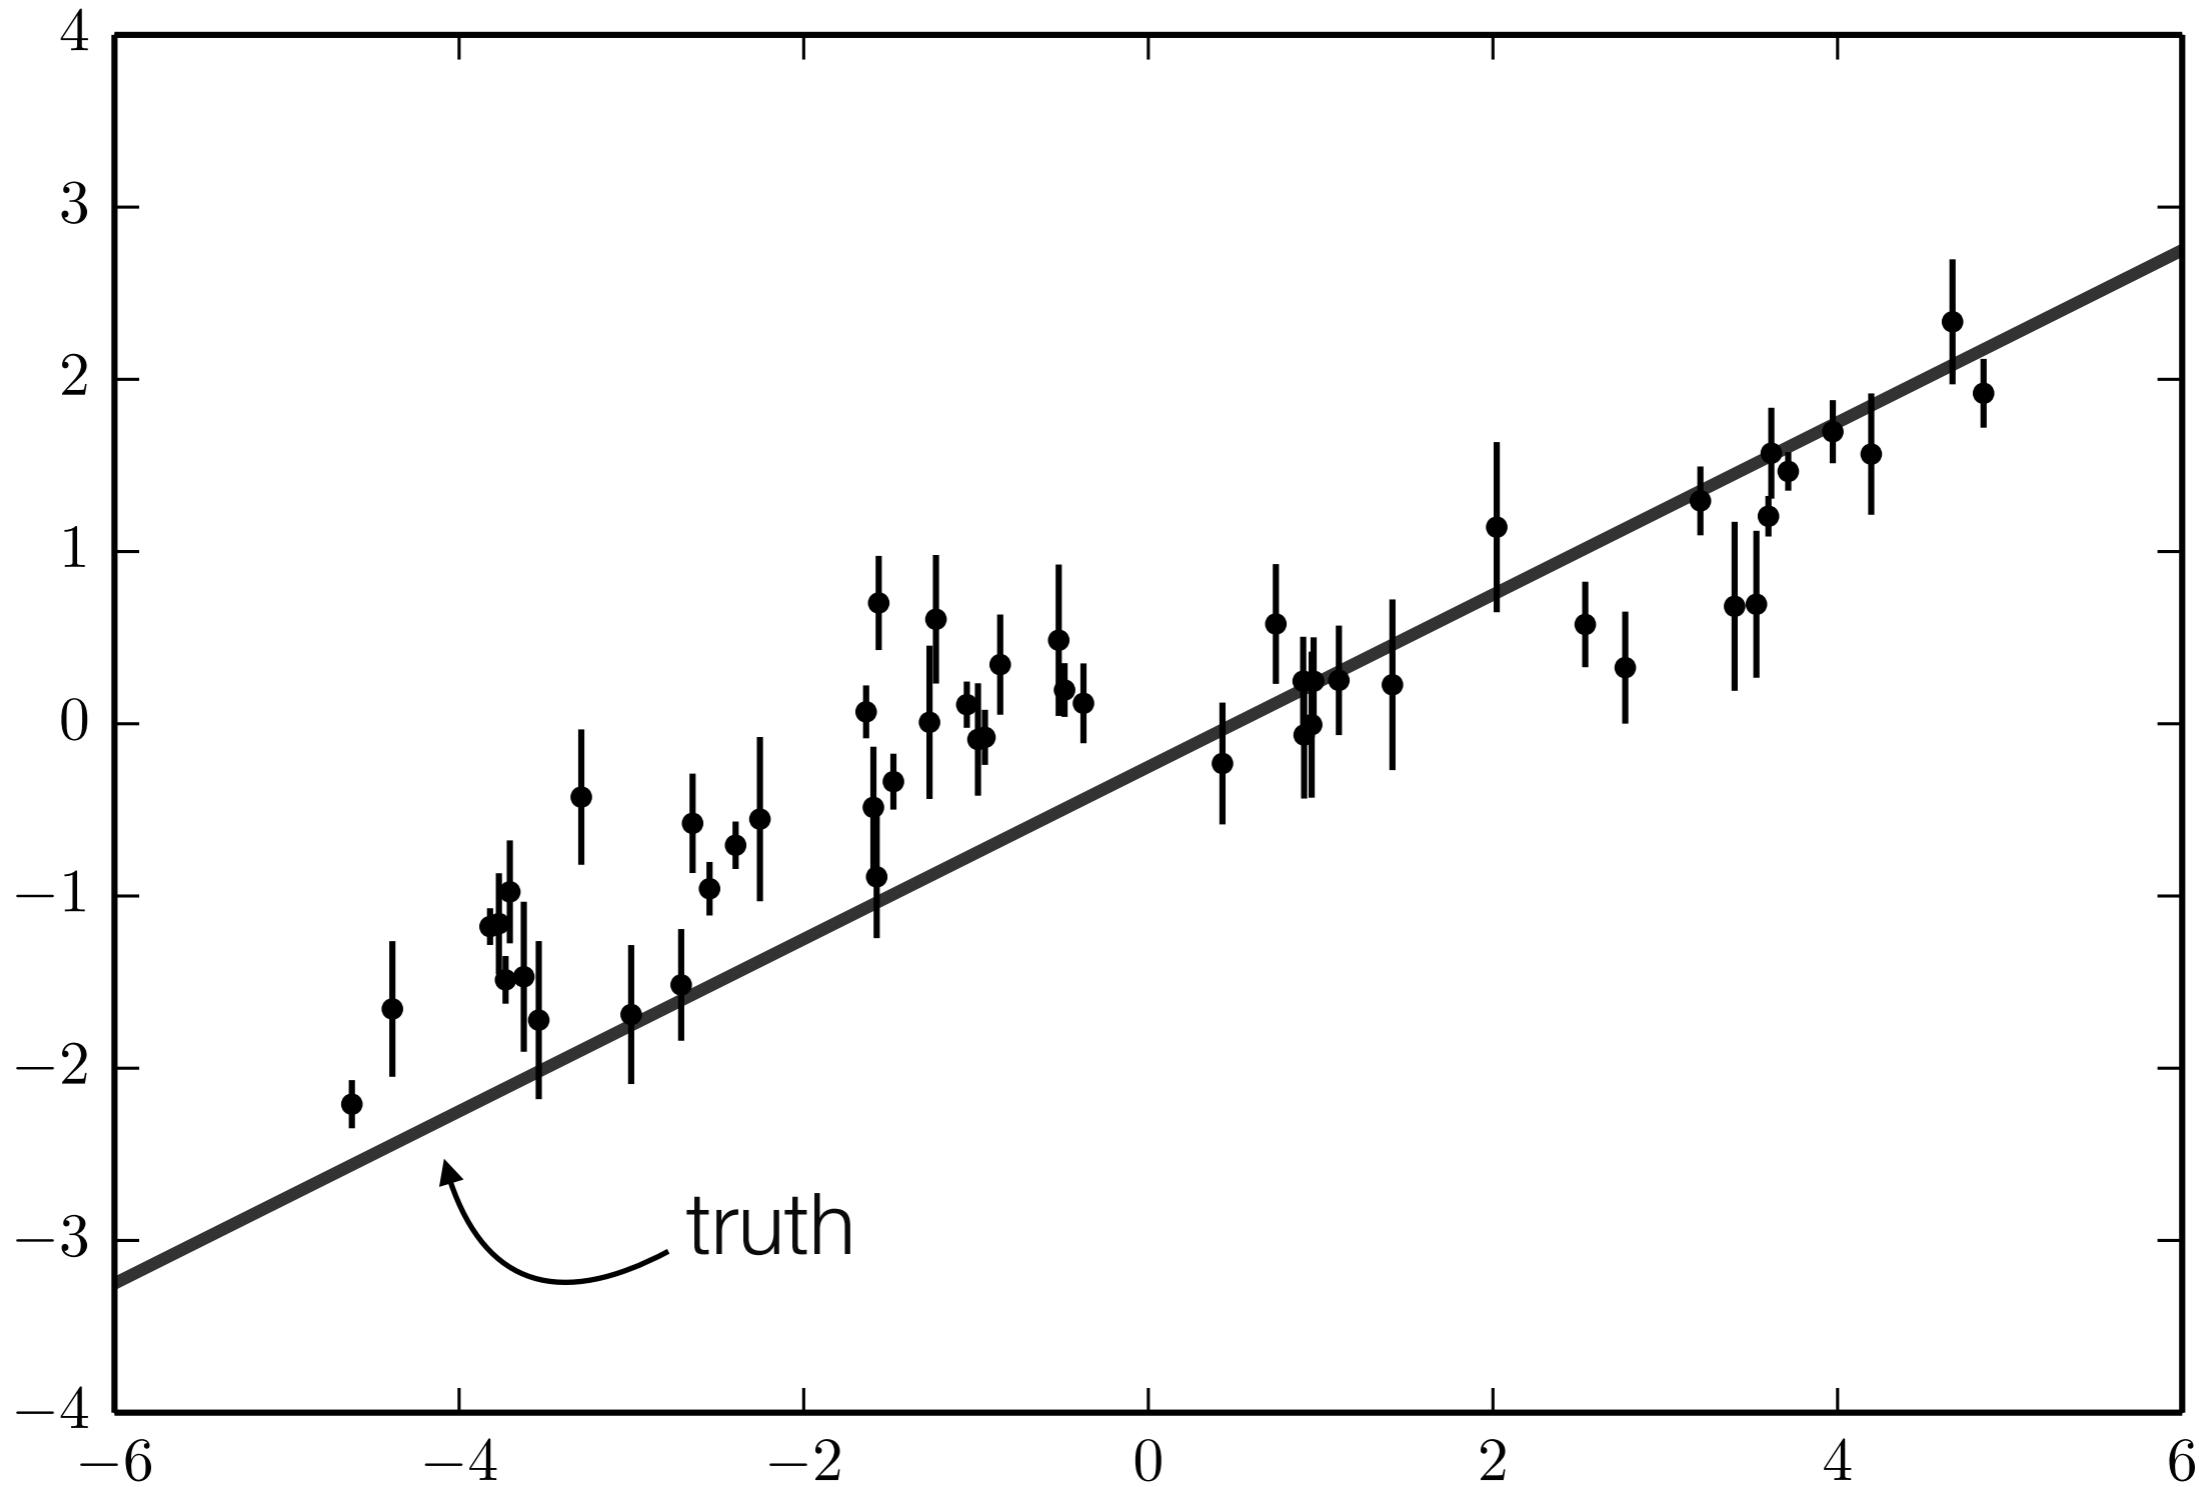
$$\begin{bmatrix} m \\ b \end{bmatrix} = \mathbf{S} \mathbf{A}^T \mathbf{C}^{-1} \mathbf{y} \quad \mathbf{S} = [\mathbf{A}^T \mathbf{C}^{-1} \mathbf{A}]^{-1}$$

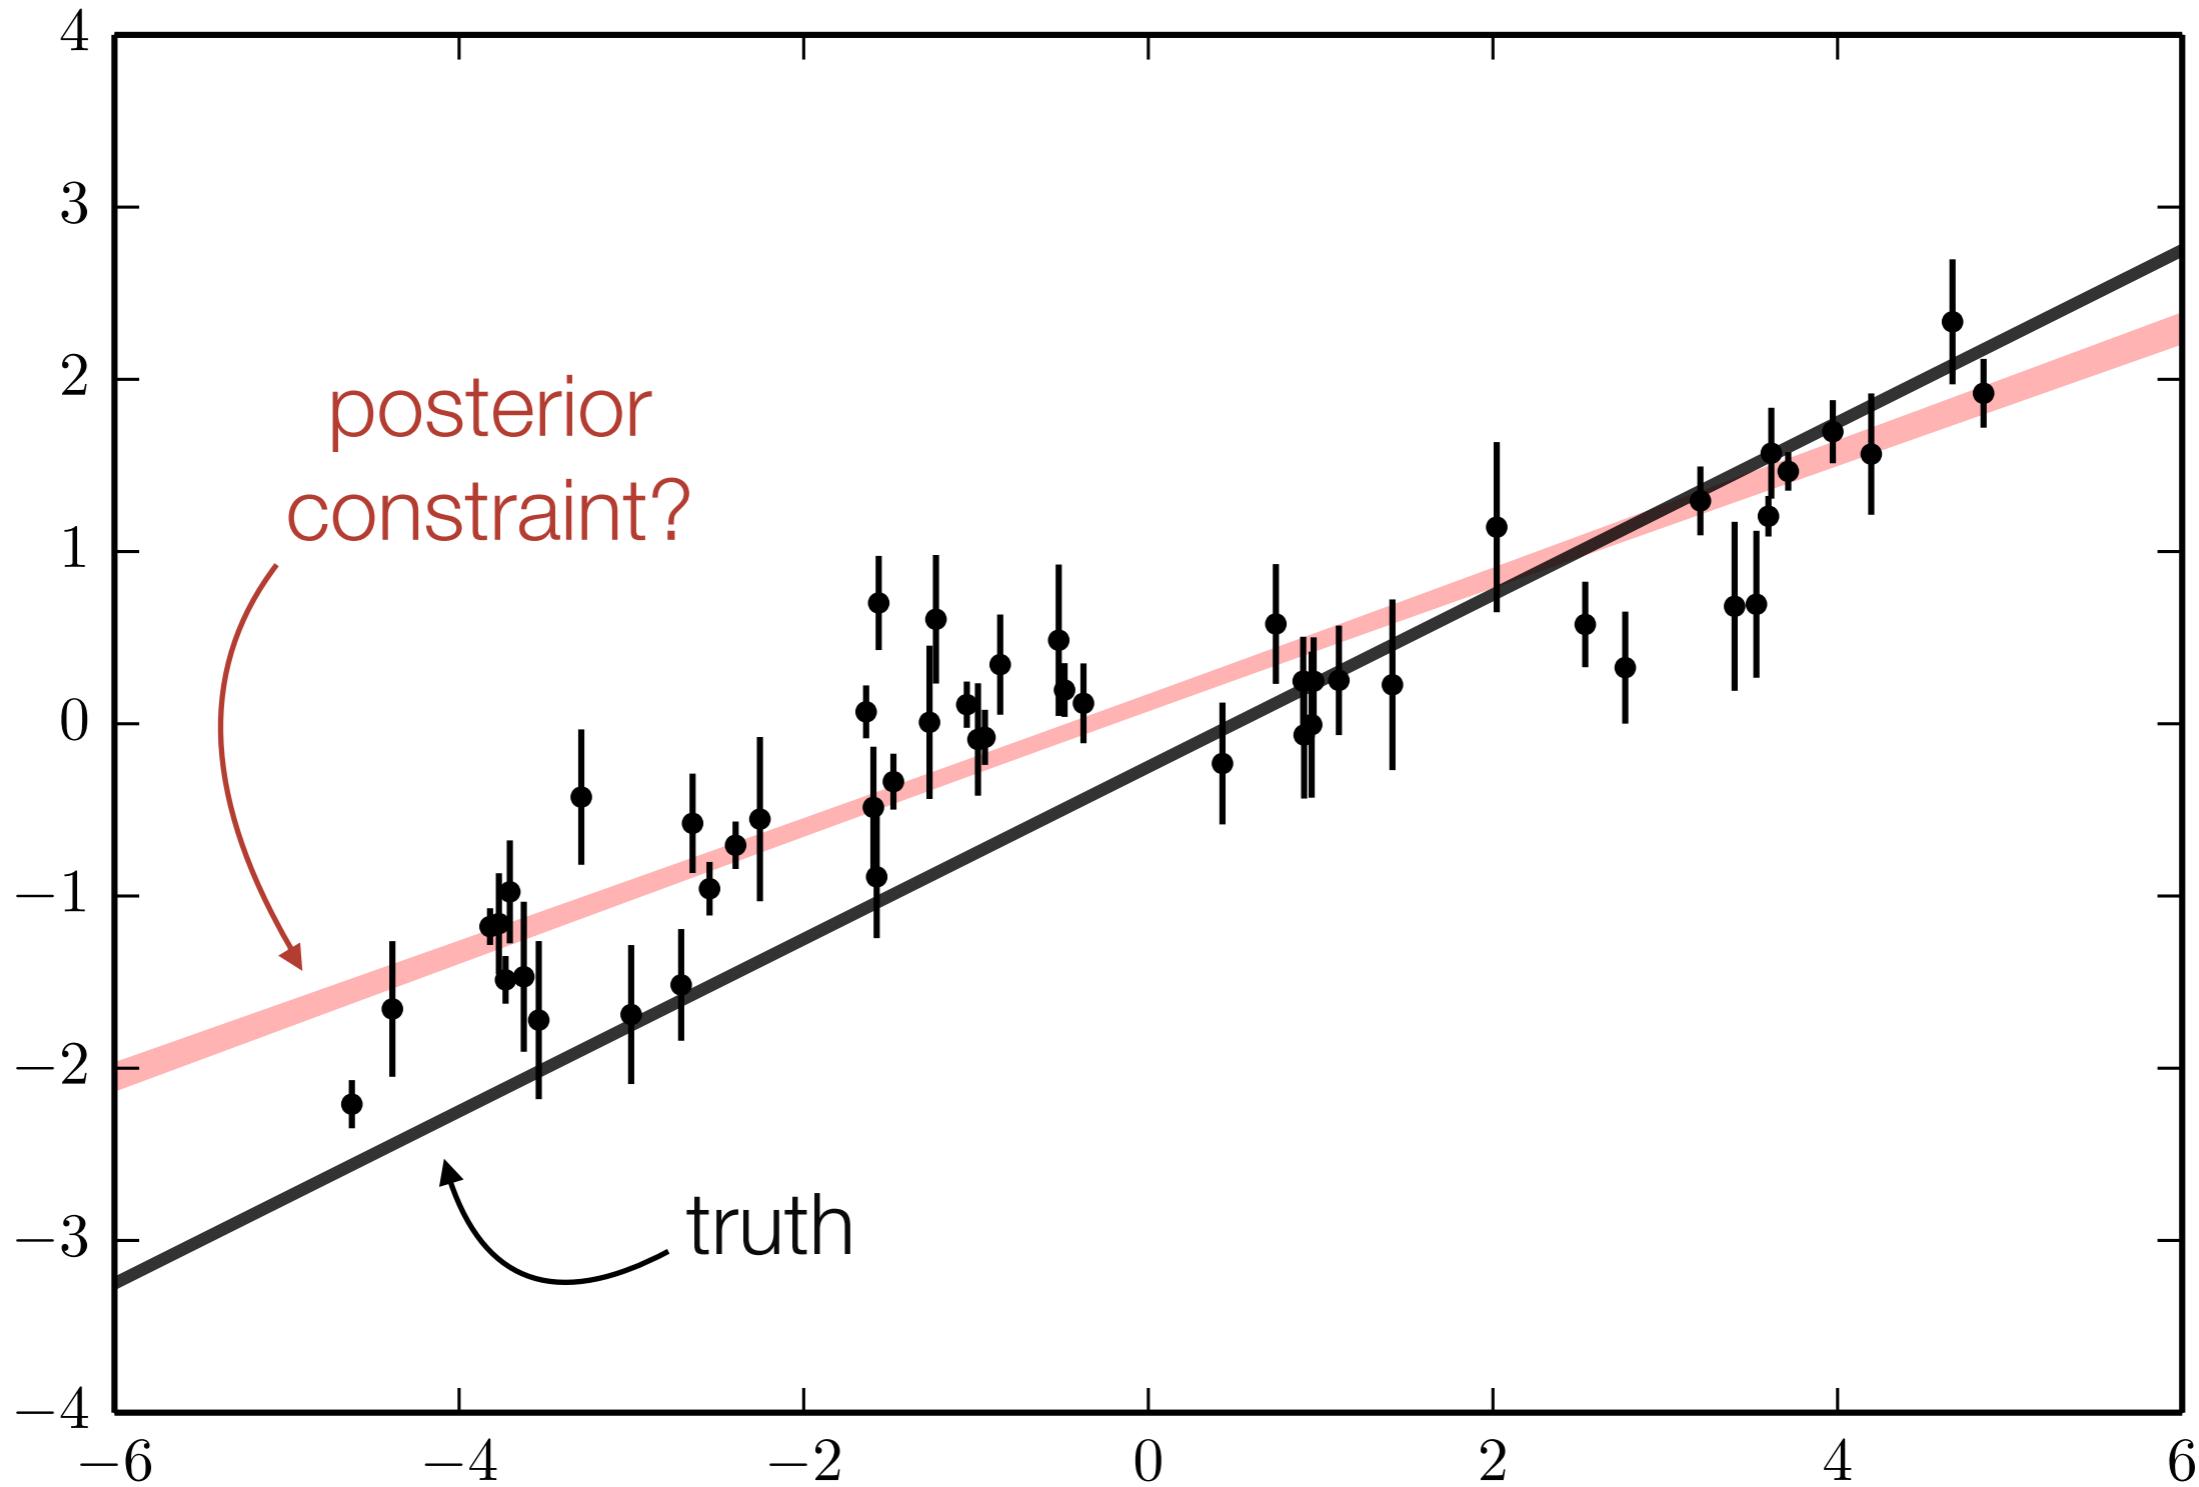
maximum likelihood &  
in this case only mean of posterior

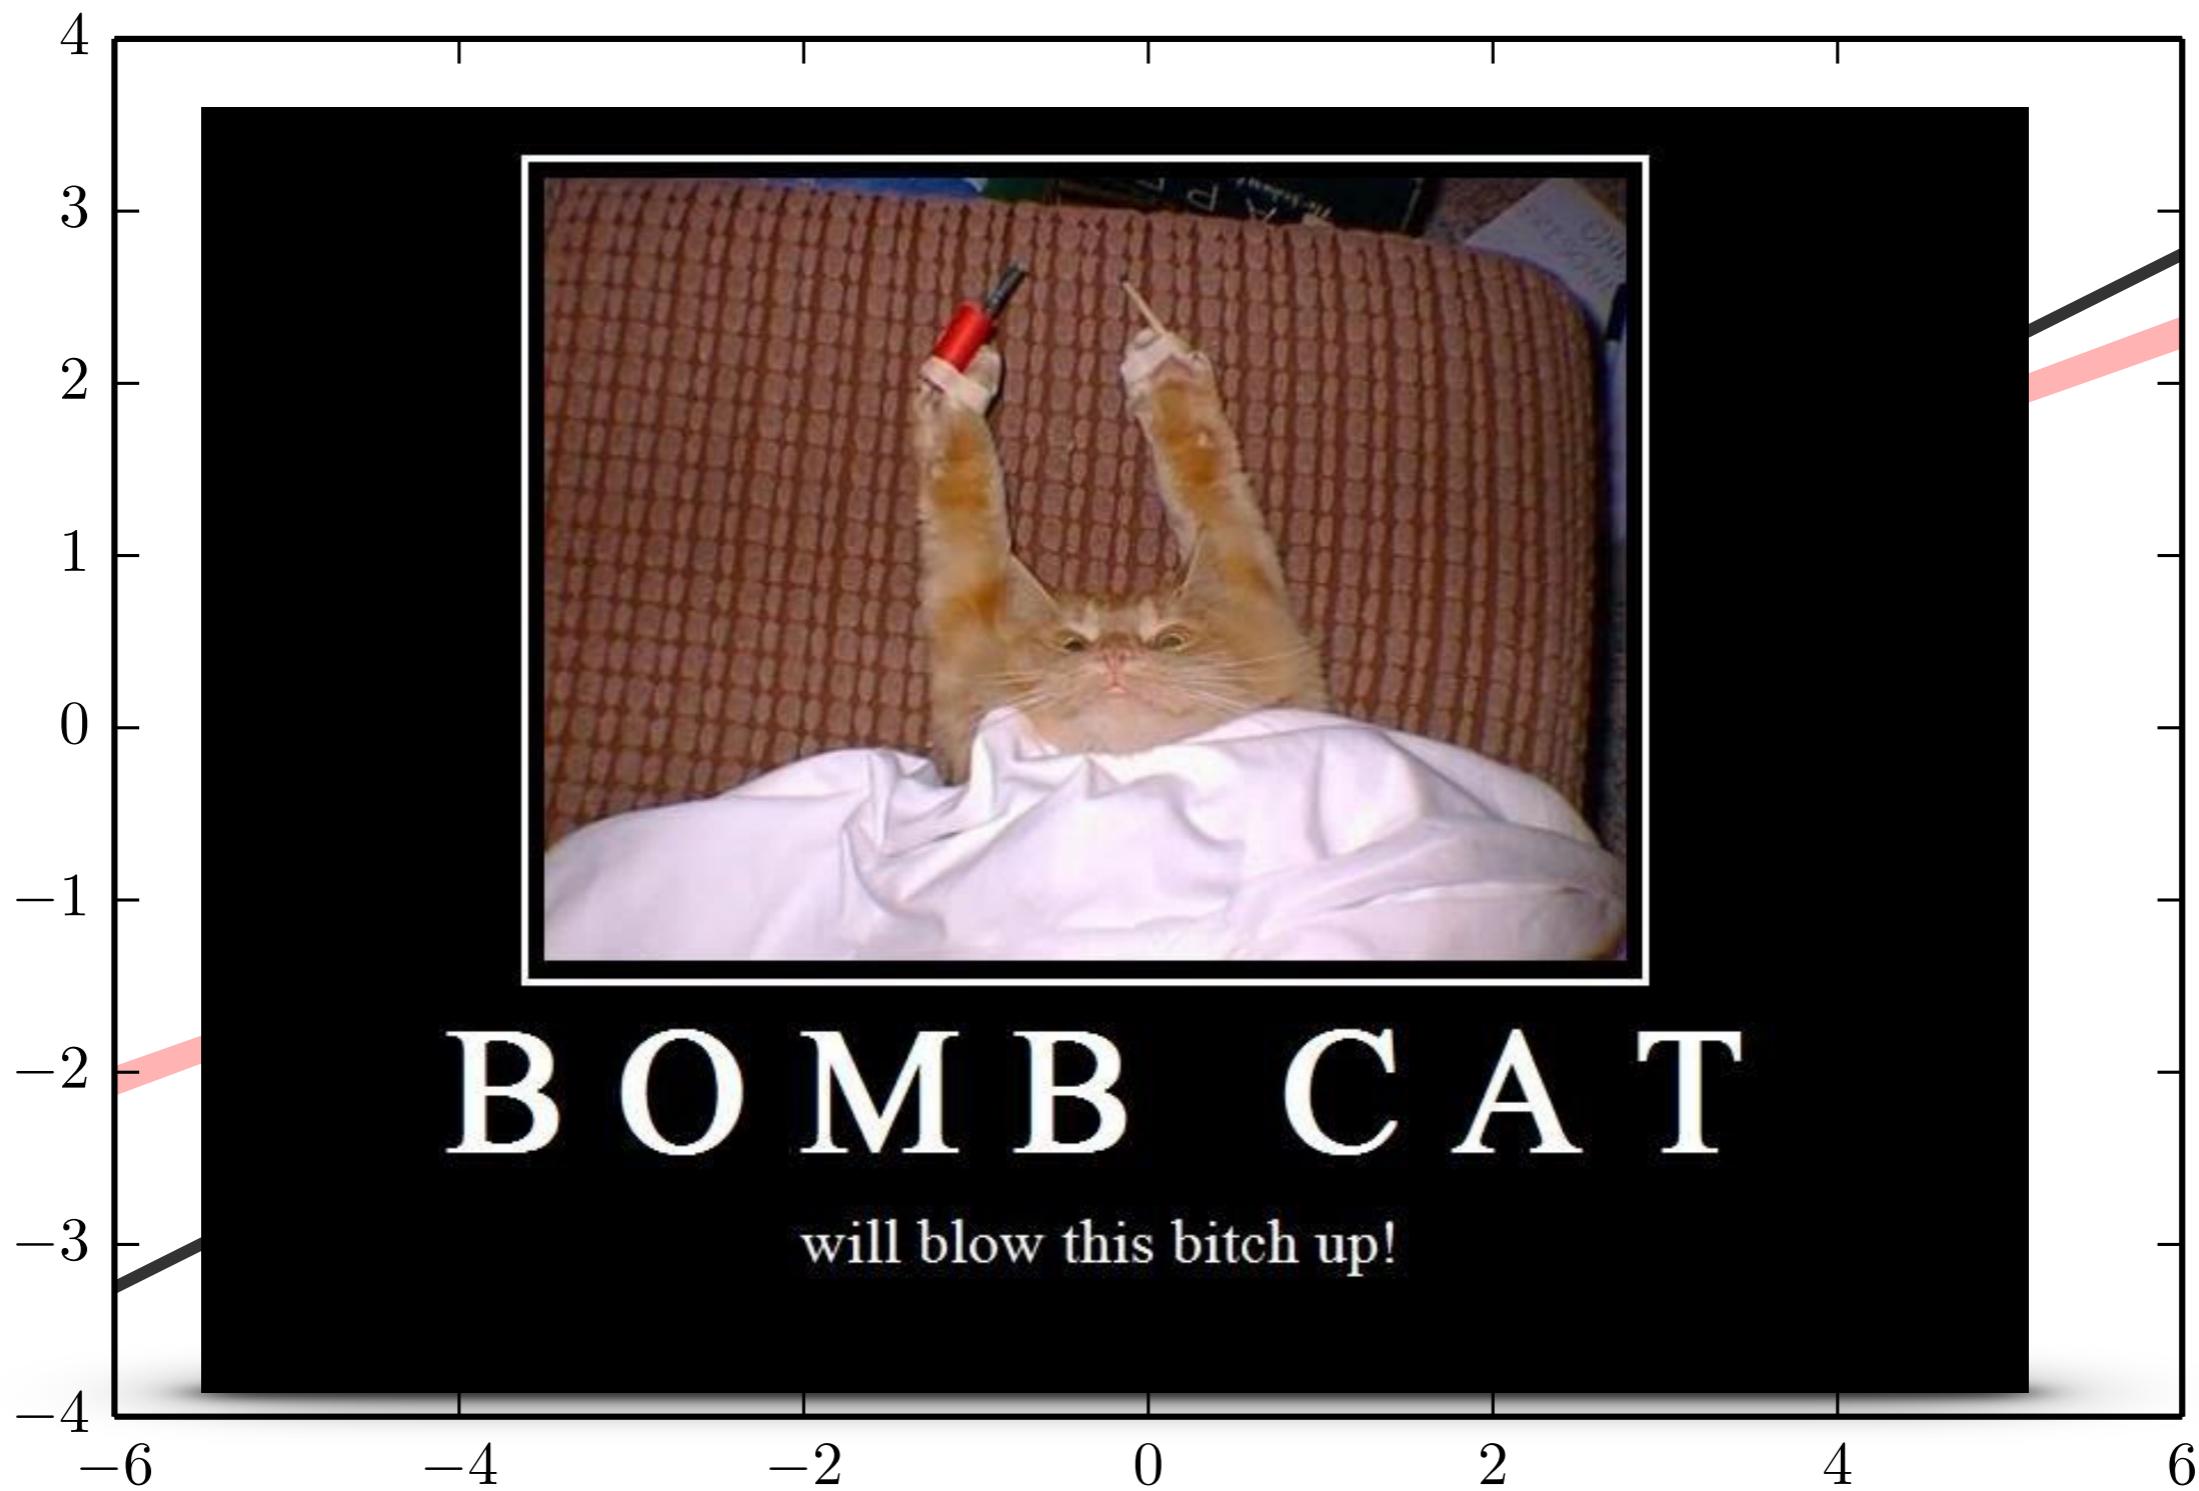
posterior covariance

assuming uniform priors

$$\mathbf{A} = \begin{bmatrix} x_1 & 1 \\ x_2 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \quad \mathbf{C} = \begin{bmatrix} \sigma_1^2 & 0 & \cdots & 0 \\ 0 & \sigma_2^2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \sigma_n^2 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

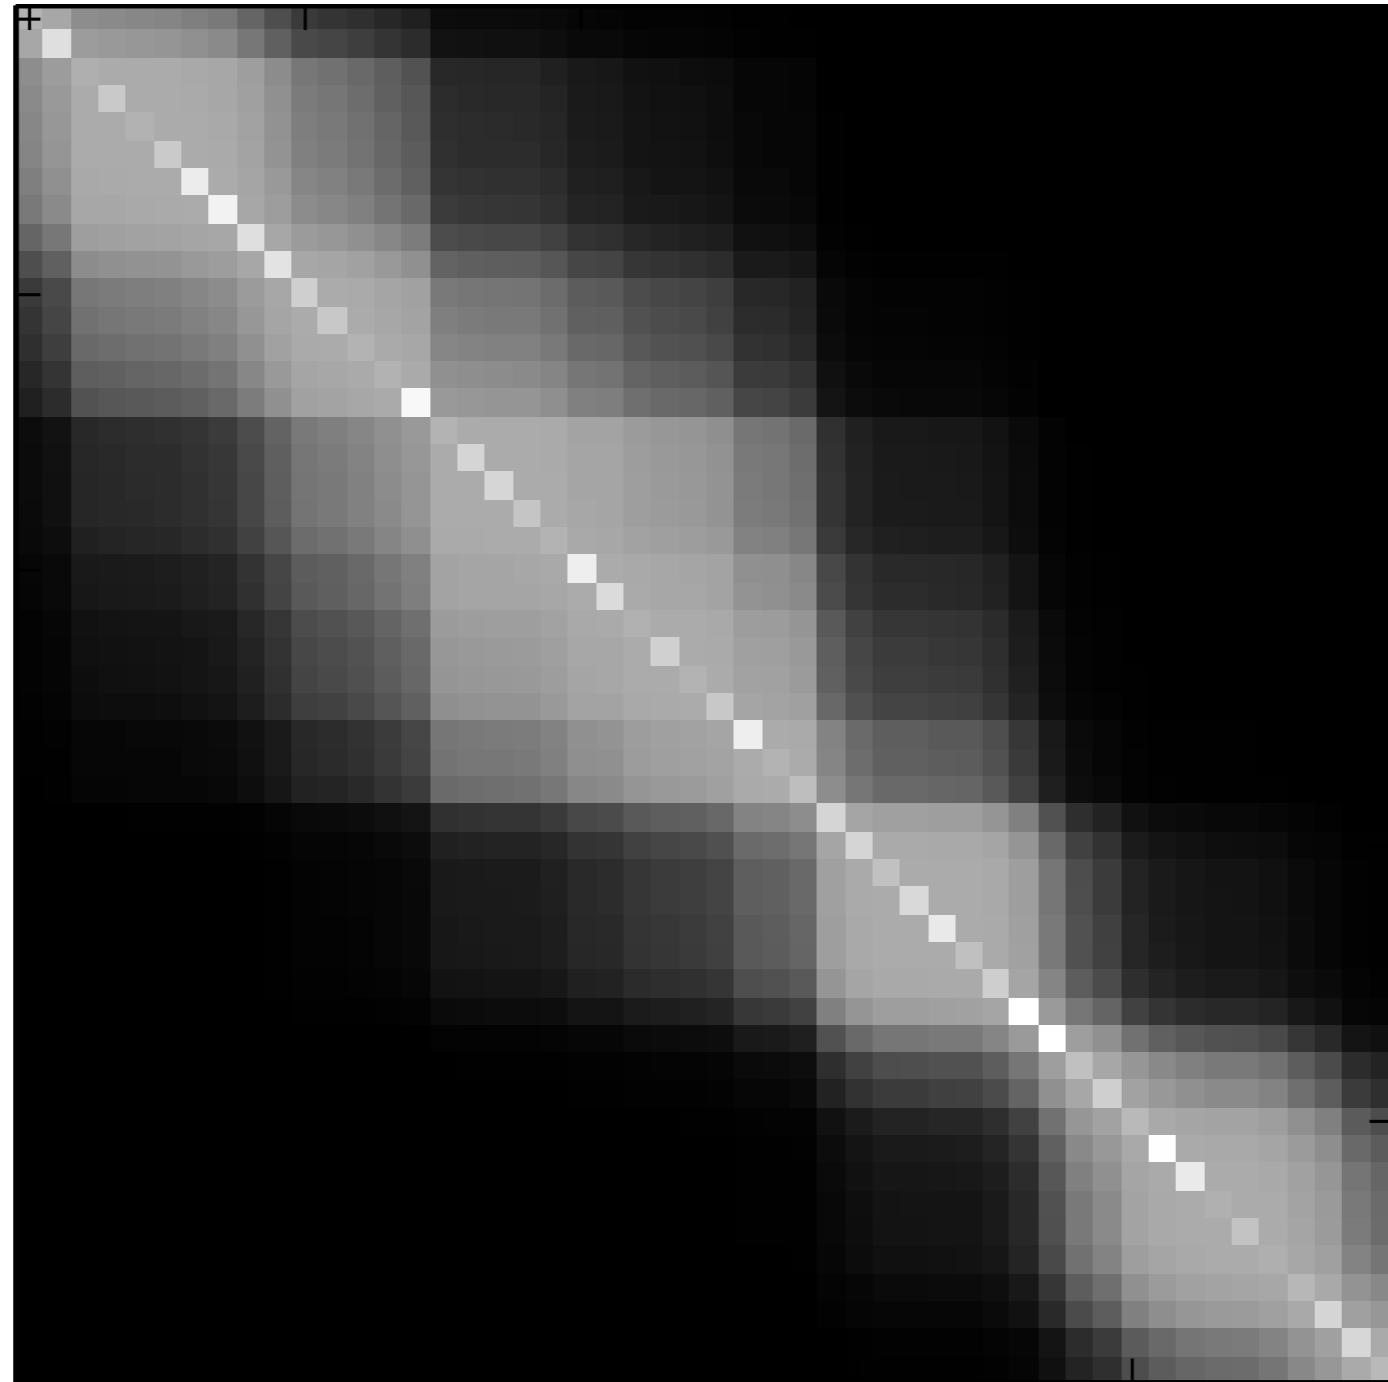




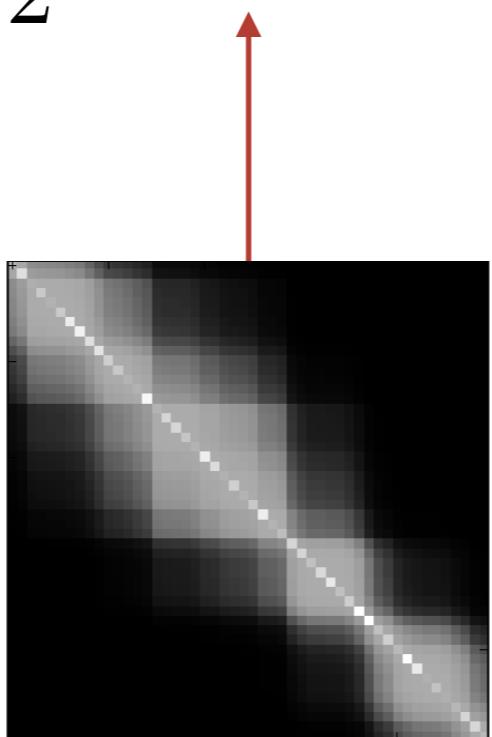


But we **know** the **true** covariance matrix.

---



$$\log p(\mathbf{y} \mid \mathbf{x}, \boldsymbol{\sigma}, \boldsymbol{\theta}) = -\frac{1}{2} \mathbf{r}^T \mathbf{C}^{-1} \mathbf{r} - \frac{1}{2} \log \det \mathbf{C} - \frac{N}{2} \log 2 \pi$$



## Linear least-squares.

---

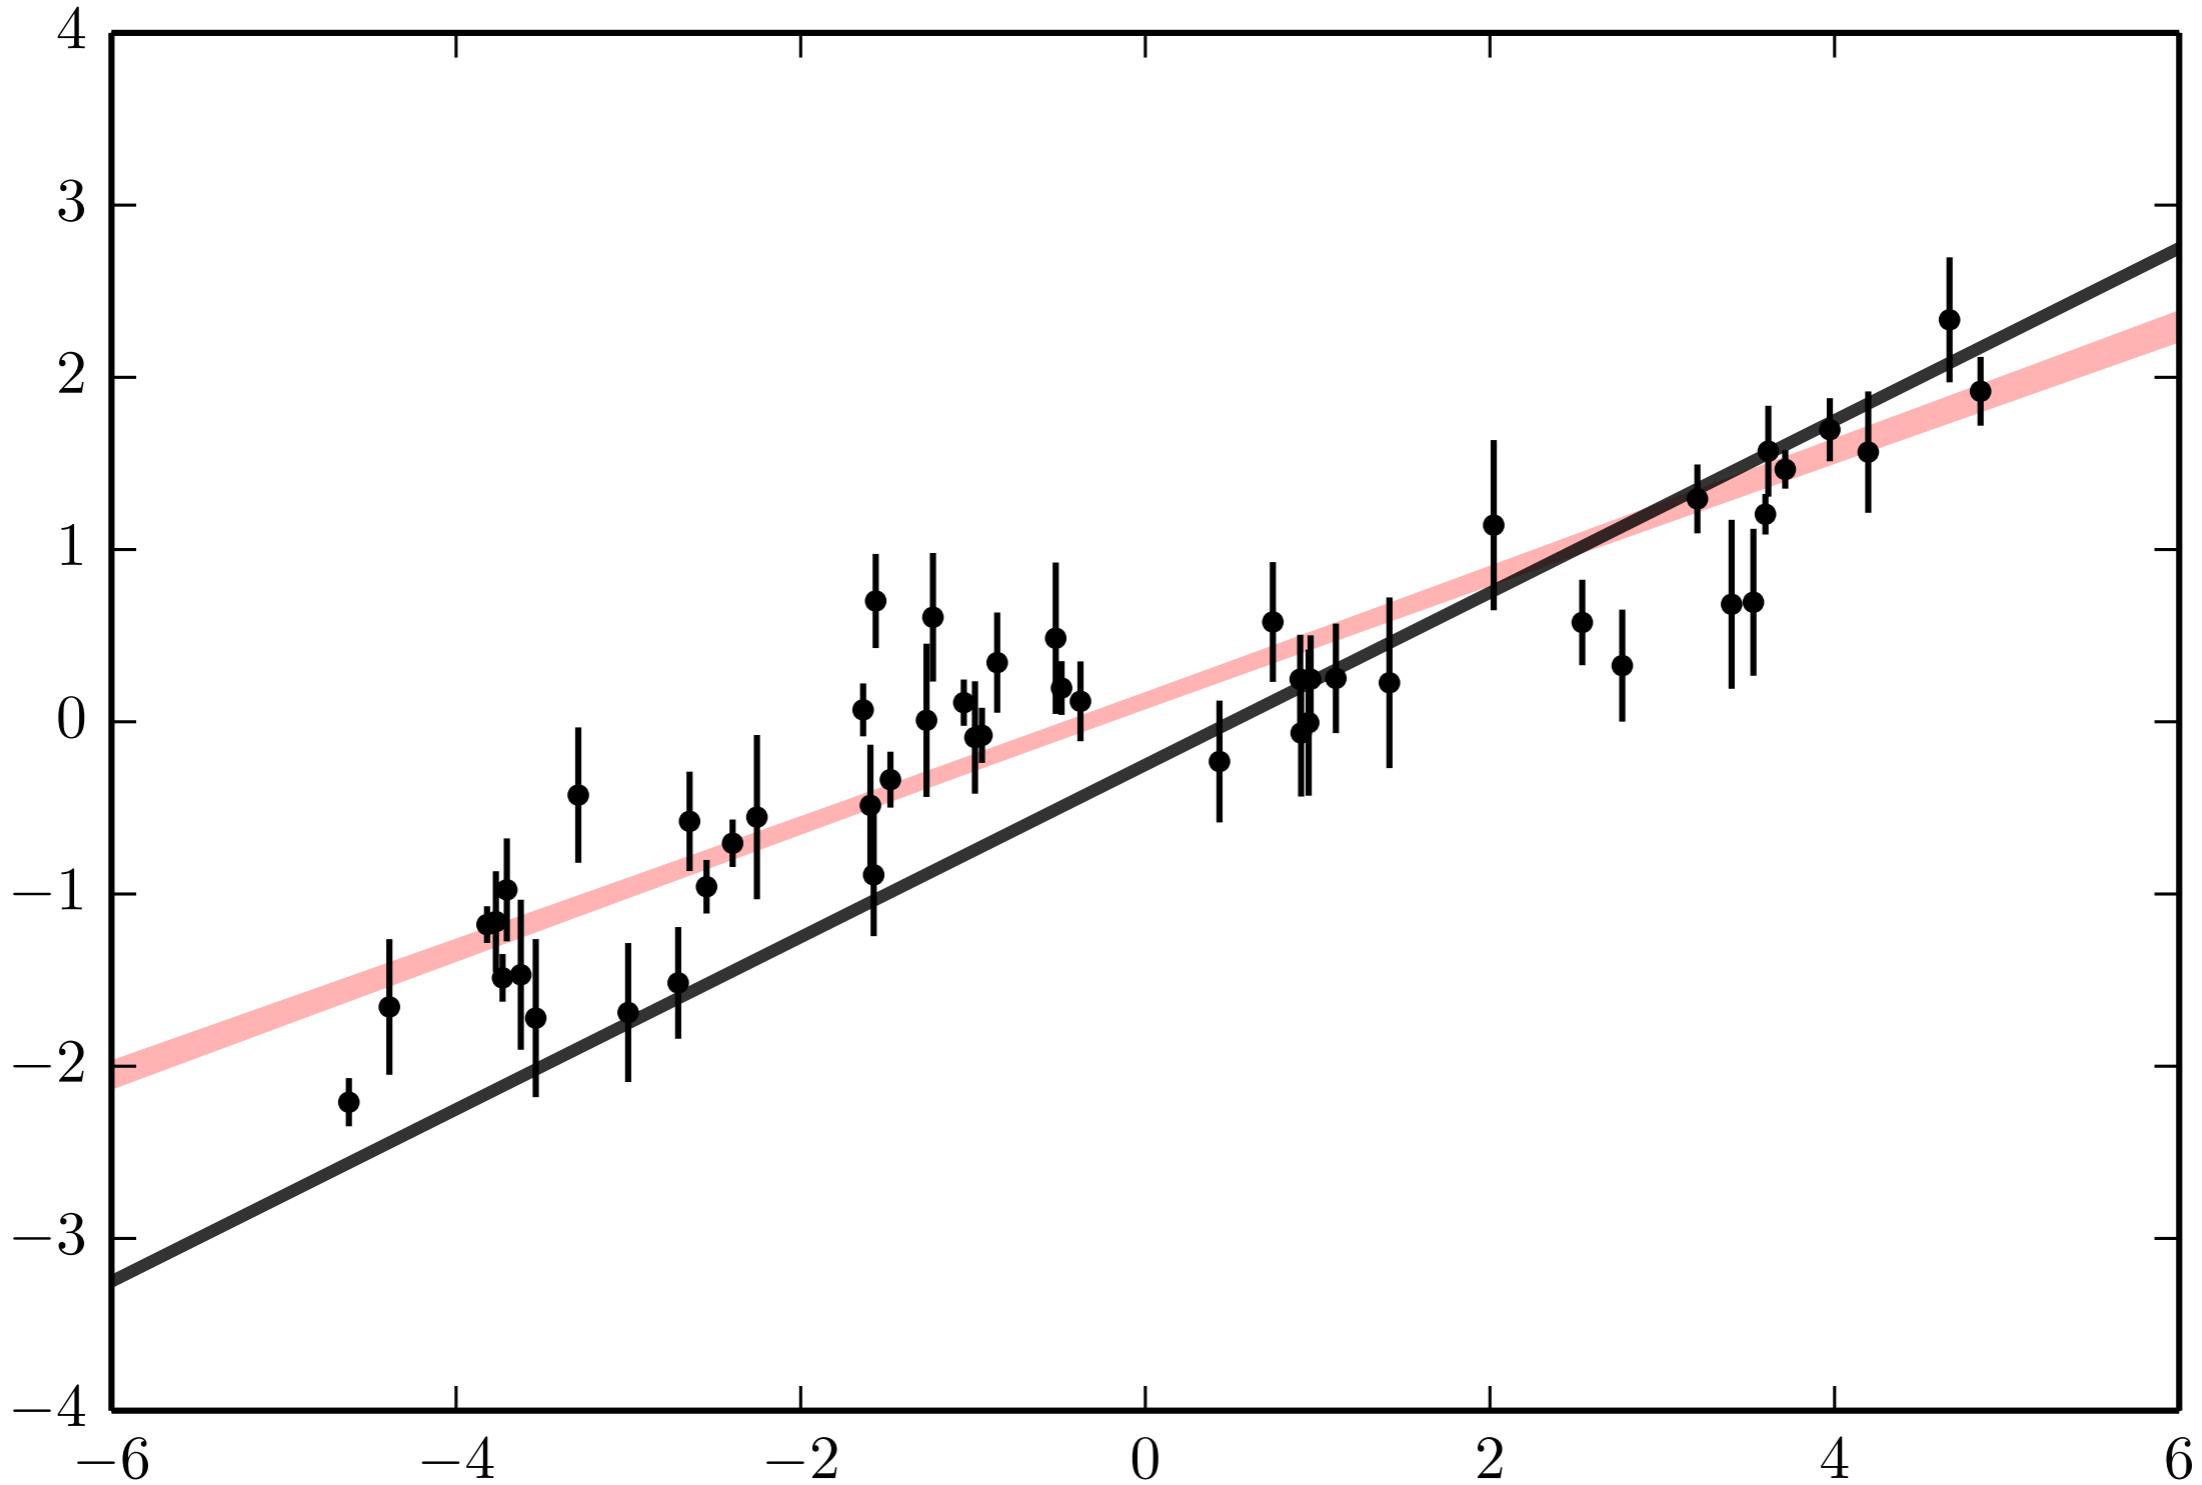
$$\begin{bmatrix} m \\ b \end{bmatrix} = \mathbf{S} \mathbf{A}^T \mathbf{C}^{-1} \mathbf{y} \quad \mathbf{S} = [\mathbf{A}^T \mathbf{C}^{-1} \mathbf{A}]^{-1}$$

maximum likelihood &  
in this case only mean of posterior

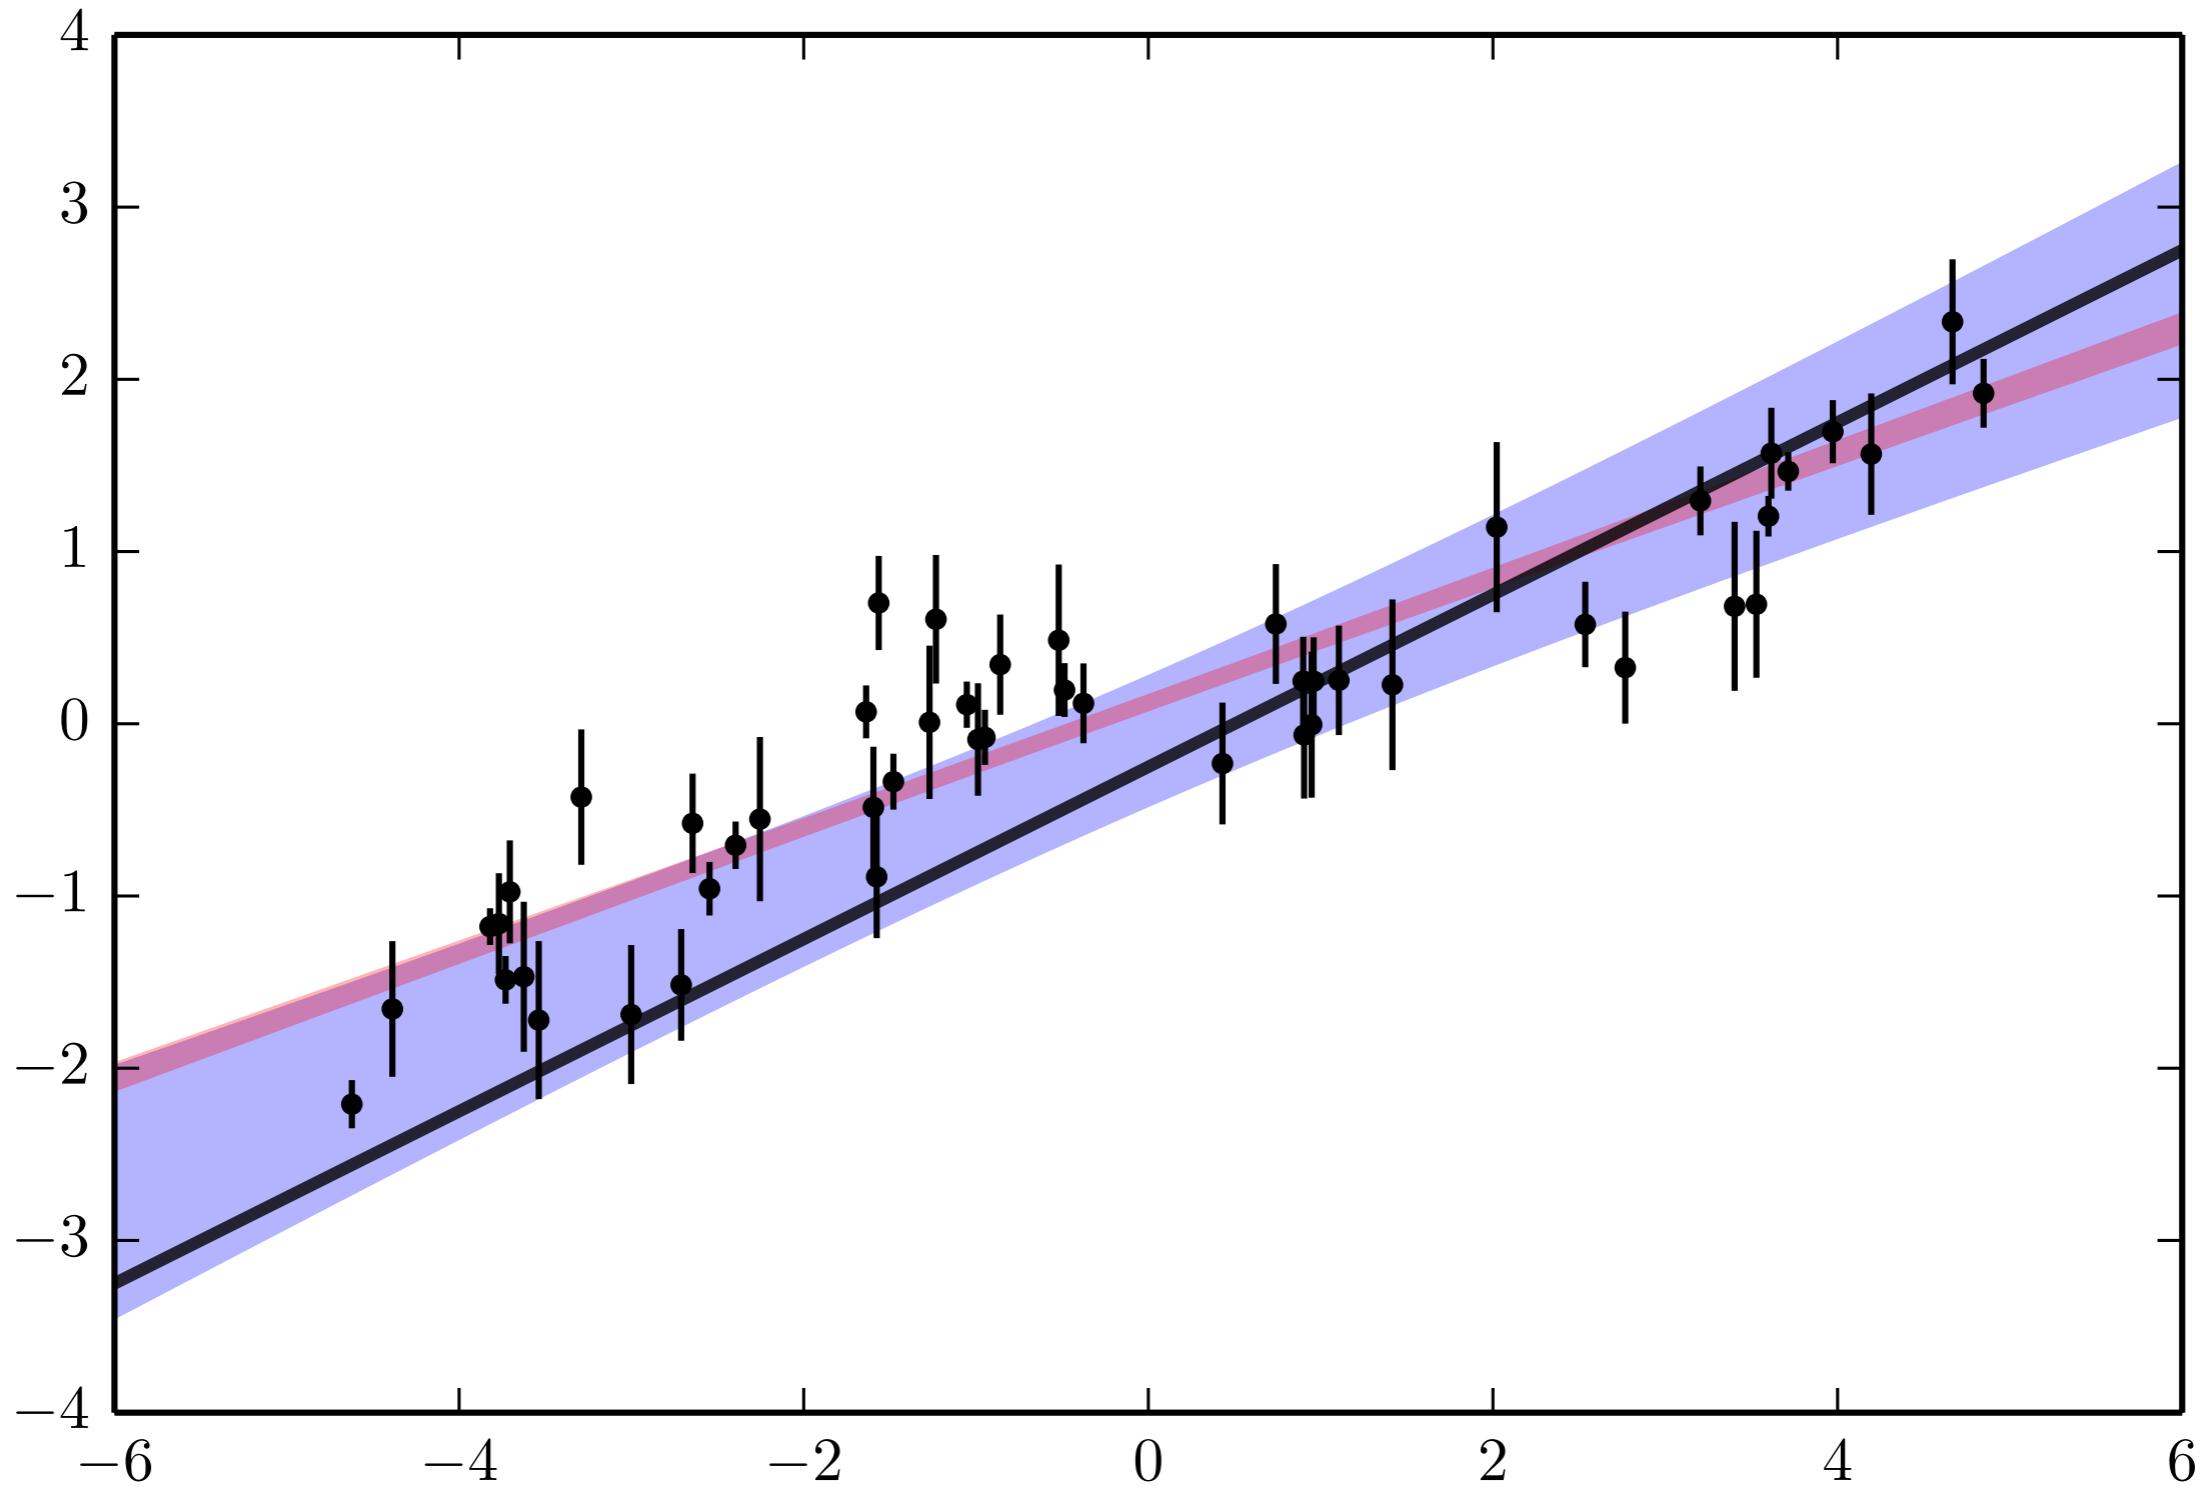
posterior covariance

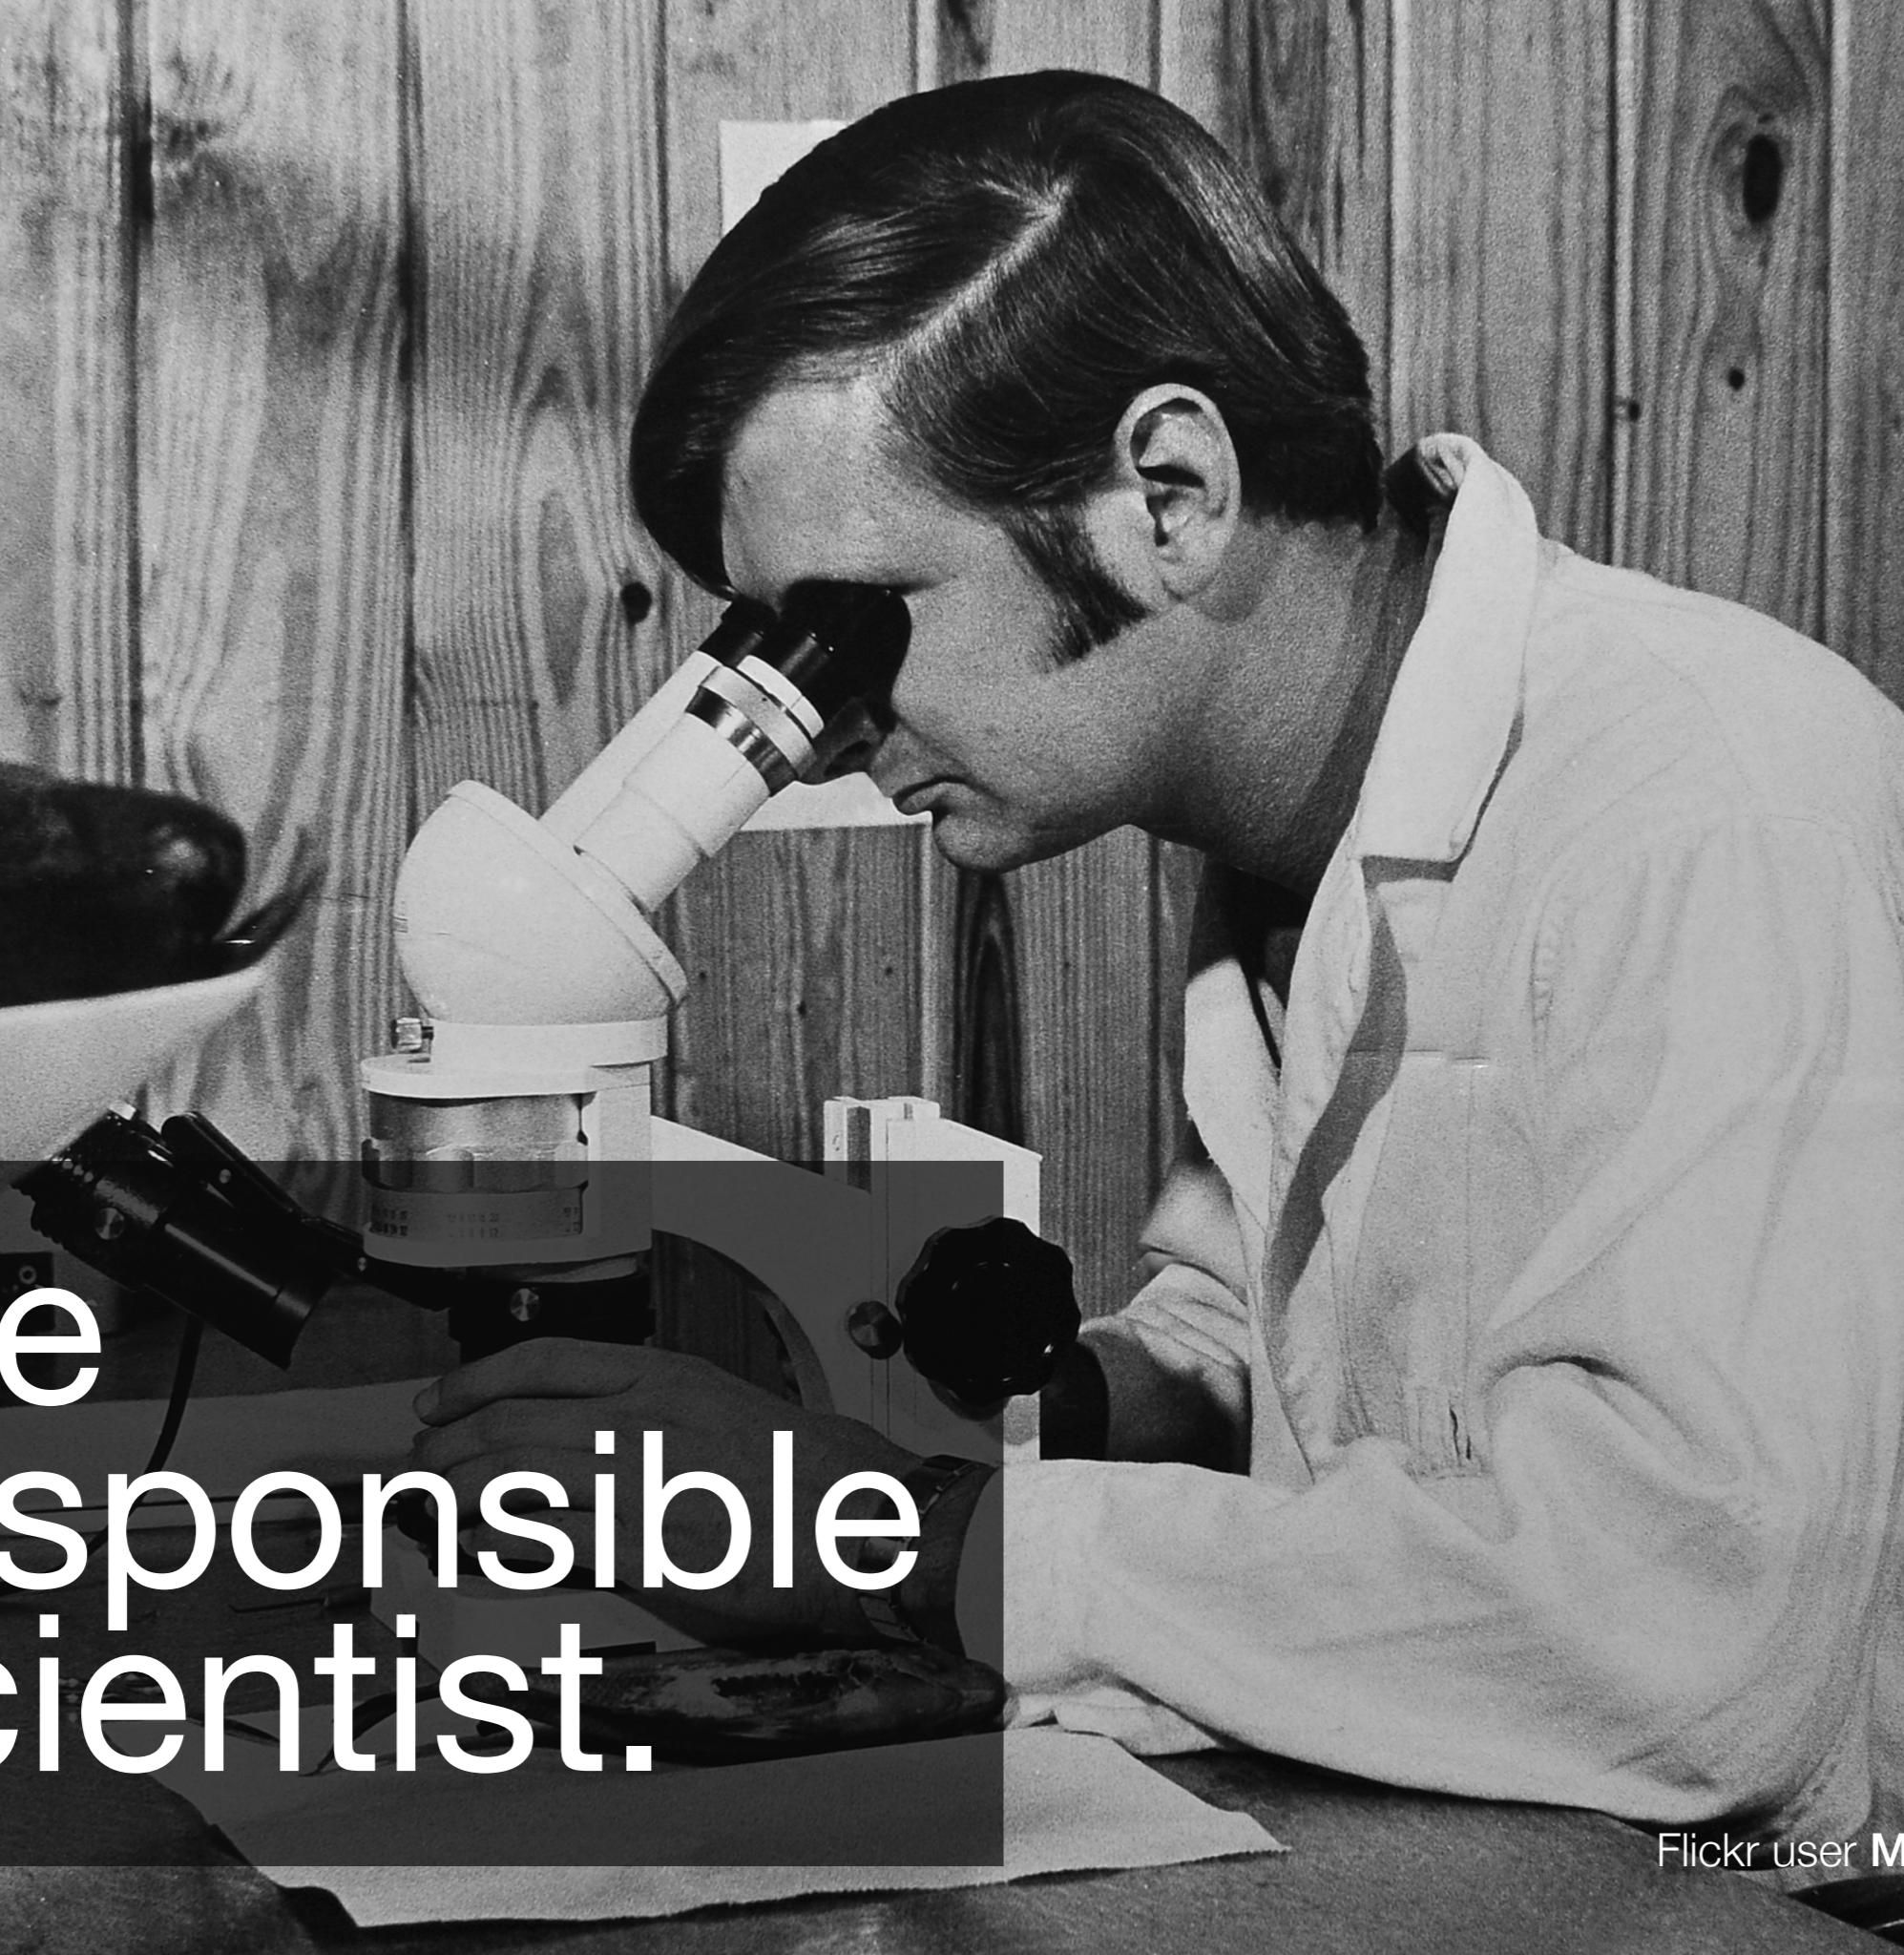
$$\mathbf{A} = \begin{bmatrix} x_1 & 1 \\ x_2 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \quad \mathbf{C} = \begin{bmatrix} \text{image} \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

Before.



After.





the  
responsible  
scientist.



Flickr user MyFWCmedia

So... we're **finished**, right?

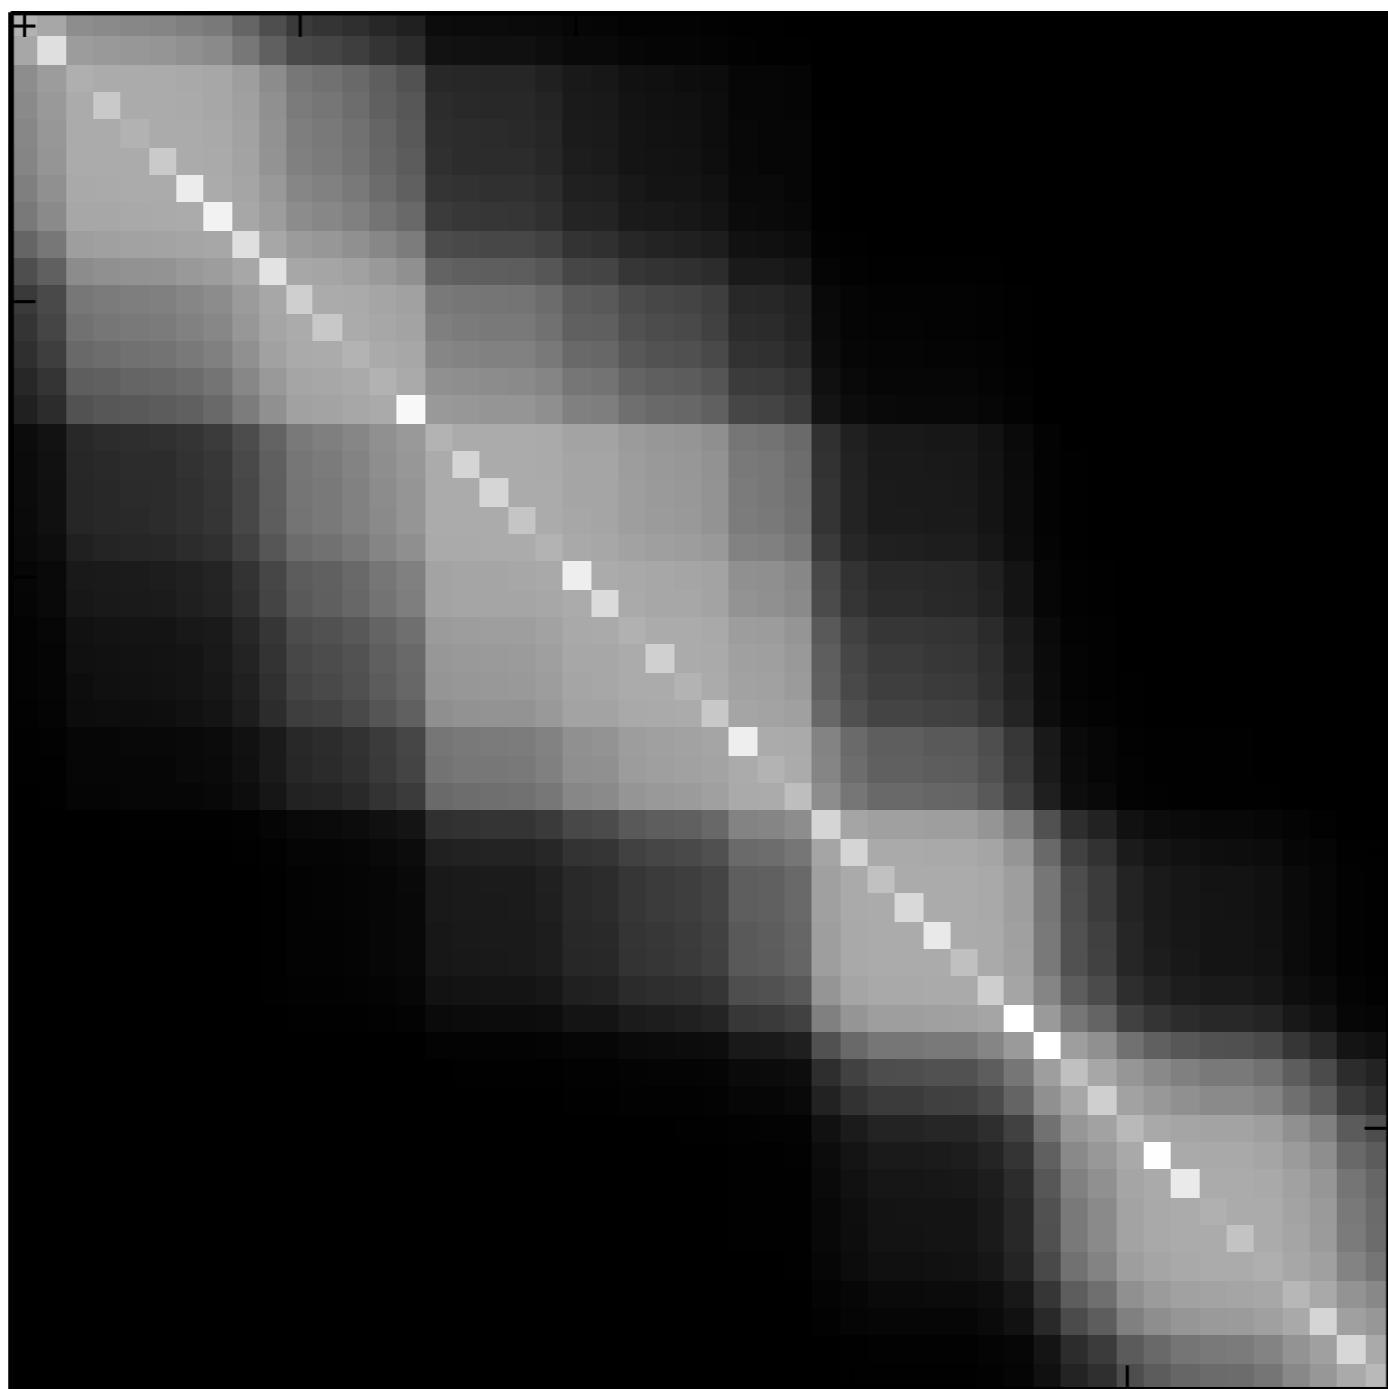
---

In The Real World™, we **never** know the noise.

---

Just gotta **model** it!

---



## Model it!

---

$$\log p(\mathbf{y} \mid \mathbf{x}, \boldsymbol{\sigma}, \boldsymbol{\theta}, \boldsymbol{\alpha}) = -\frac{1}{2} \mathbf{r}^T \mathbf{K}_{\boldsymbol{\alpha}}^{-1} \mathbf{r} - \frac{1}{2} \log \det \mathbf{K}_{\boldsymbol{\alpha}} + C$$

where

$$K_{ij} = \sigma_i^2 \delta_{ij} + k_{\boldsymbol{\alpha}}(\mathbf{x}_i, \mathbf{x}_j)$$

for example

$$k_{\boldsymbol{\alpha}}(\mathbf{x}_i, \mathbf{x}_j) = a^2 \exp\left(-\frac{[\mathbf{x}_i - \mathbf{x}_j]^2}{2 l^2}\right)$$

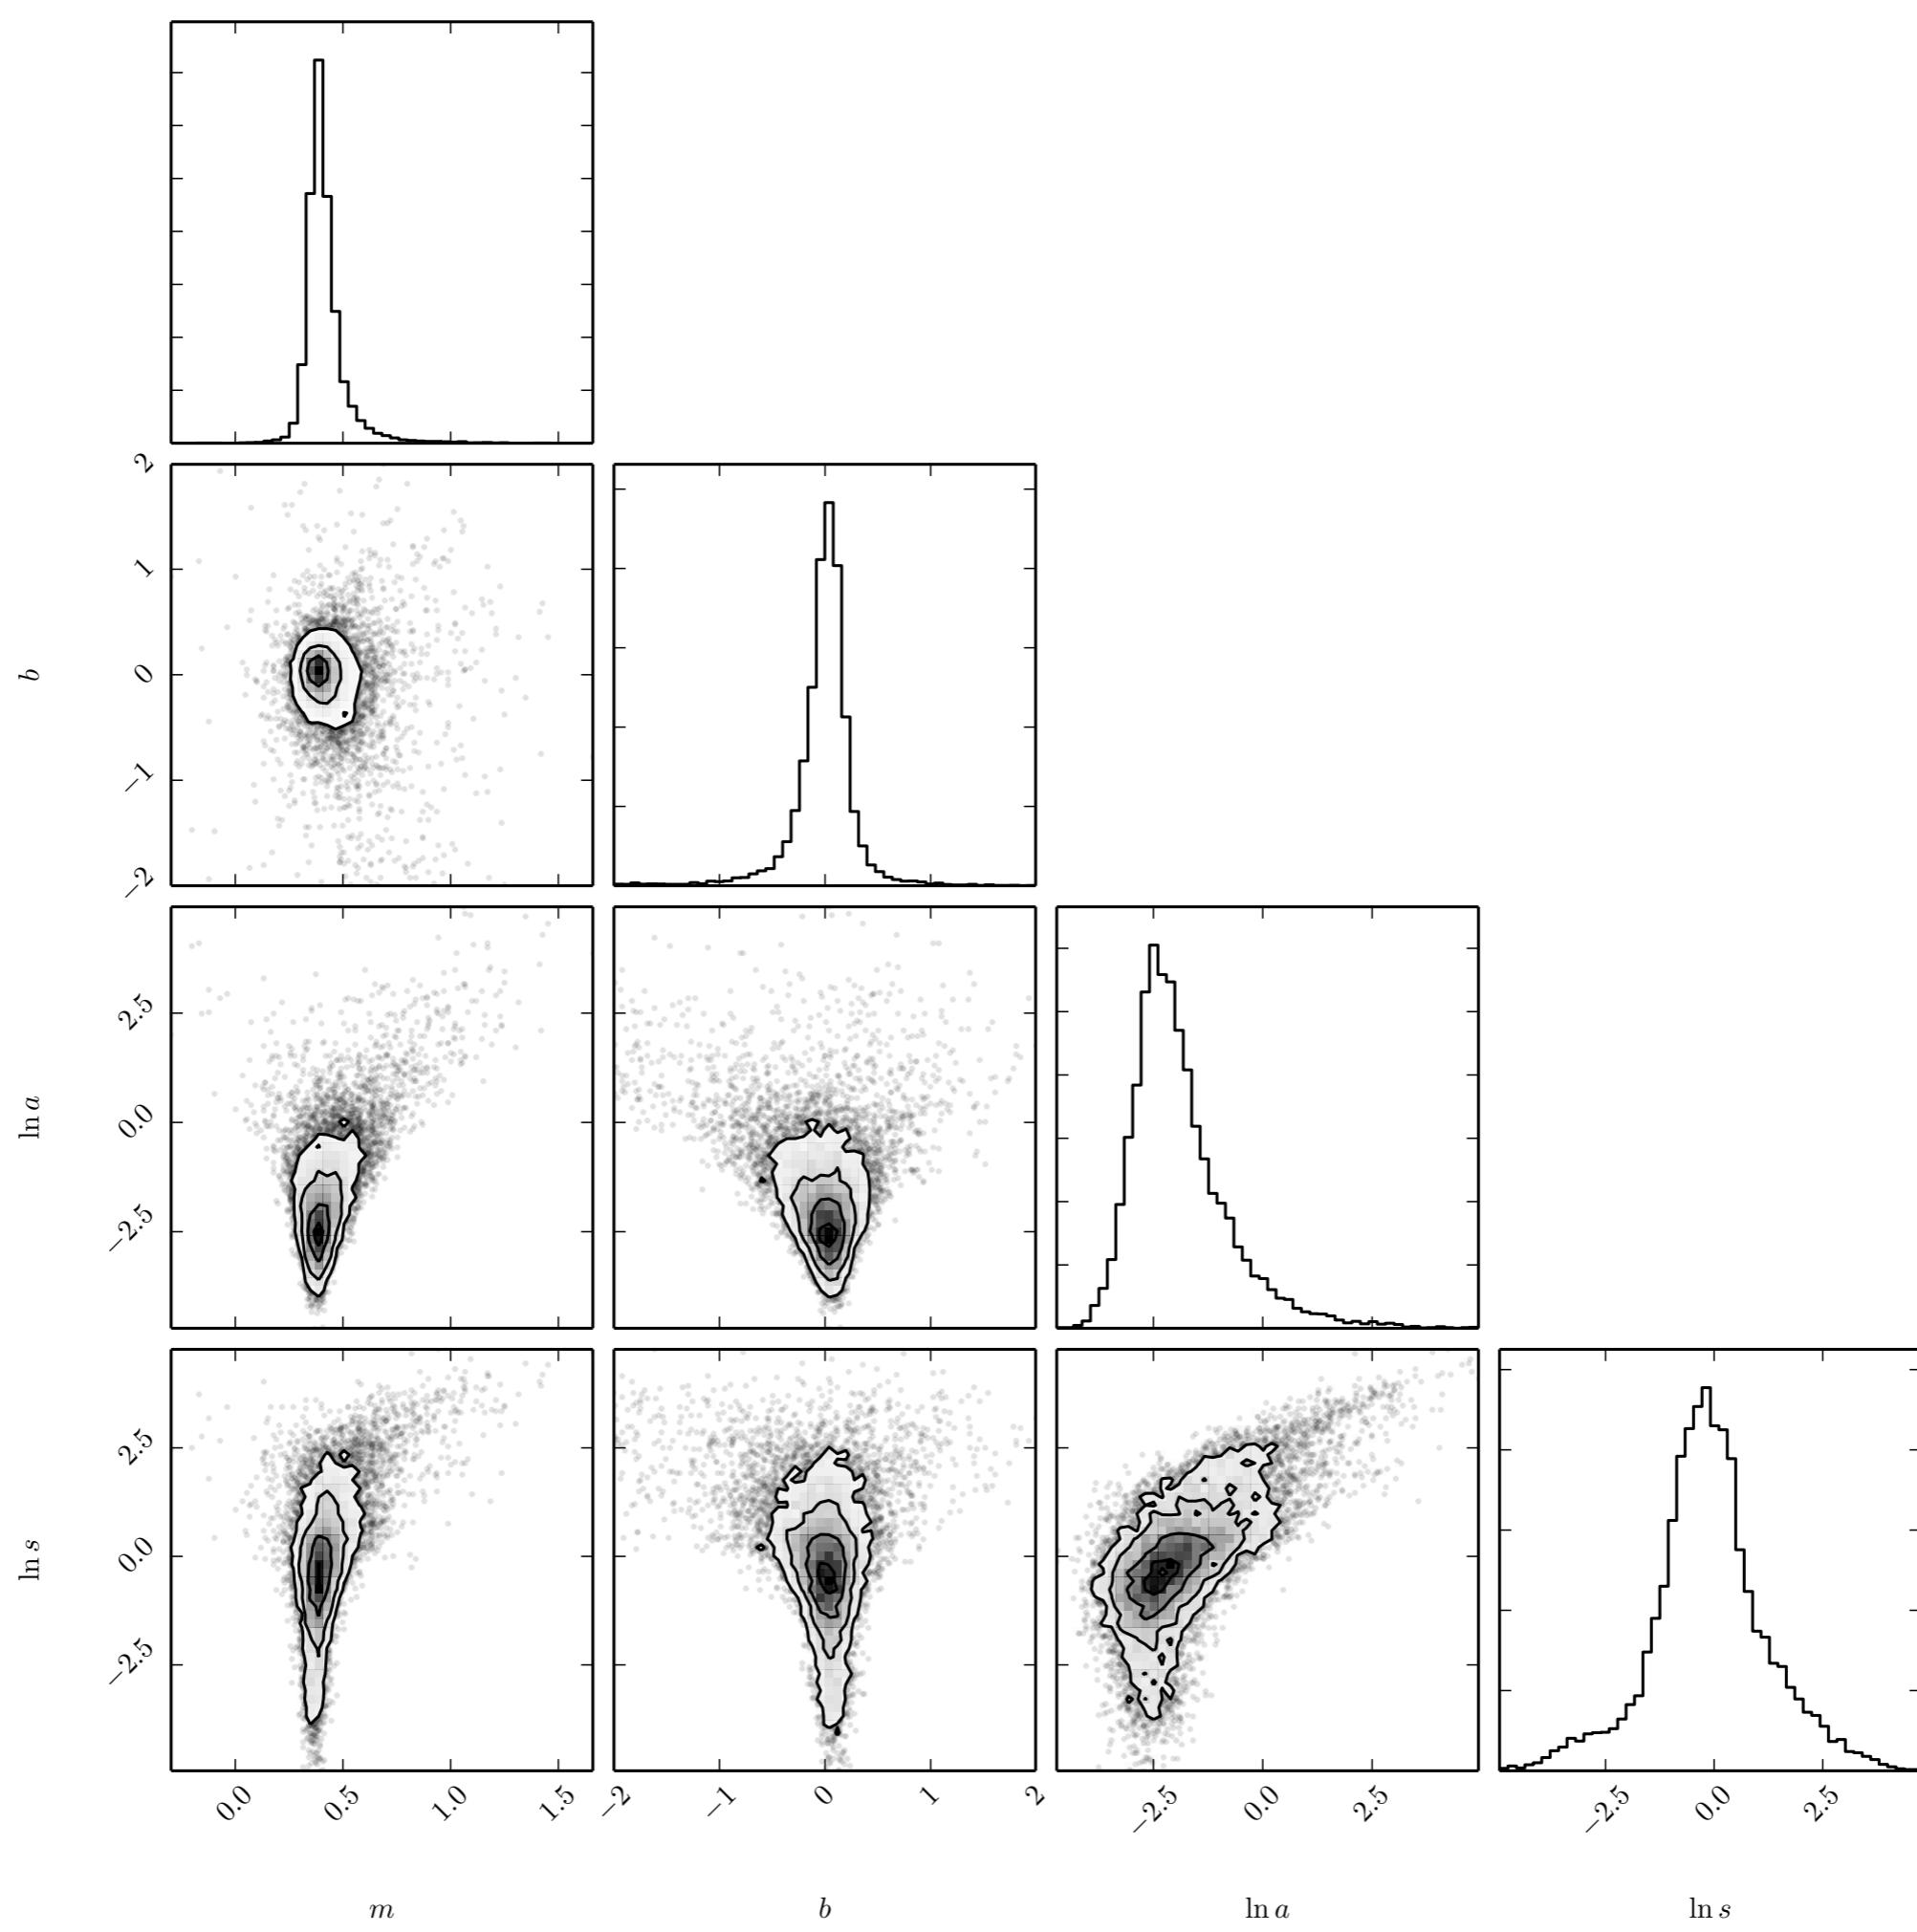
**drop-in replacement** for your current log-likelihood function!

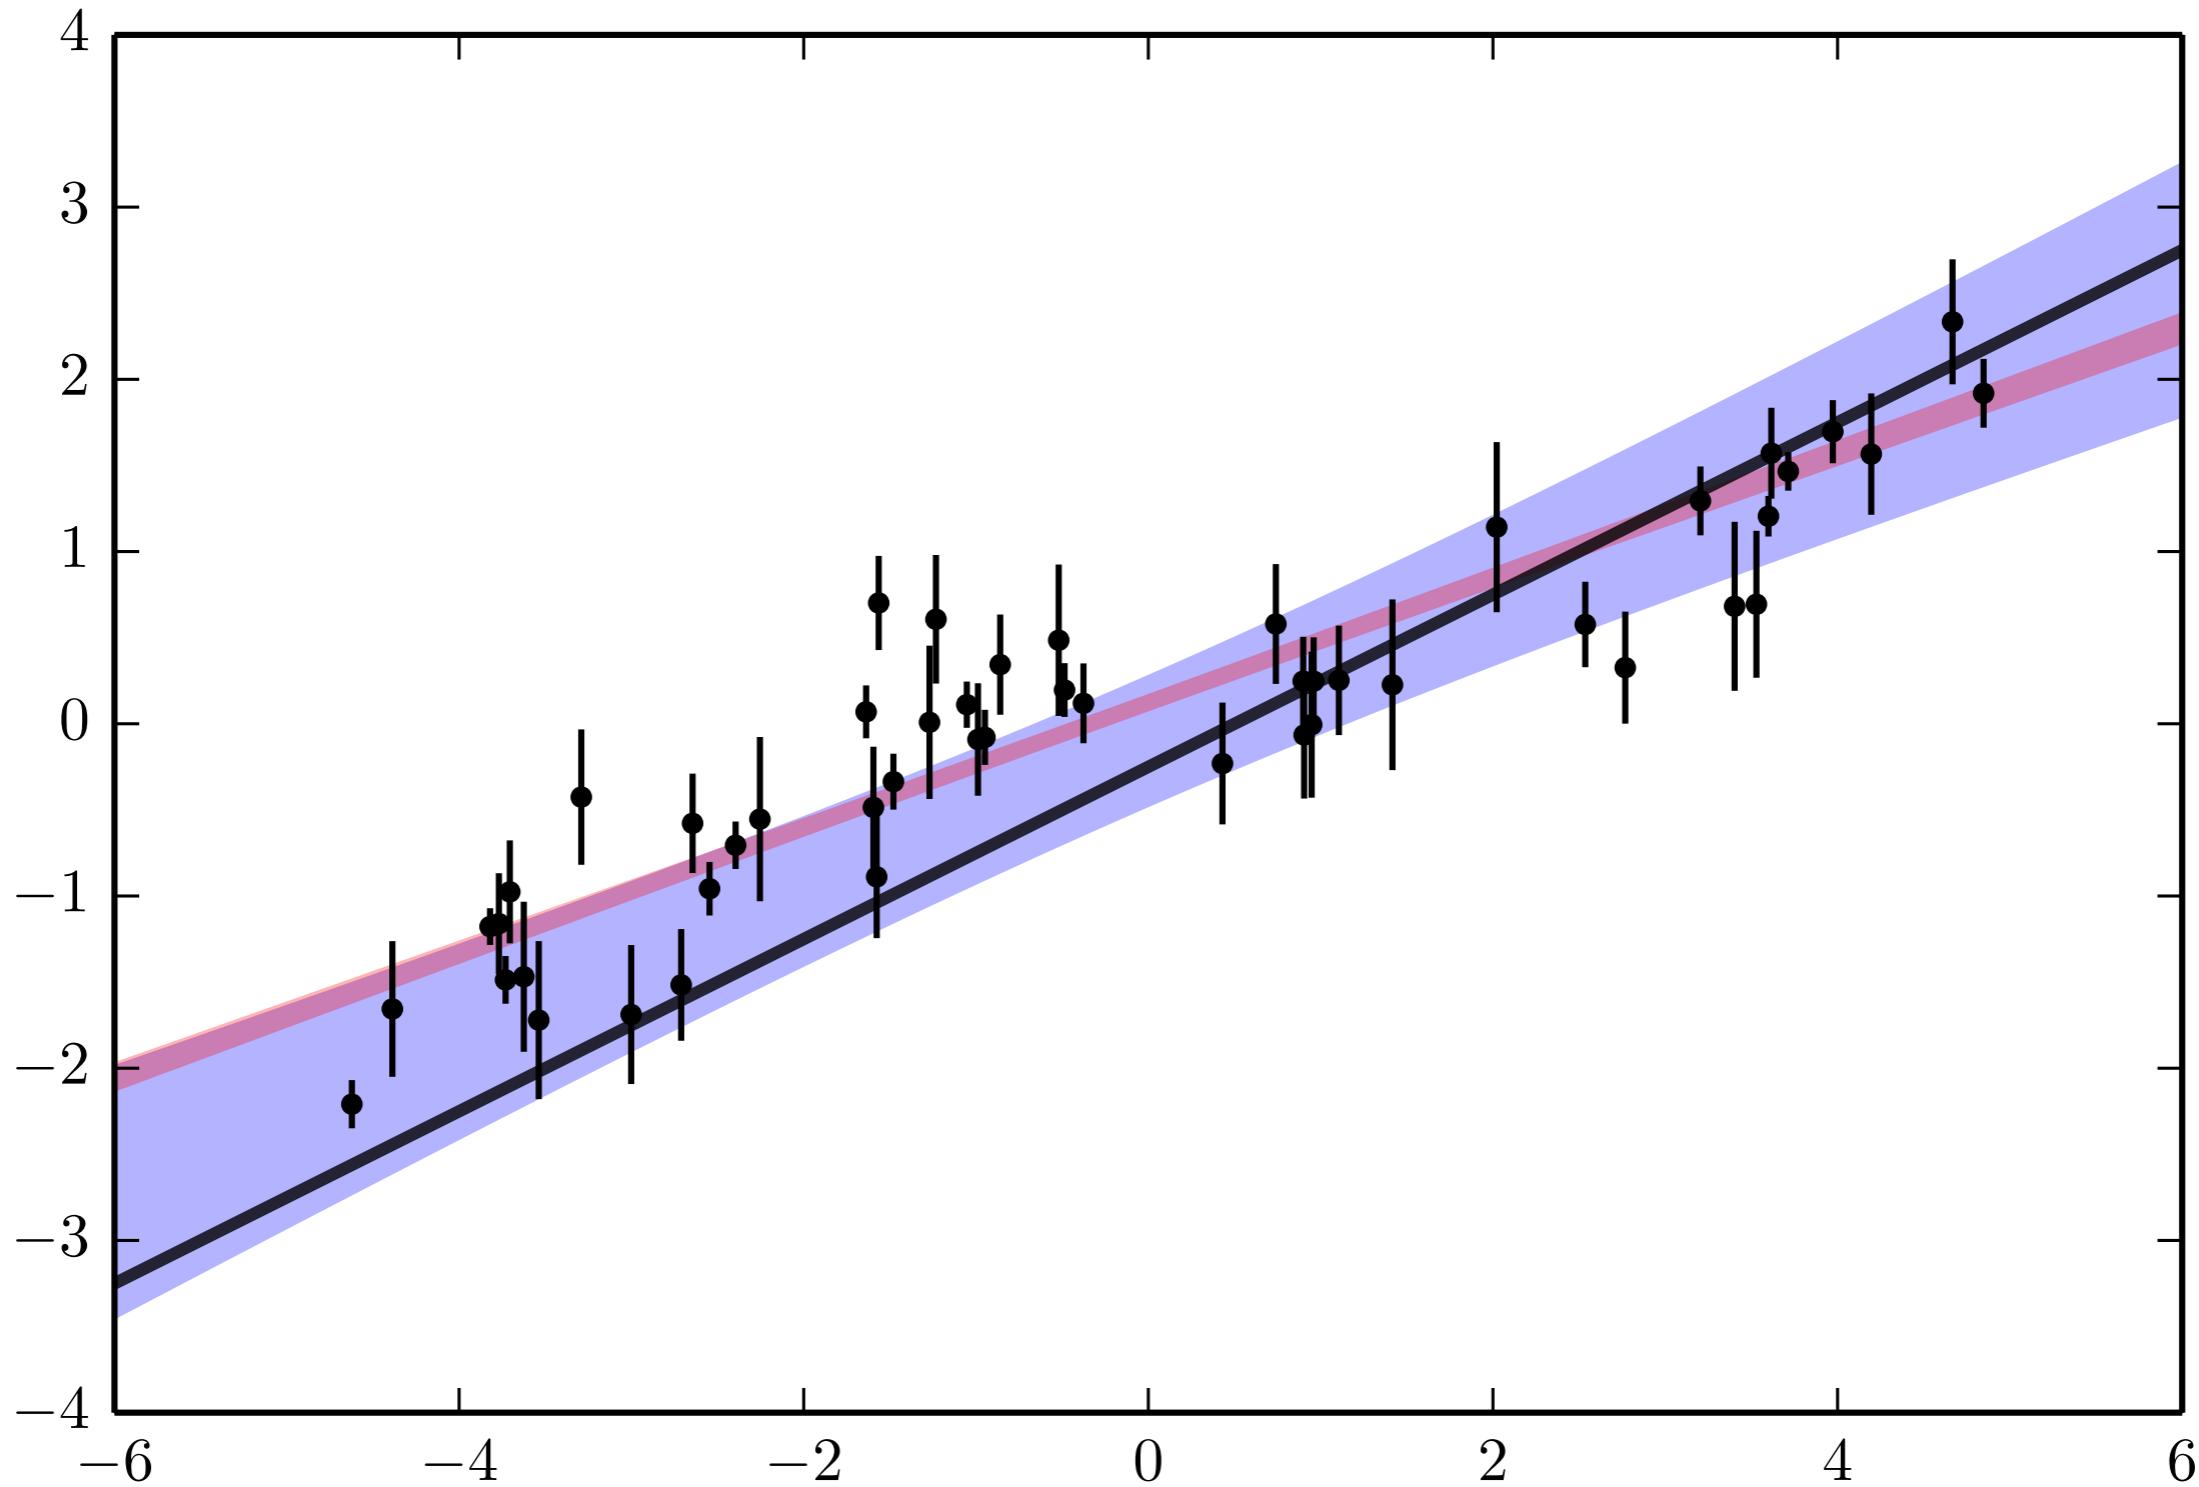
# emcee

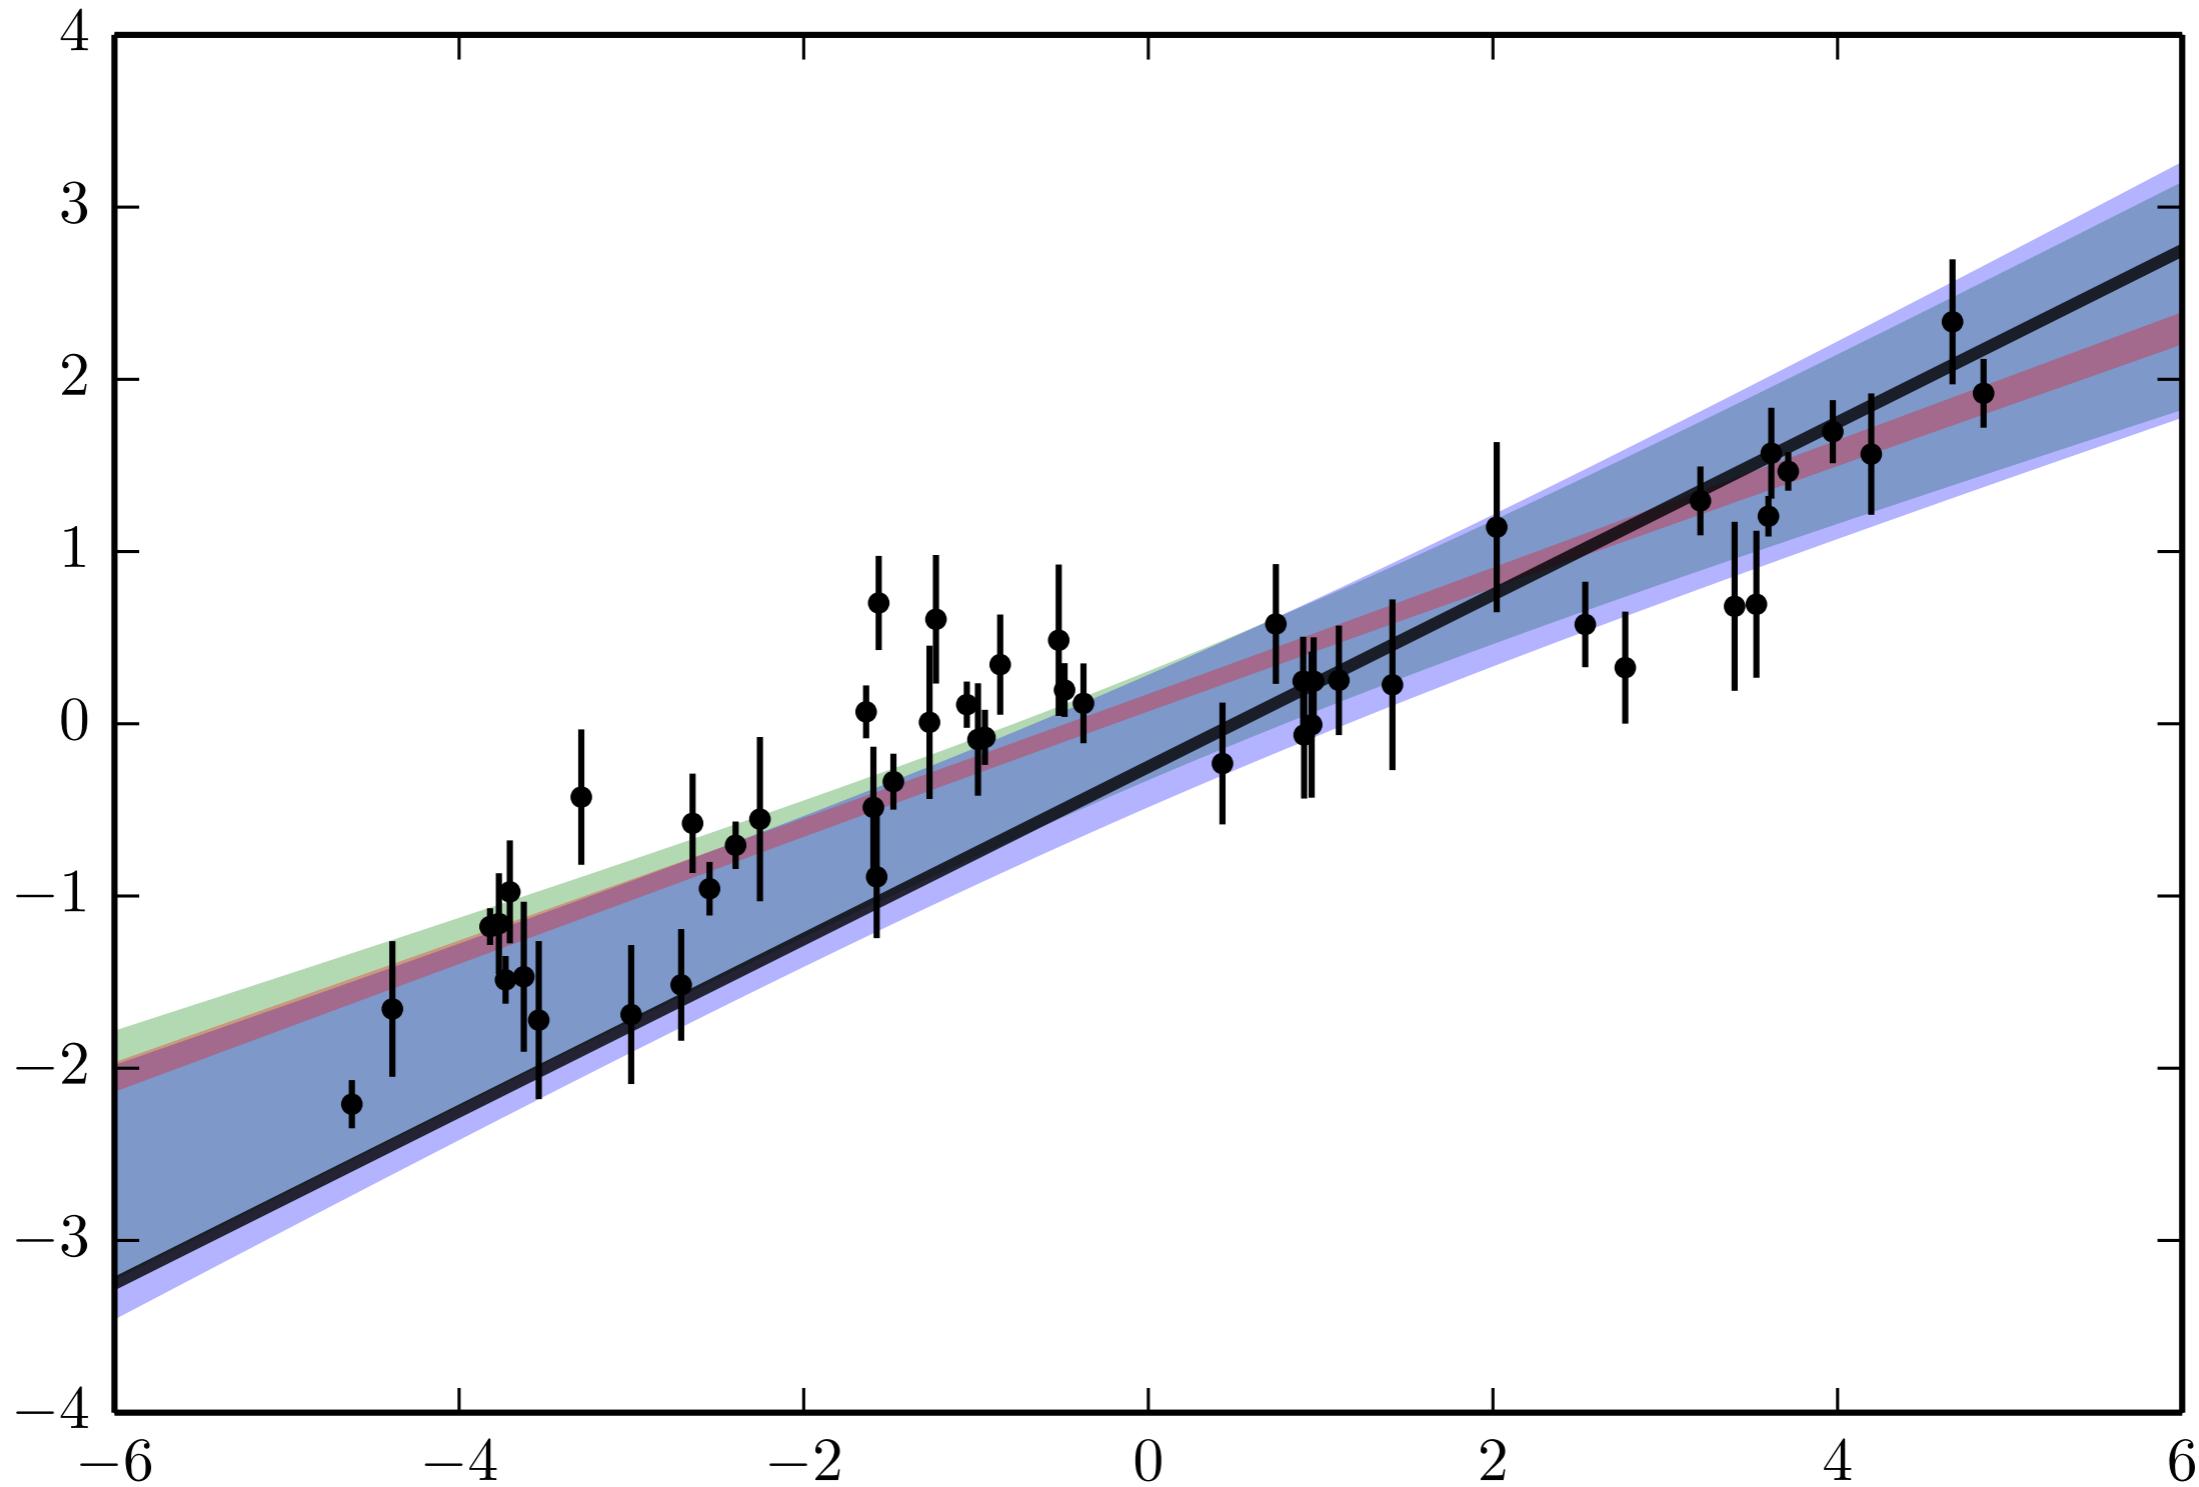
the  
MCMC  
Hammer

[arxiv.org/abs/1202.3665](https://arxiv.org/abs/1202.3665)  
[dan.iel.fm/emcee](https://dan.iel.fm/emcee)



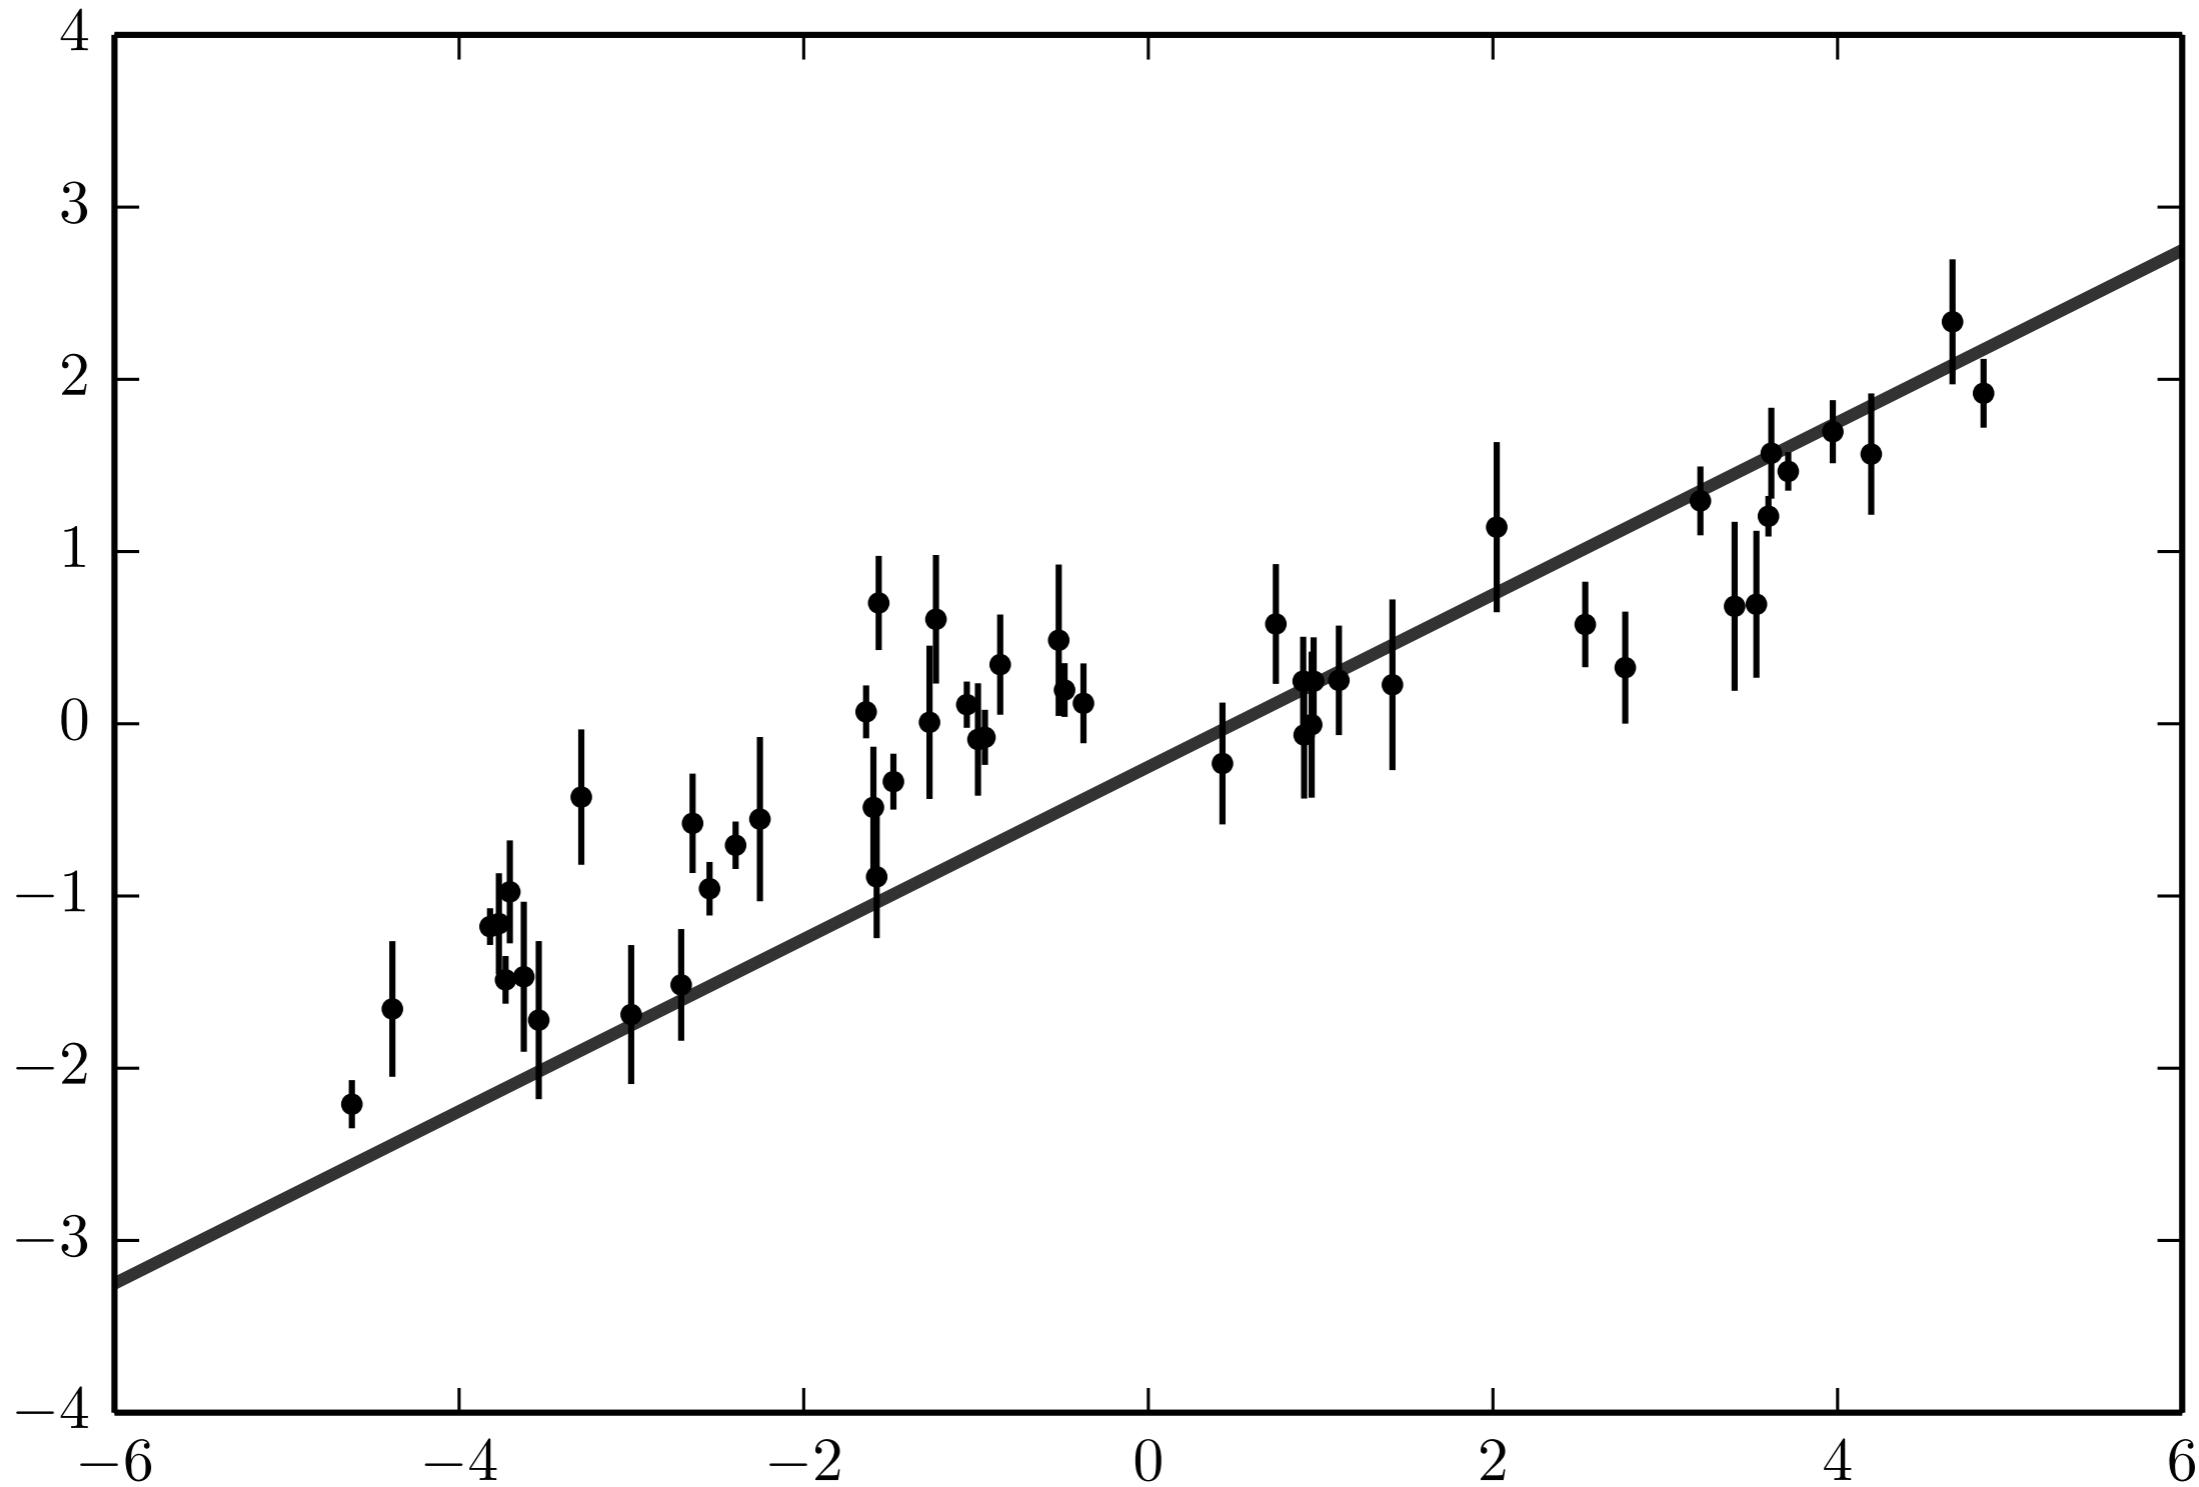


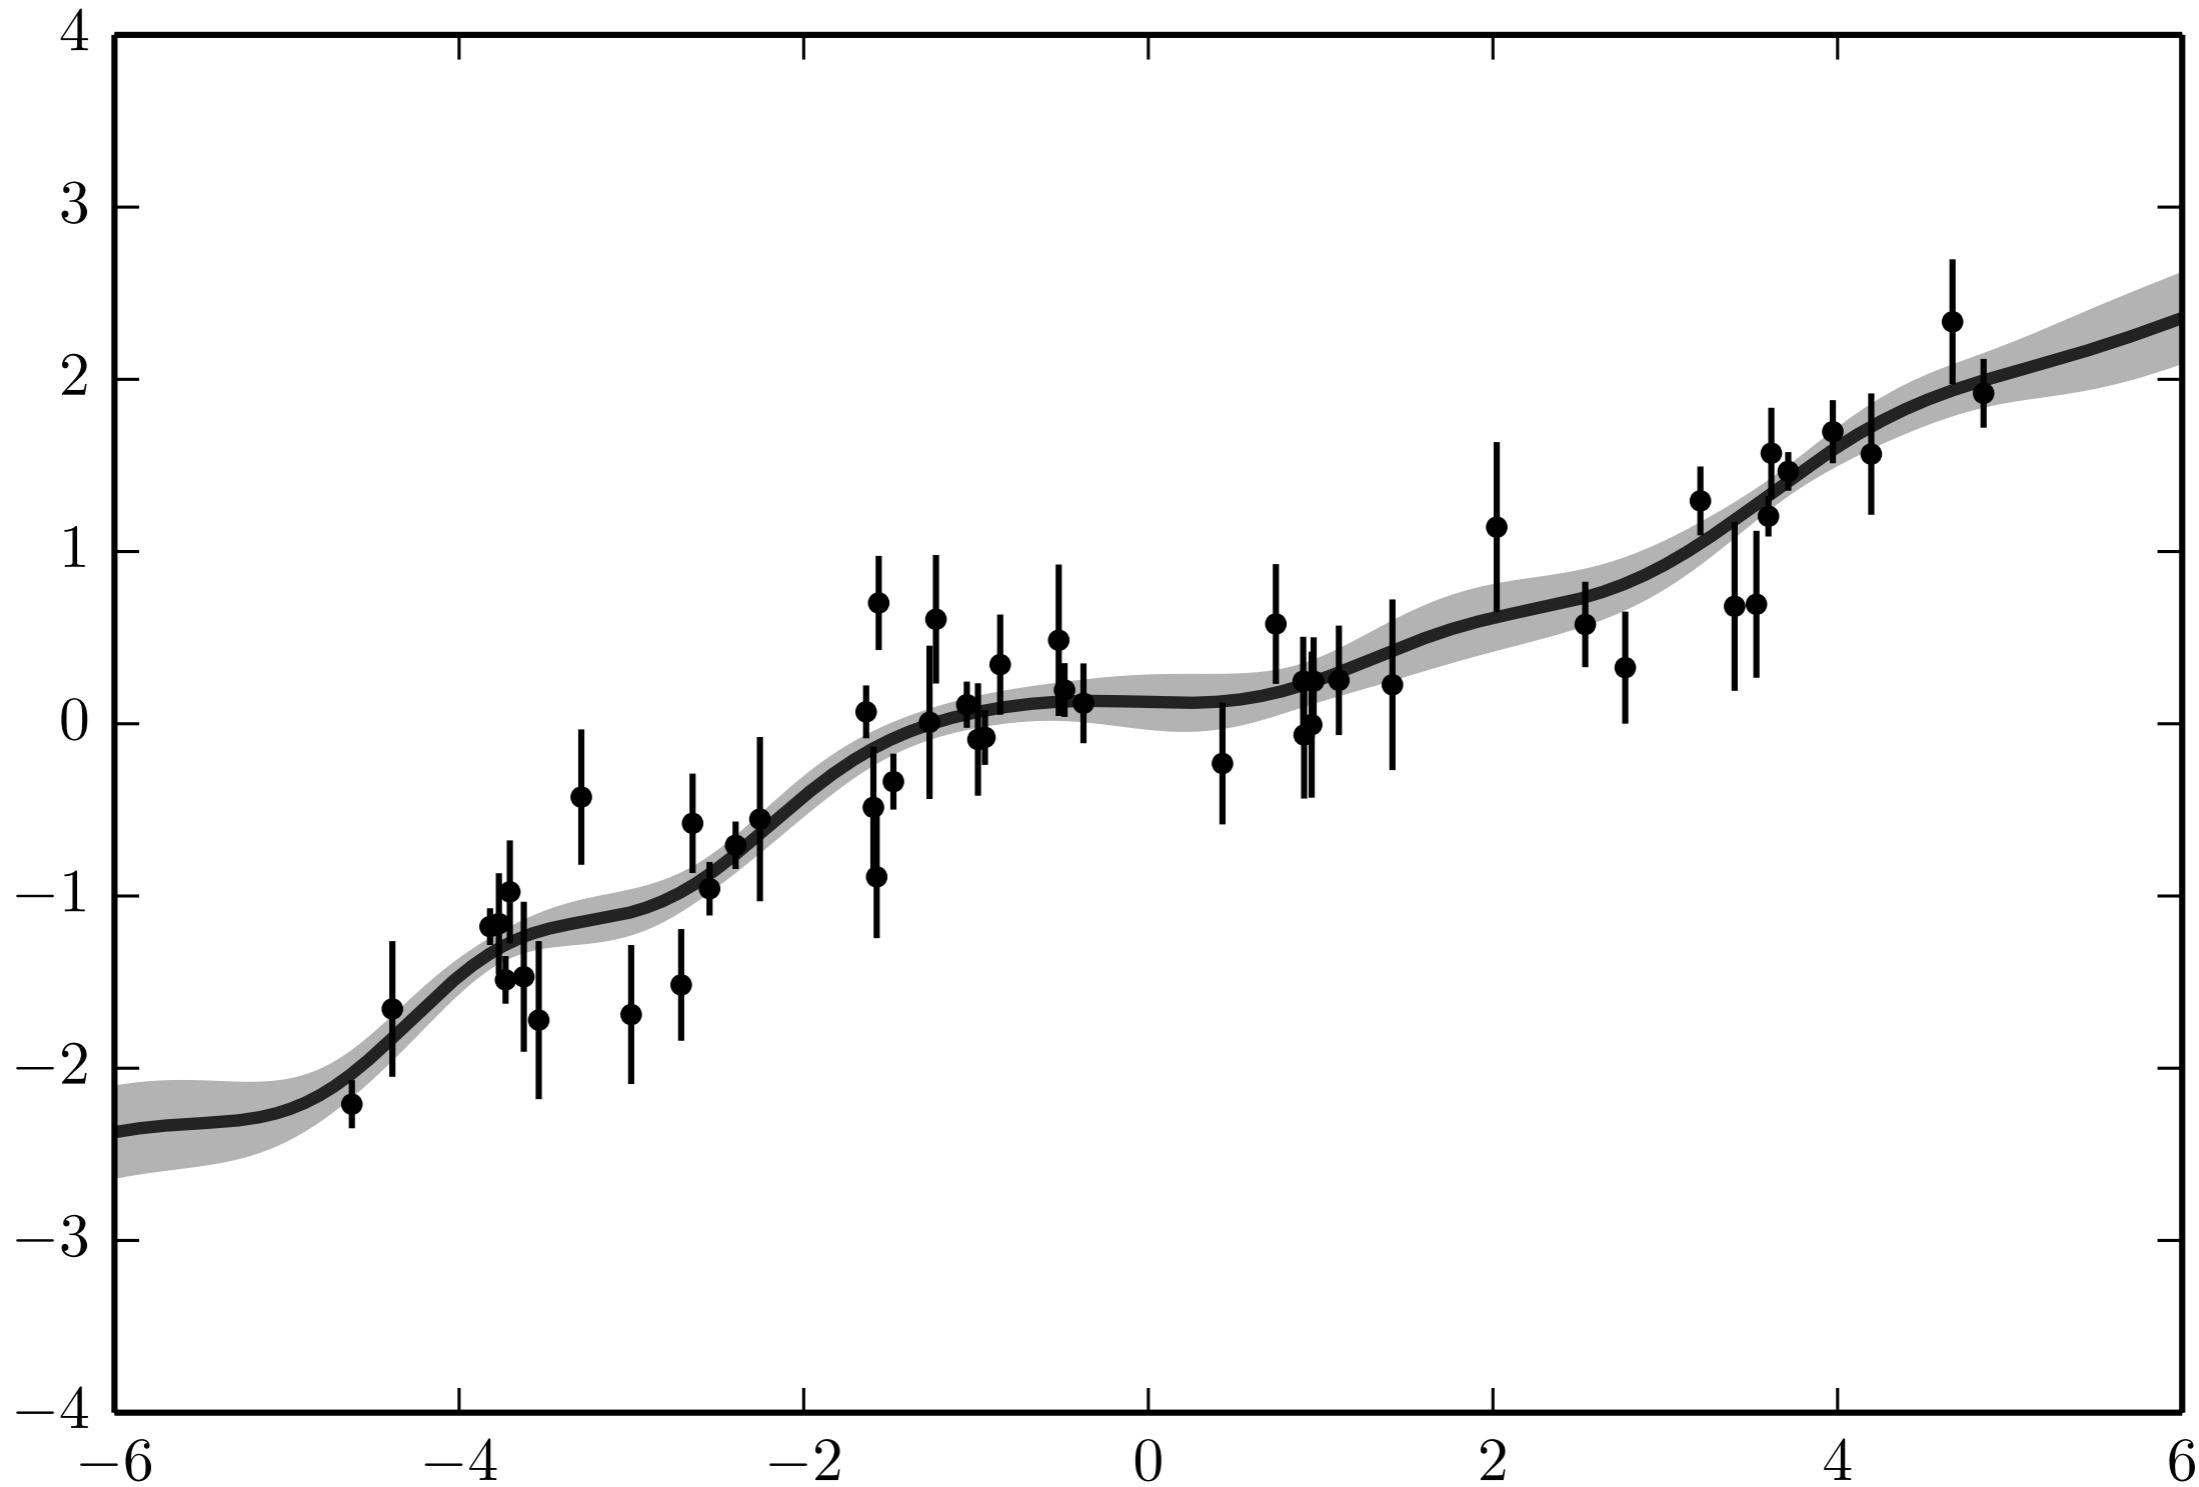




# Prediction?

---





A close-up photograph of a young kitten with a white and light brown coat, sitting in a field of green grass. The kitten has large, expressive blue eyes and is looking directly at the camera. The background is filled with blades of grass.

take a deep  
breath.



Flickr user kpjas

2

The formal Gaussian process.

---

## The model.

---

$$\begin{aligned}\log p(\mathbf{y} | \mathbf{x}, \boldsymbol{\sigma}, \boldsymbol{\theta}, \boldsymbol{\alpha}) = & -\frac{1}{2} [\mathbf{y} - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x})]^T K_{\boldsymbol{\alpha}}(\mathbf{x}, \boldsymbol{\sigma})^{-1} [\mathbf{y} - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x})] \\ & -\frac{1}{2} \log \det K_{\boldsymbol{\alpha}}(\mathbf{x}, \boldsymbol{\sigma}) - \frac{N}{2} \log 2 \pi\end{aligned}$$

where

$$[K_{\boldsymbol{\alpha}}(\mathbf{x}, \boldsymbol{\sigma})]_{ij} = \sigma_i^2 \delta_{ij} + k_{\boldsymbol{\alpha}}(x_i, x_j)$$

drop-in replacement for your current log-likelihood function!

---

## The model.

---

$$y \sim \mathcal{N}(f_{\theta}(x), K_{\alpha}(x, \sigma))$$

where

$$[K_{\alpha}(x, \sigma)]_{ij} = \sigma_i^2 \delta_{ij} + k_{\alpha}(x_i, x_j)$$

drop-in replacement for your current log-likelihood function!

---

the data are drawn from one

**huge  
Gaussian**

\*

\* the dimension is the number of data points.

## A generative model

---

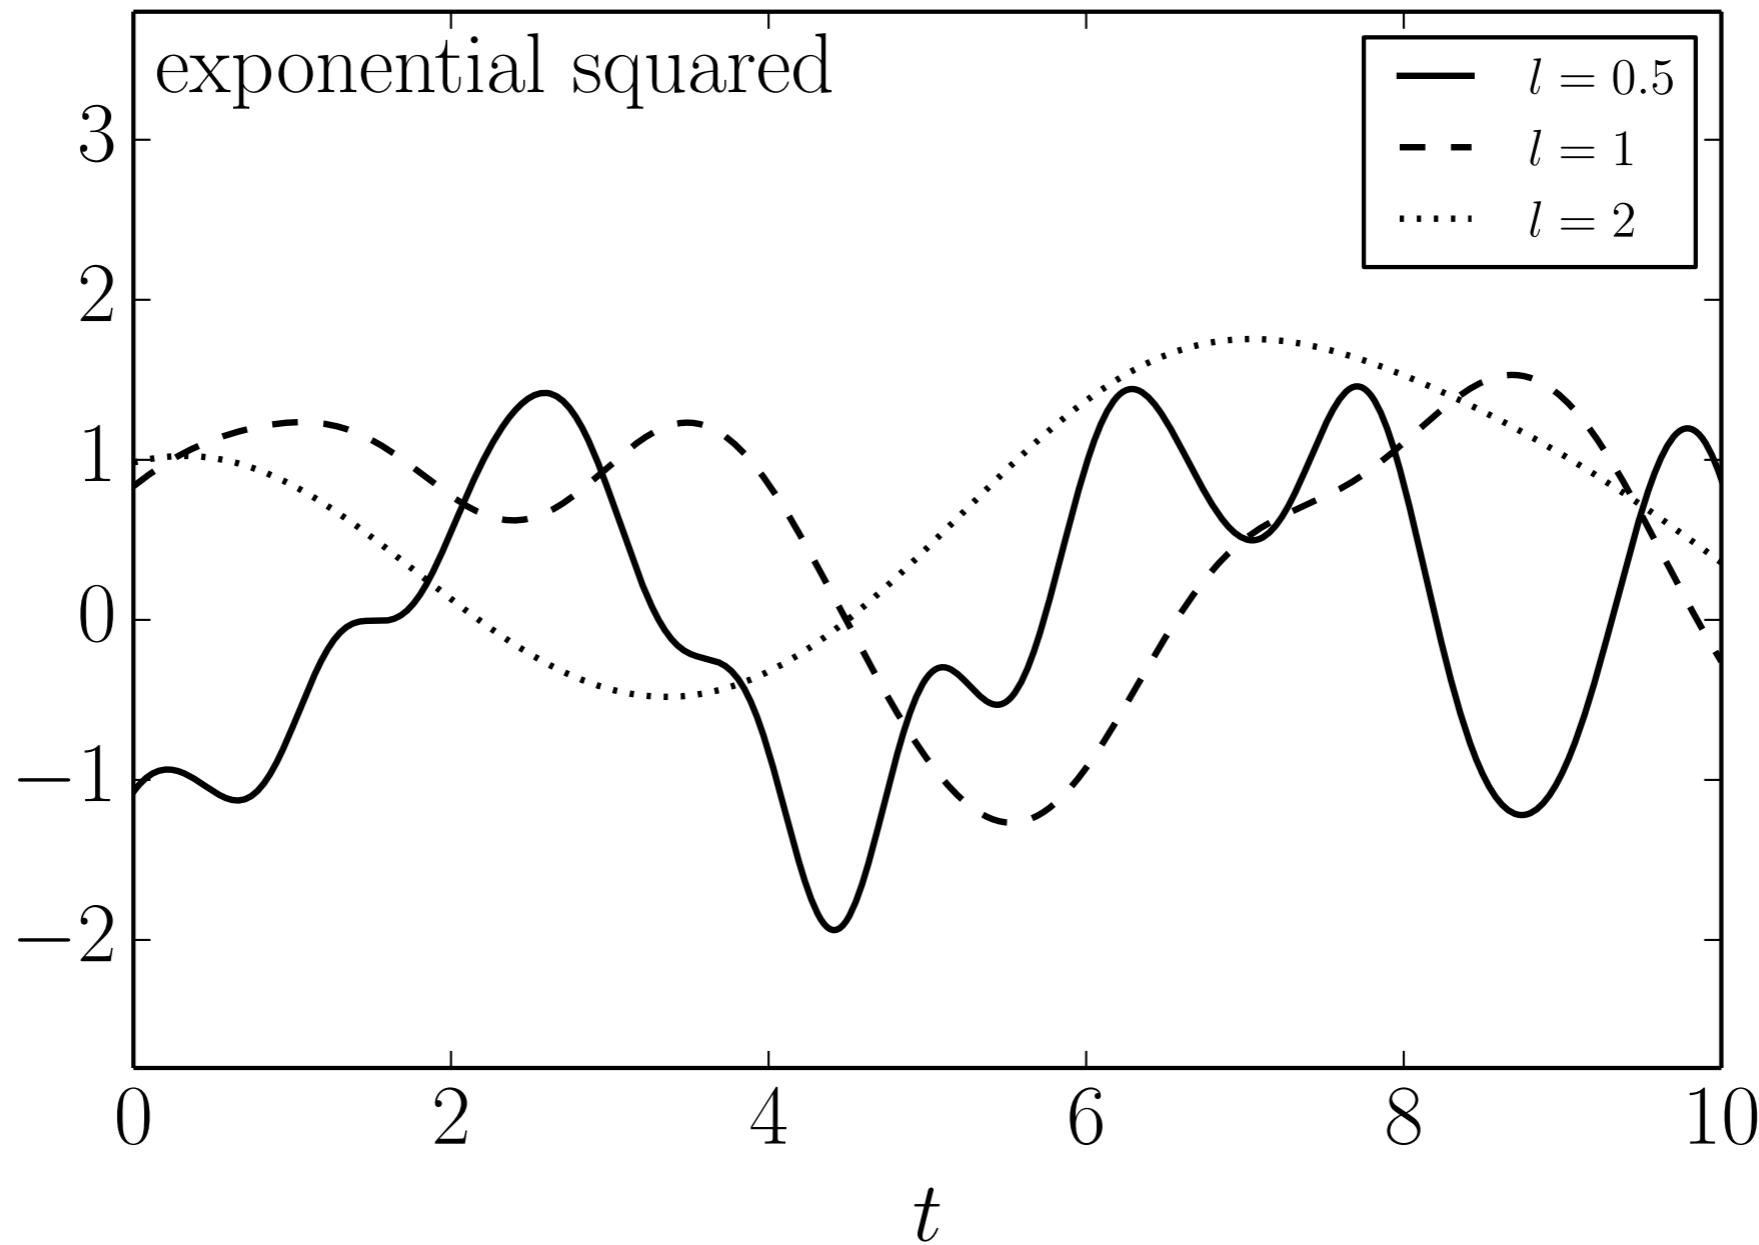
$$y \sim \mathcal{N}(f_{\theta}(x), K_{\alpha}(x, \sigma))$$

a probability distribution for  $y$  values

## “Likelihood” samples.

---

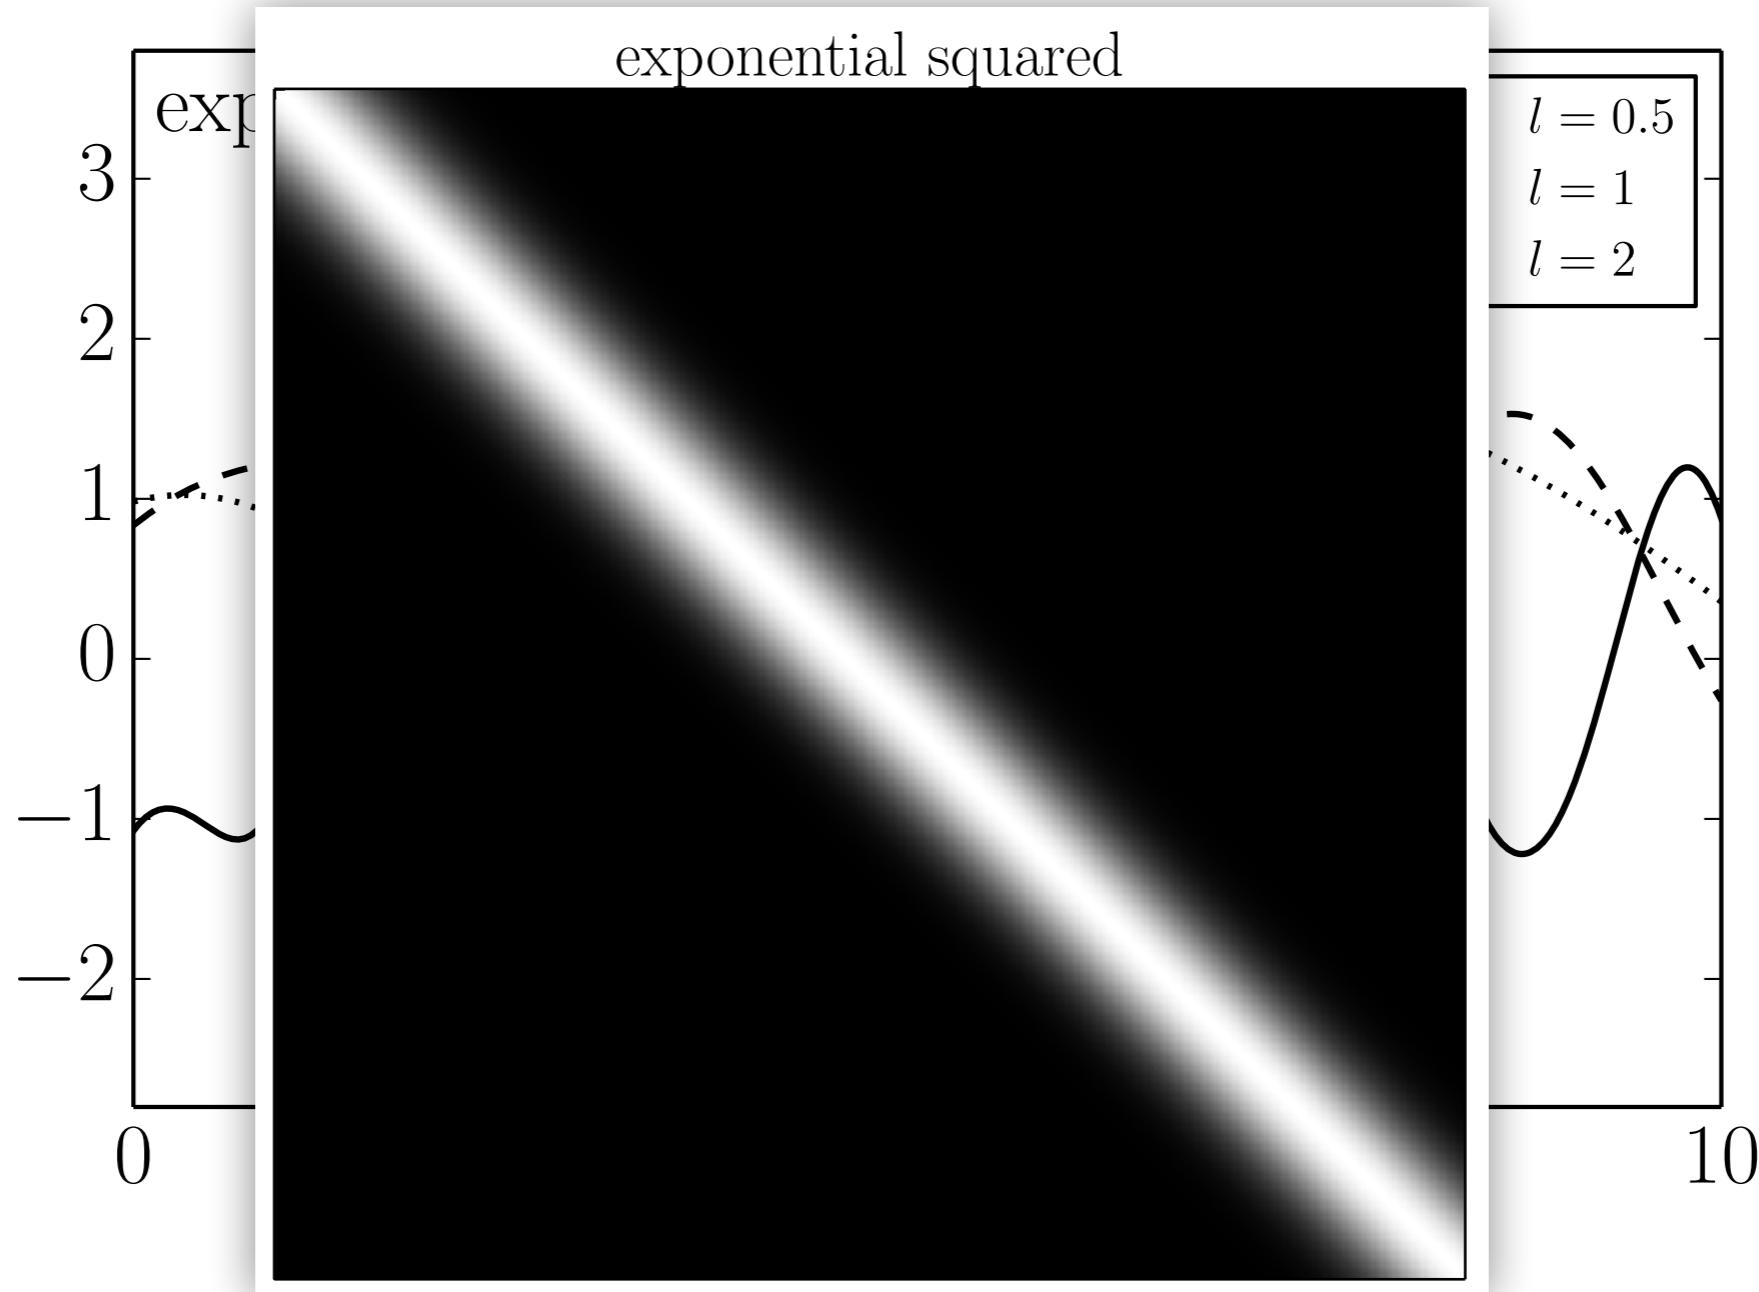
$$k_{\alpha}(x_i, x_j) = \exp\left(-\frac{[x_i - x_j]^2}{2 \ell^2}\right)$$



## “Likelihood” samples.

---

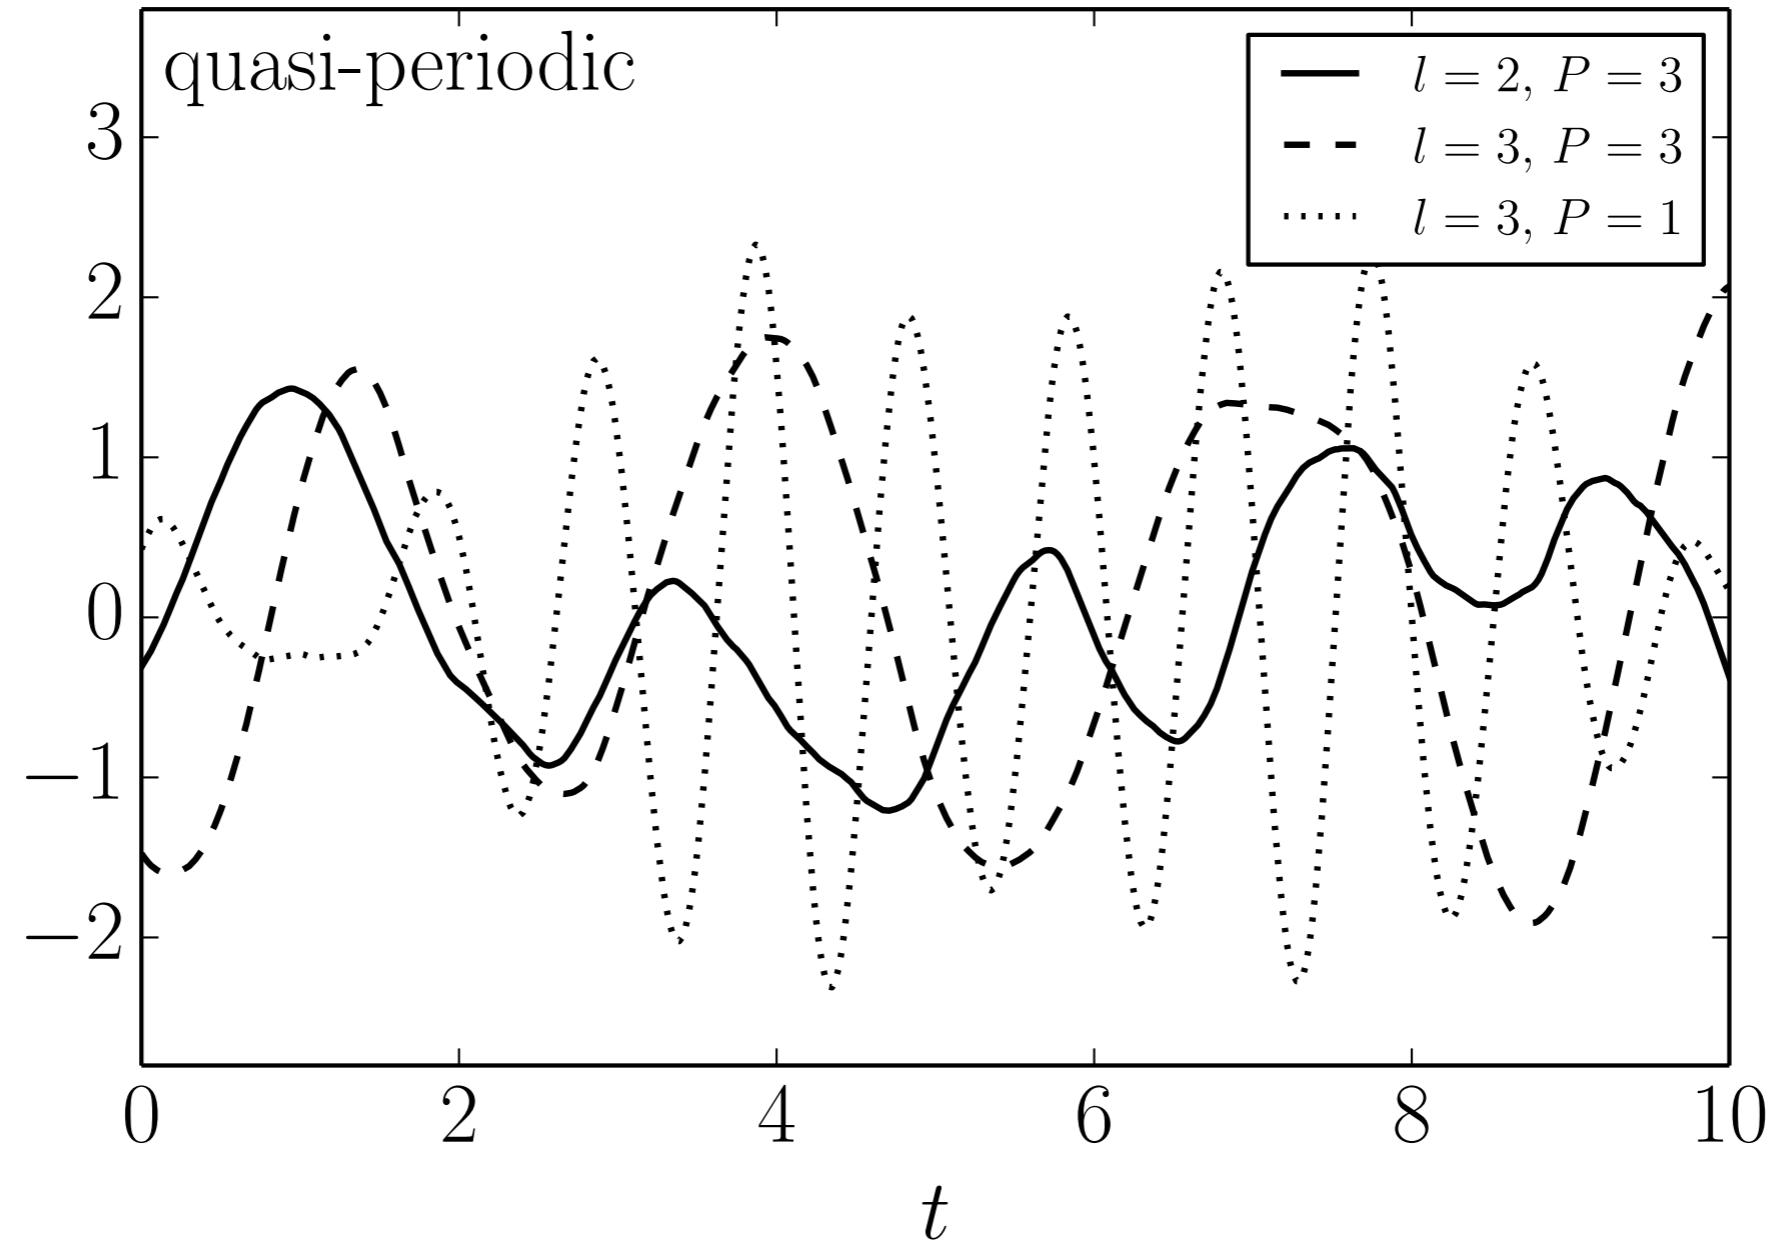
$$k_{\alpha}(x_i, x_j) = \exp\left(-\frac{[x_i - x_j]^2}{2 \ell^2}\right)$$



## “Likelihood” samples.

---

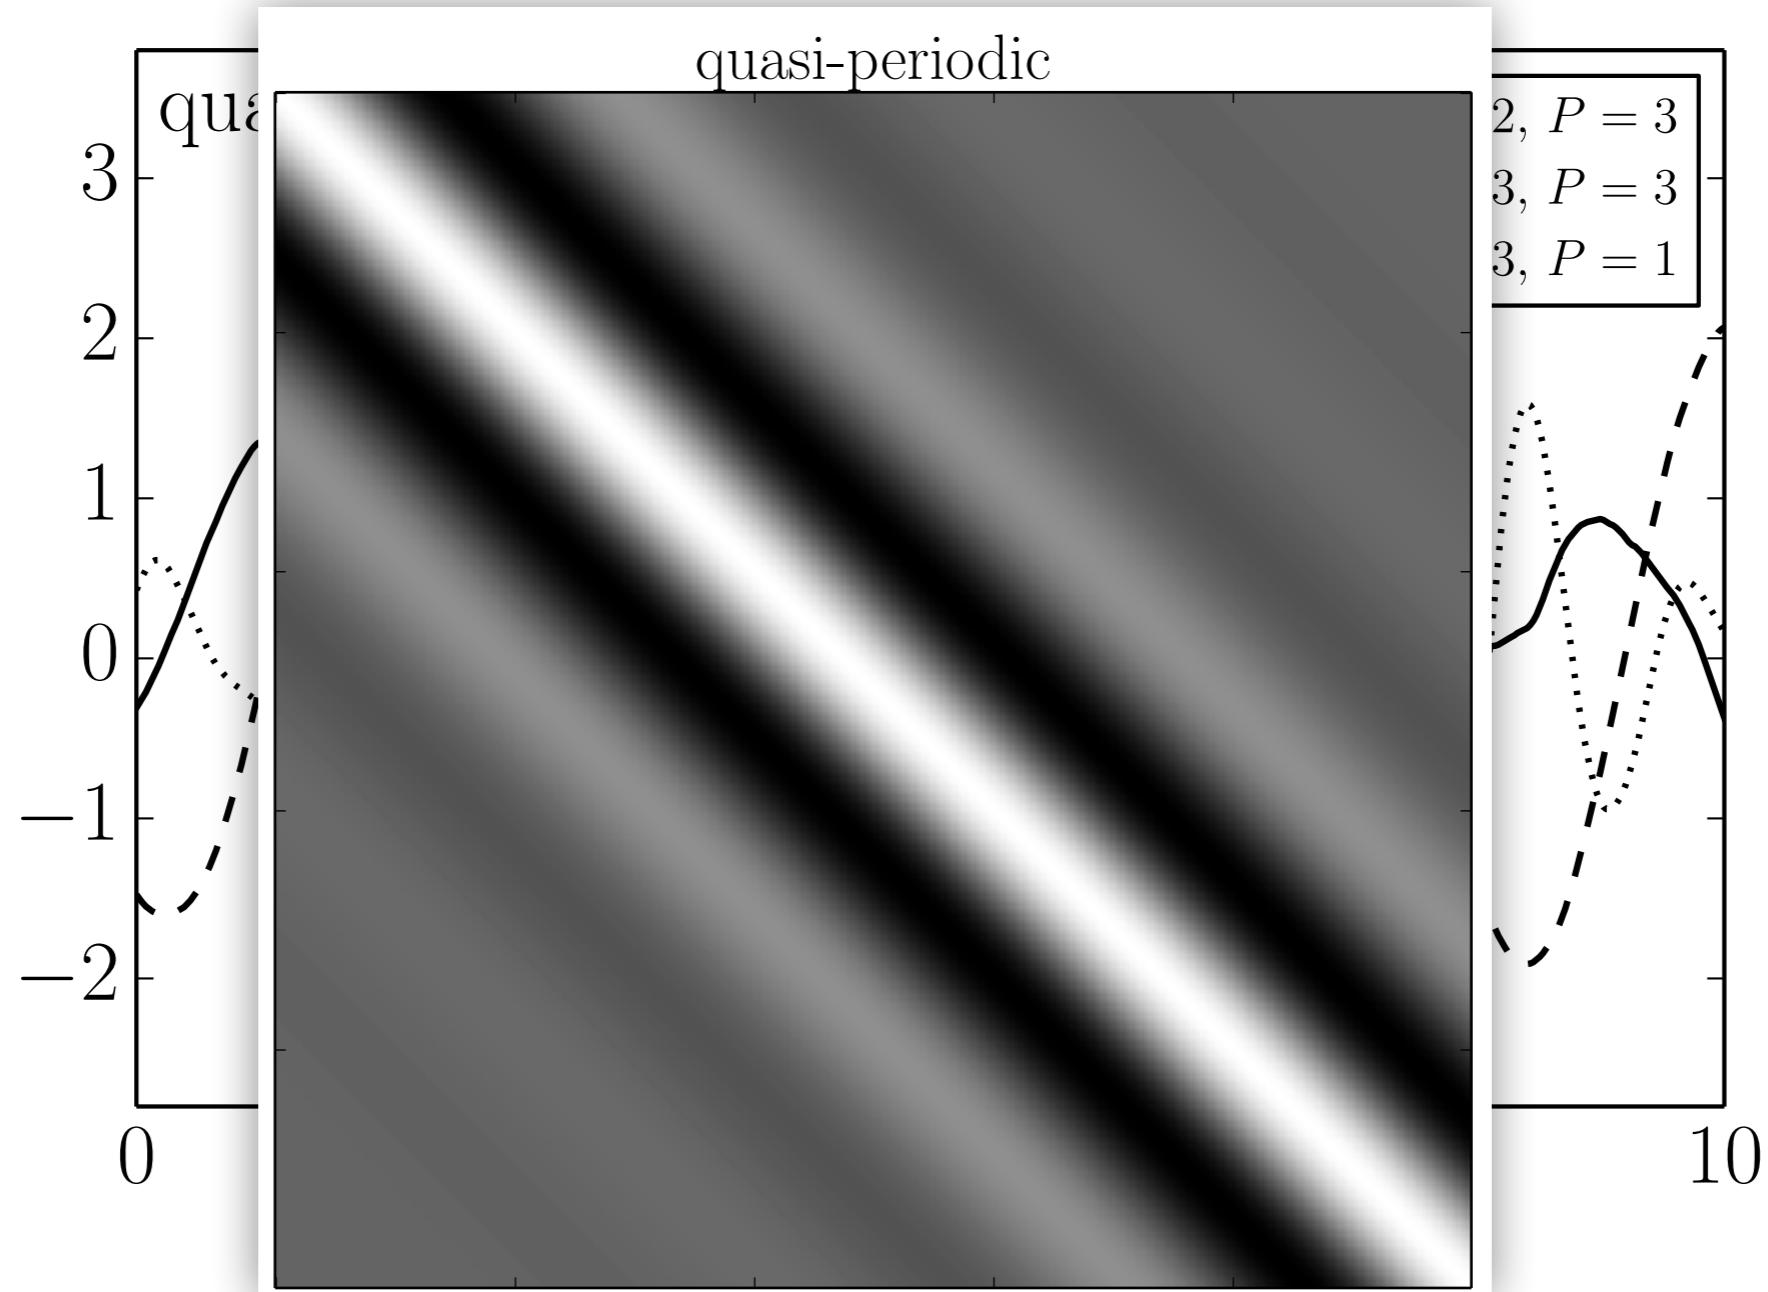
$$k_{\alpha}(x_i, x_j) = \left[ 1 + \frac{\sqrt{3} |x_i - x_j|}{\ell} \right] \exp \left( -\frac{|x_i - x_j|}{\ell} \right) \cos \left( \frac{2 \pi |x_i - x_j|}{P} \right)$$



## “Likelihood” samples.

---

$$k_{\alpha}(x_i, x_j) = \left[ 1 + \frac{\sqrt{3} |x_i - x_j|}{\ell} \right] \exp \left( -\frac{|x_i - x_j|}{\ell} \right) \cos \left( \frac{2 \pi |x_i - x_j|}{P} \right)$$



## The conditional distribution

---

$$\mathbf{y} \sim \mathcal{N}(f_{\boldsymbol{\theta}}(\mathbf{x}), K_{\boldsymbol{\alpha}}(\mathbf{x}, \boldsymbol{\sigma}))$$



$$\begin{bmatrix} \mathbf{y} \\ \mathbf{y}_* \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} f_{\boldsymbol{\theta}}(\mathbf{x}) \\ f_{\boldsymbol{\theta}}(\mathbf{x}_*) \end{bmatrix}, \begin{bmatrix} K_{\boldsymbol{\alpha}, \mathbf{x}, \mathbf{x}} & K_{\boldsymbol{\alpha}, \mathbf{x}, *} \\ K_{\boldsymbol{\alpha}, *, \mathbf{x}} & K_{\boldsymbol{\alpha}, *, *} \end{bmatrix}\right)$$



$$\begin{aligned} \mathbf{y}_* | \mathbf{y} &\sim \mathcal{N}\left(K_{\boldsymbol{\alpha}, *, \mathbf{x}} K_{\boldsymbol{\alpha}, \mathbf{x}, \mathbf{x}}^{-1} [\mathbf{y} - f_{\boldsymbol{\theta}}(\mathbf{x})] + f_{\boldsymbol{\theta}}(\mathbf{x}_*), \right. \\ &\quad \left. K_{\boldsymbol{\alpha}, *, *} - K_{\boldsymbol{\alpha}, *, \mathbf{x}} K_{\boldsymbol{\alpha}, \mathbf{x}, \mathbf{x}}^{-1} K_{\boldsymbol{\alpha}, \mathbf{x}, *} \right) \end{aligned}$$

just see Rasmussen & Williams (Chapter 2)

What's the **catch**?

---

*Kepler*

=

*Big Data*

(by some definition)

*Note:* I hate myself for this slide too...

## Computational complexity.

---

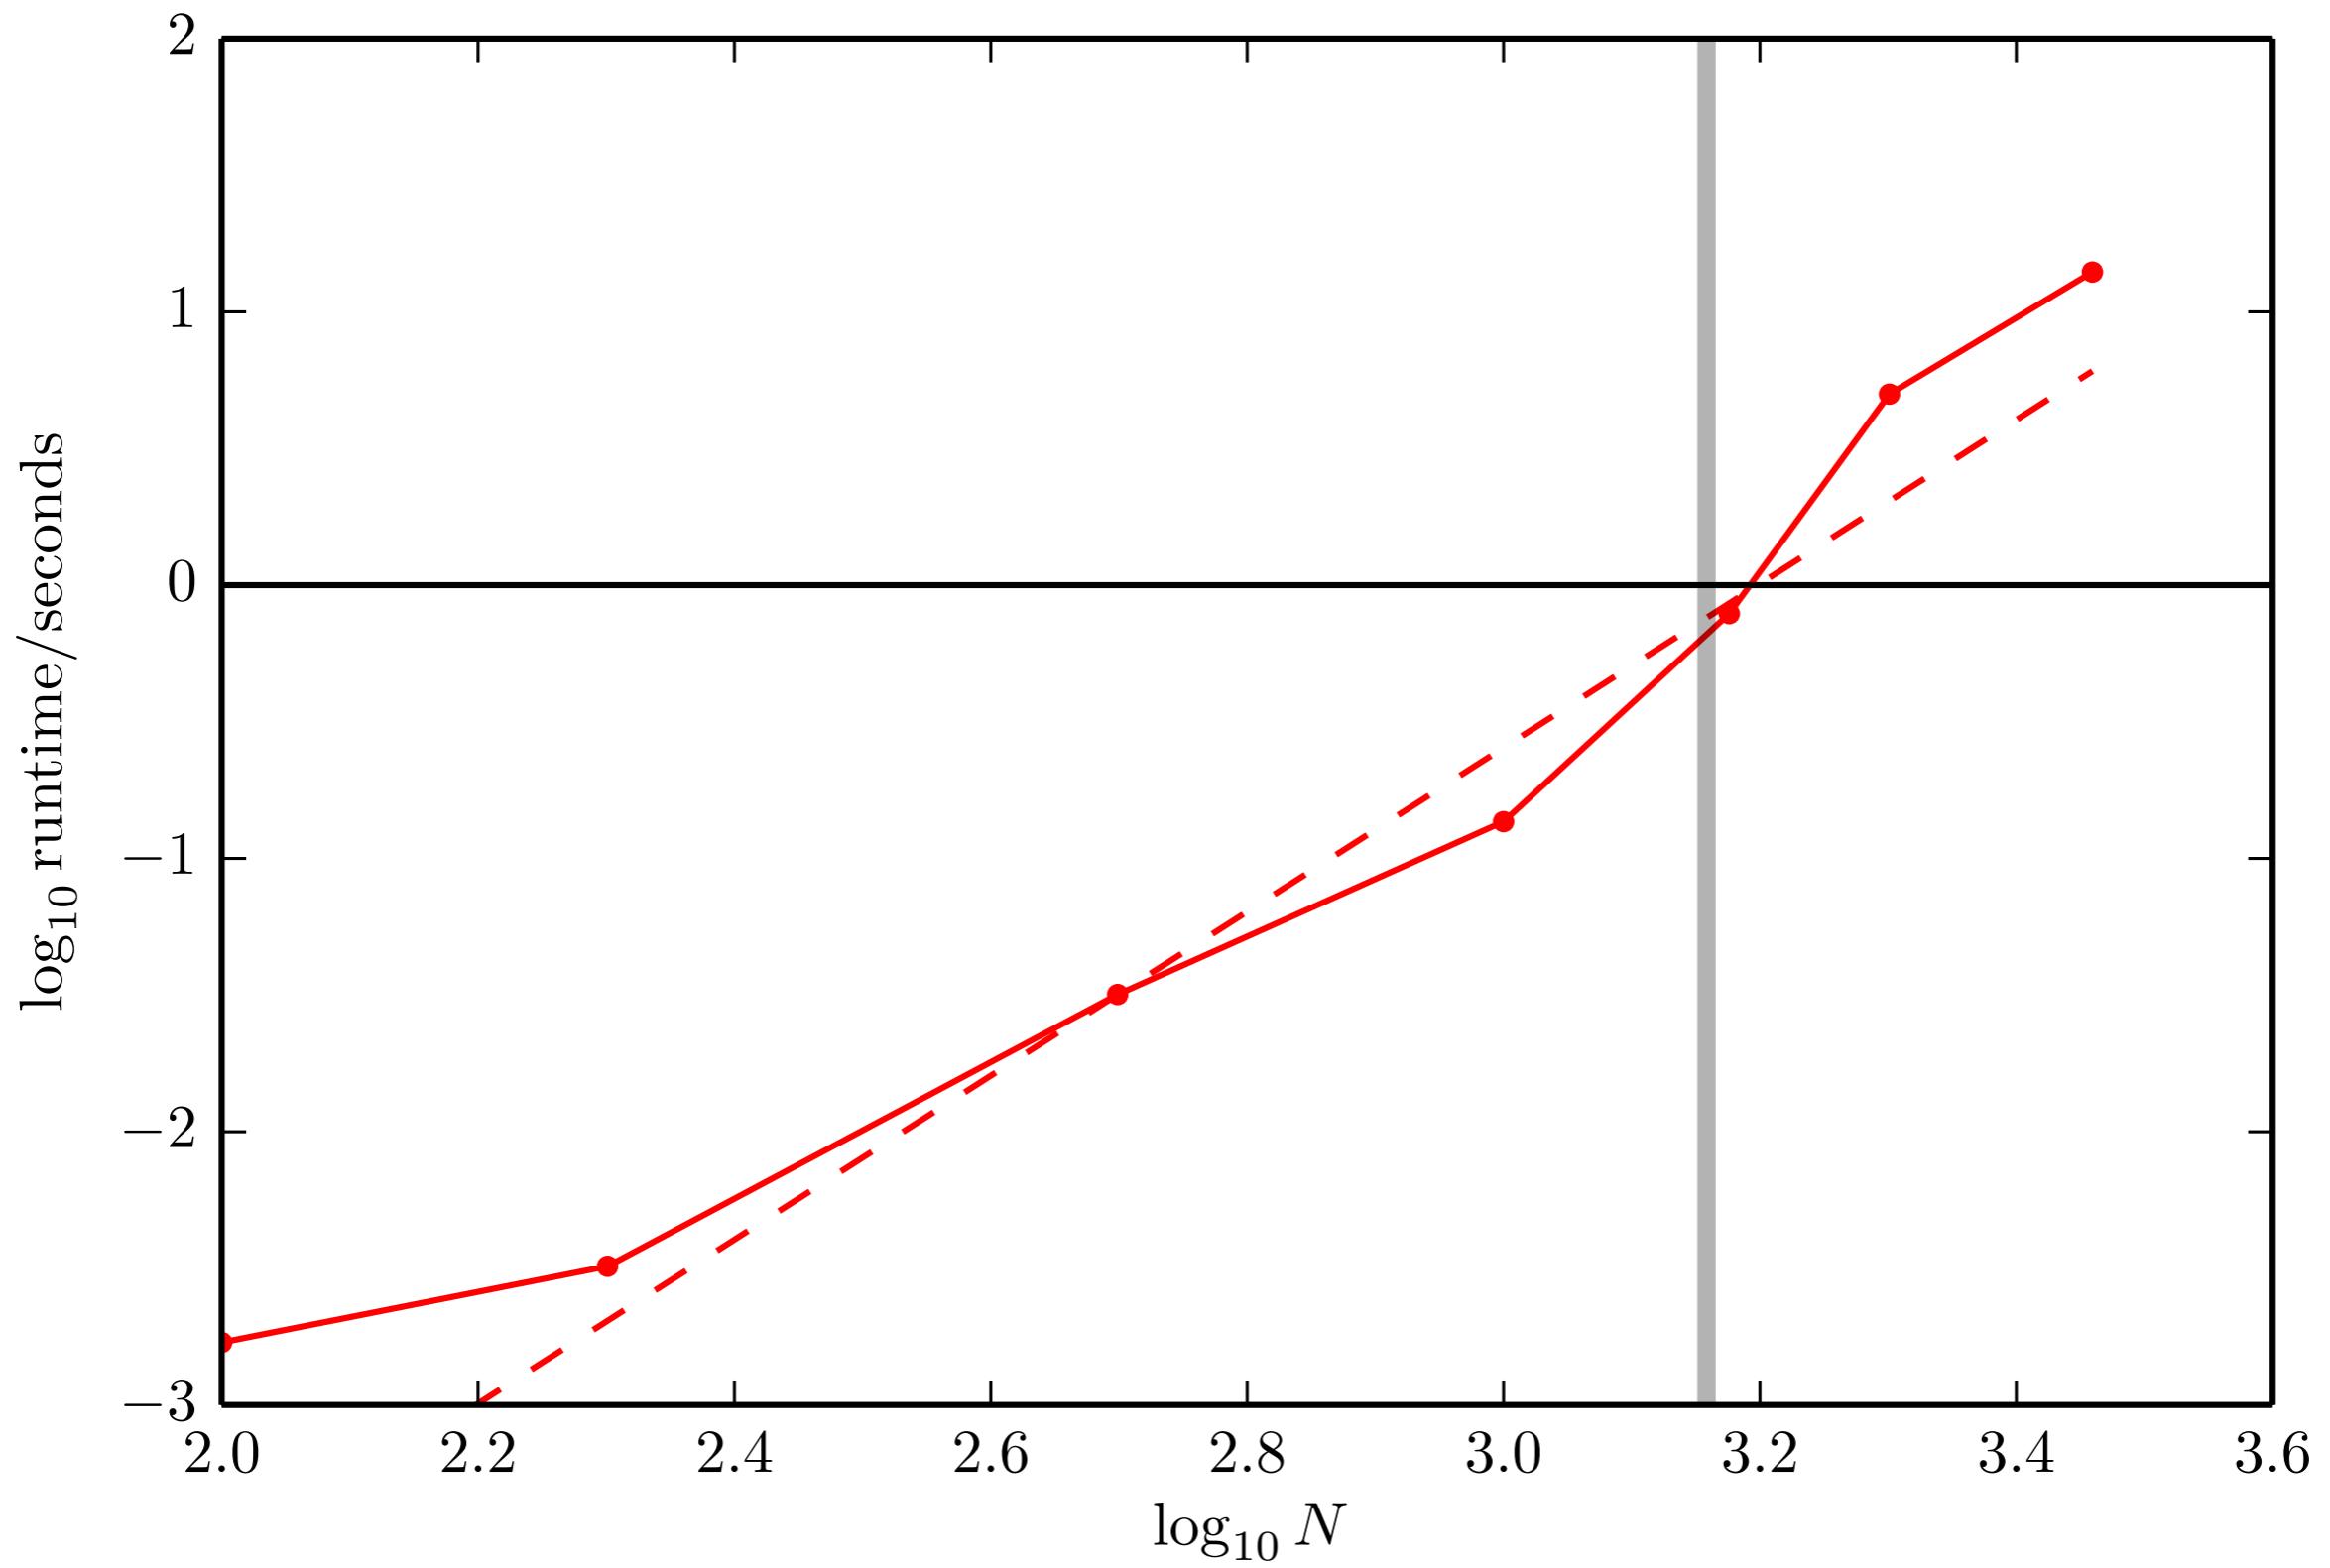
$$\begin{aligned}\log p(\mathbf{y} \mid \mathbf{x}, \boldsymbol{\sigma}, \boldsymbol{\theta}, \boldsymbol{\alpha}) = & -\frac{1}{2} [\mathbf{y} - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x})]^T K_{\boldsymbol{\alpha}}(\mathbf{x}, \boldsymbol{\sigma})^{-1} [\mathbf{y} - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x})] \\ & -\frac{1}{2} \log \det K_{\boldsymbol{\alpha}}(\mathbf{x}, \boldsymbol{\sigma}) - \frac{N}{2} \log 2 \pi\end{aligned}$$

compute **factorization** // evaluate **log-det** // apply **inverse**

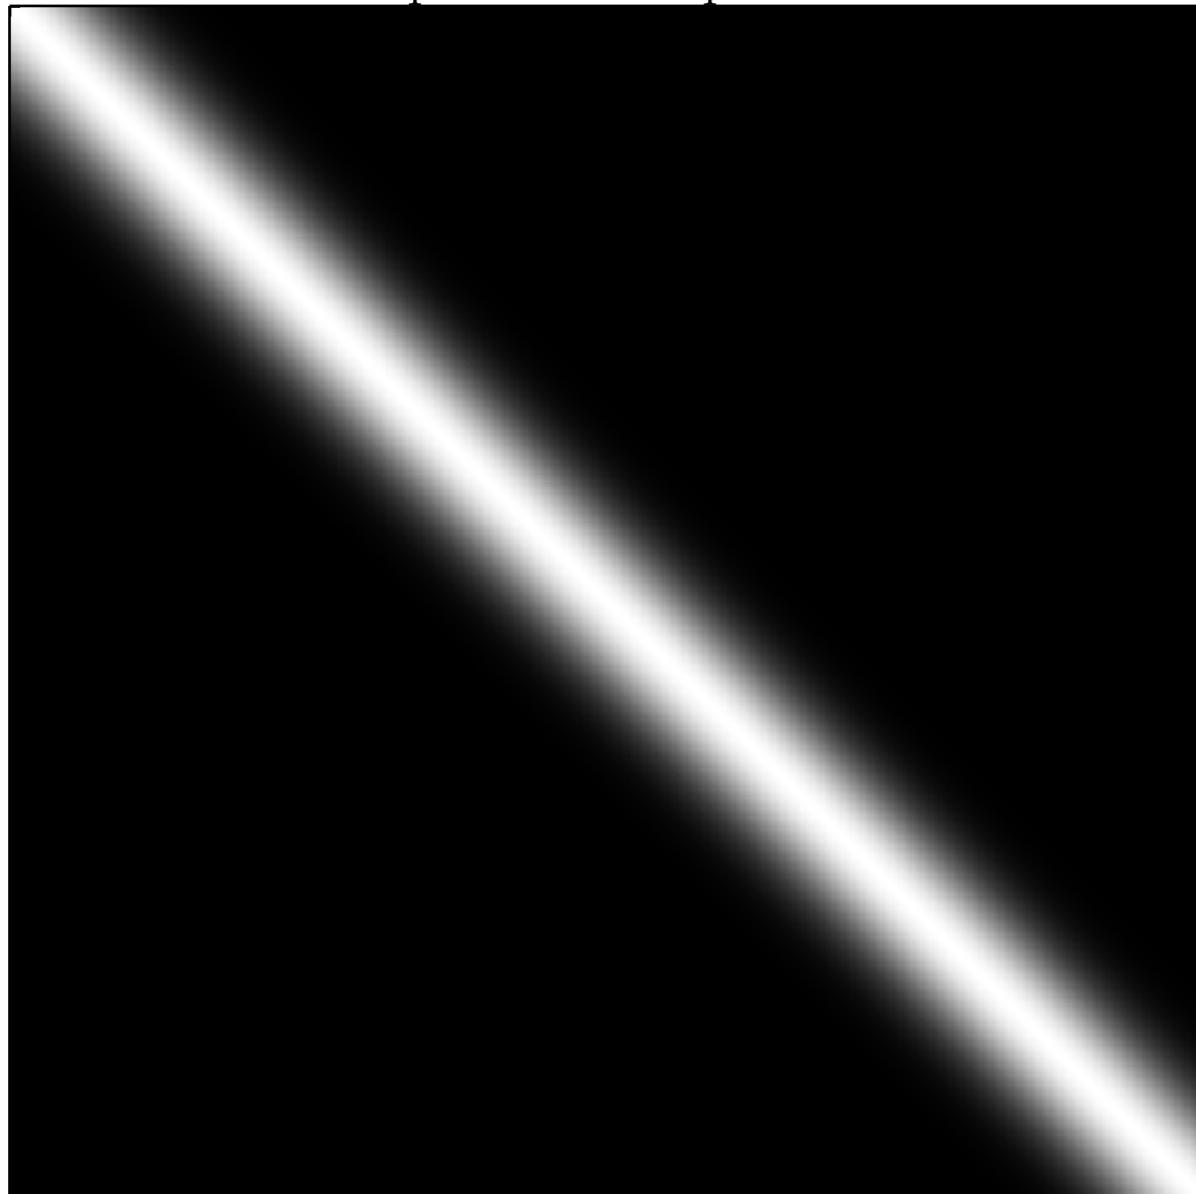
naïvely:  $\mathcal{O}(N^3)$

```
import numpy as np
from scipy.linalg import cho_factor, cho_solve

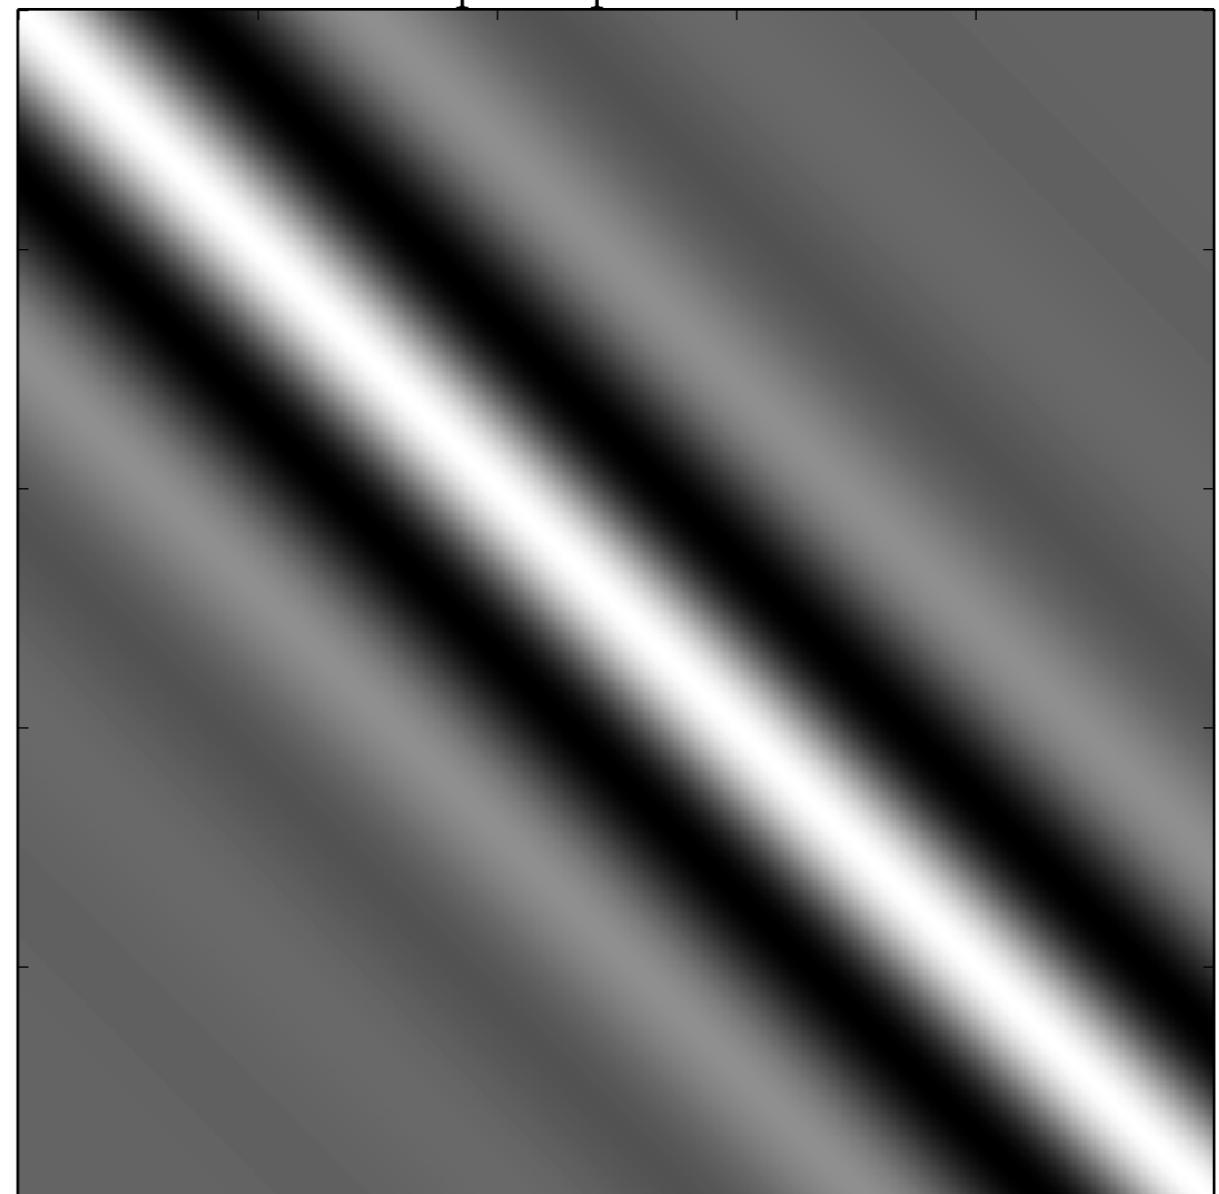
def simple_gp_lnlike(x, y, yerr, a, s):
    r = x[:, None] - x[None, :]
    C = np.diag(yerr**2) + a*np.exp(-0.5*r**2/(s*s))
    factor, flag = cho_factor(C)
    logdet = 2*np.sum(np.log(np.diag(factor)))
    return -0.5 * (np.dot(y, cho_solve((factor, flag), y)) +
                   logdet + len(x)*np.log(2*np.pi))
```



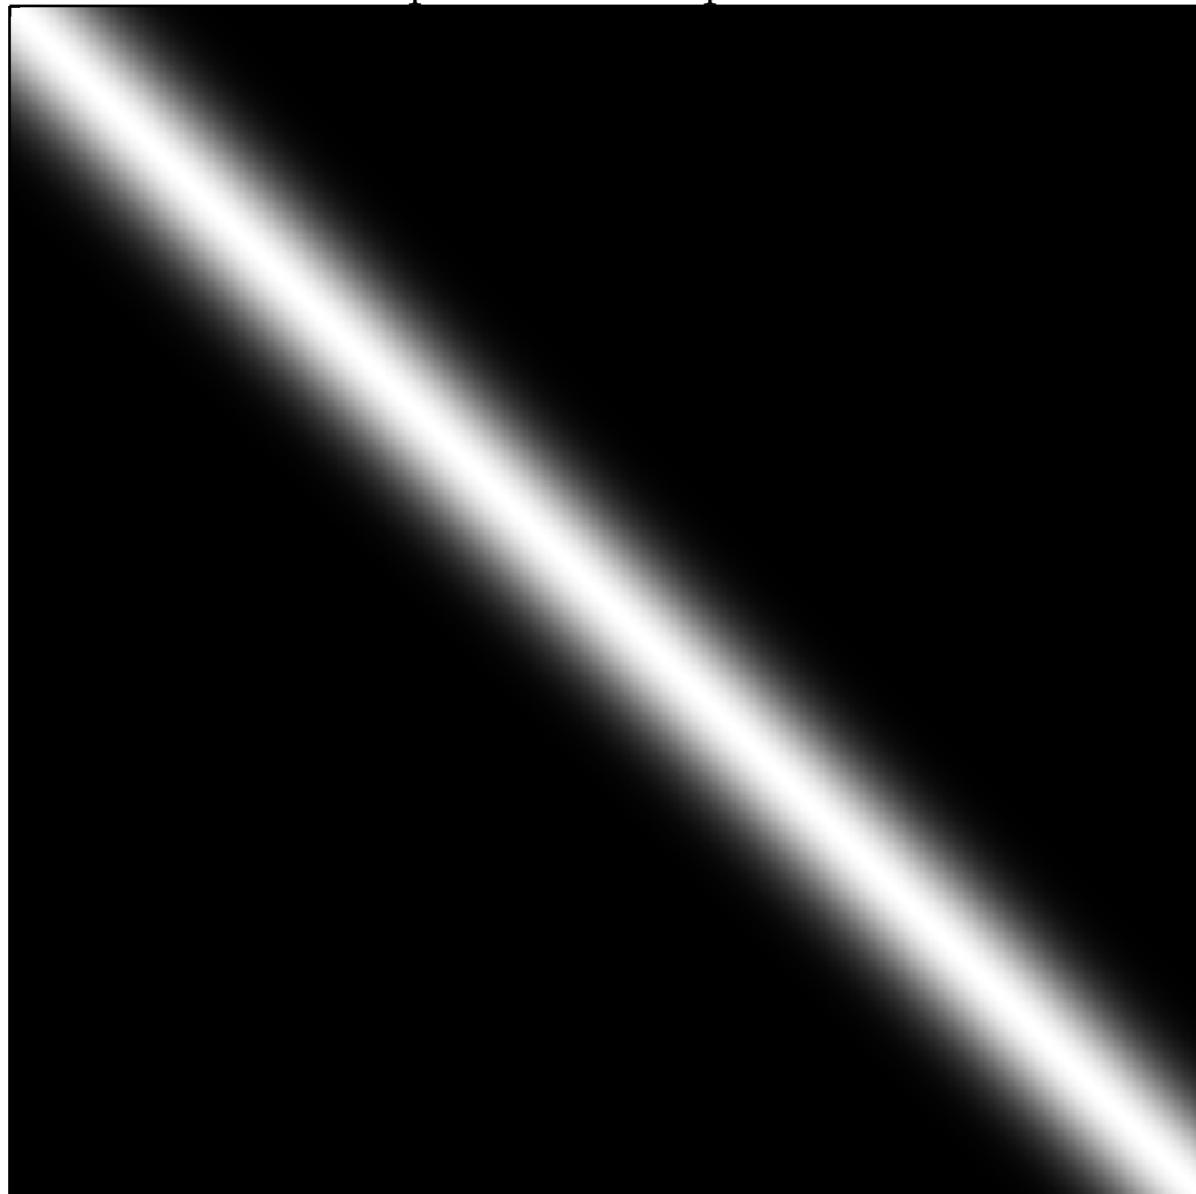
exponential squared



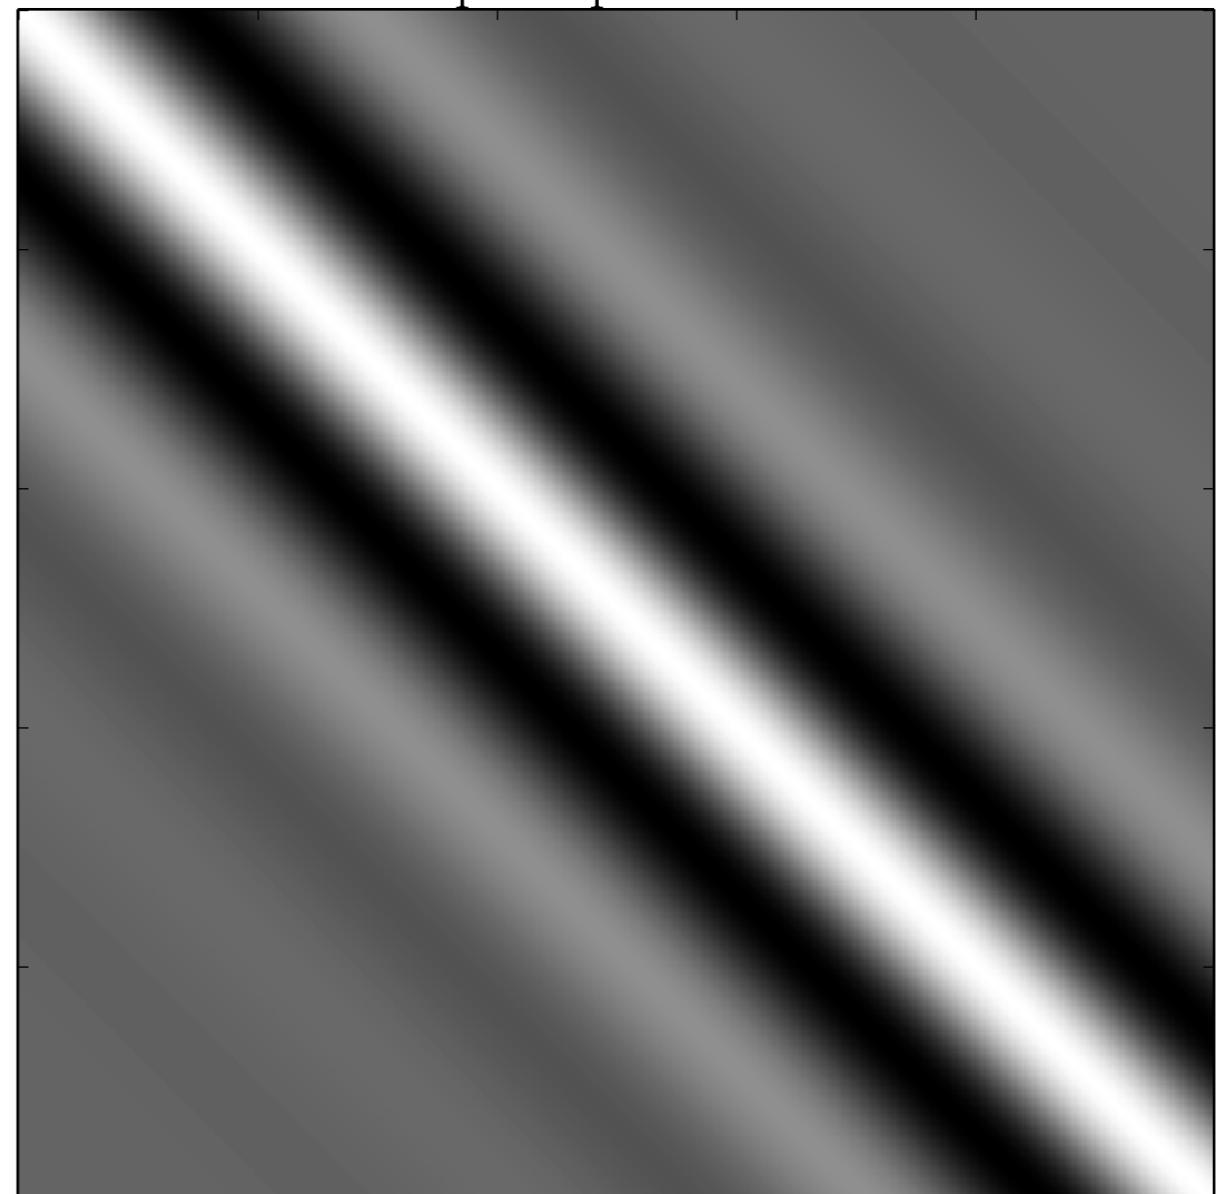
quasi-periodic

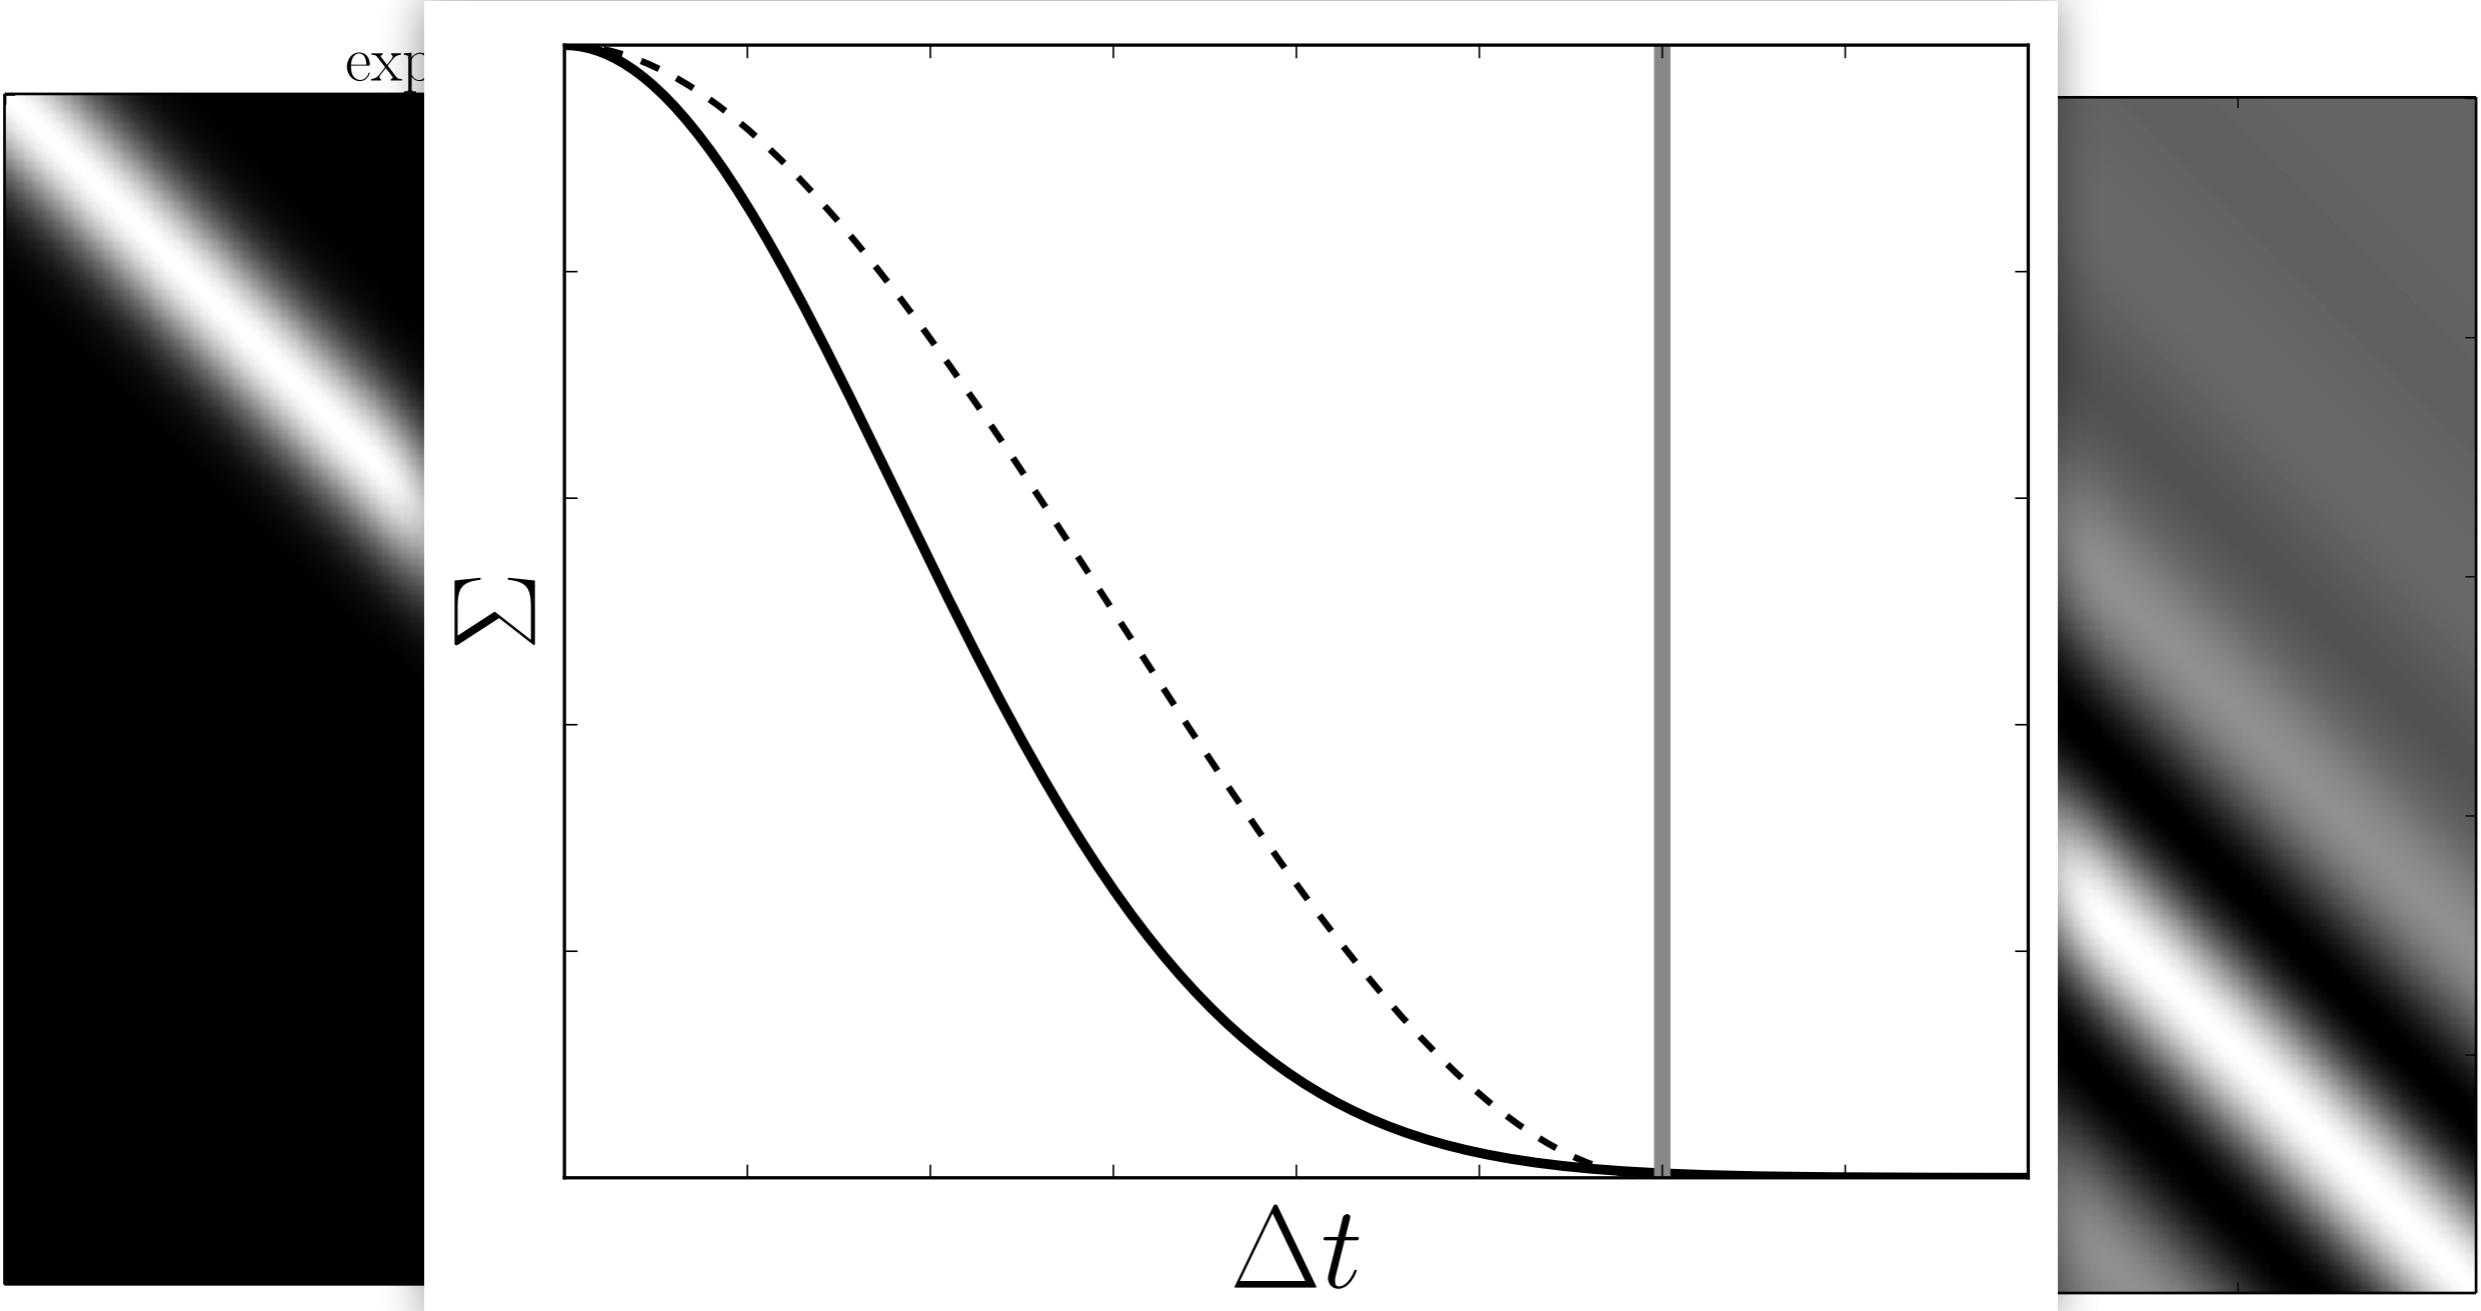


exponential squared



quasi-periodic



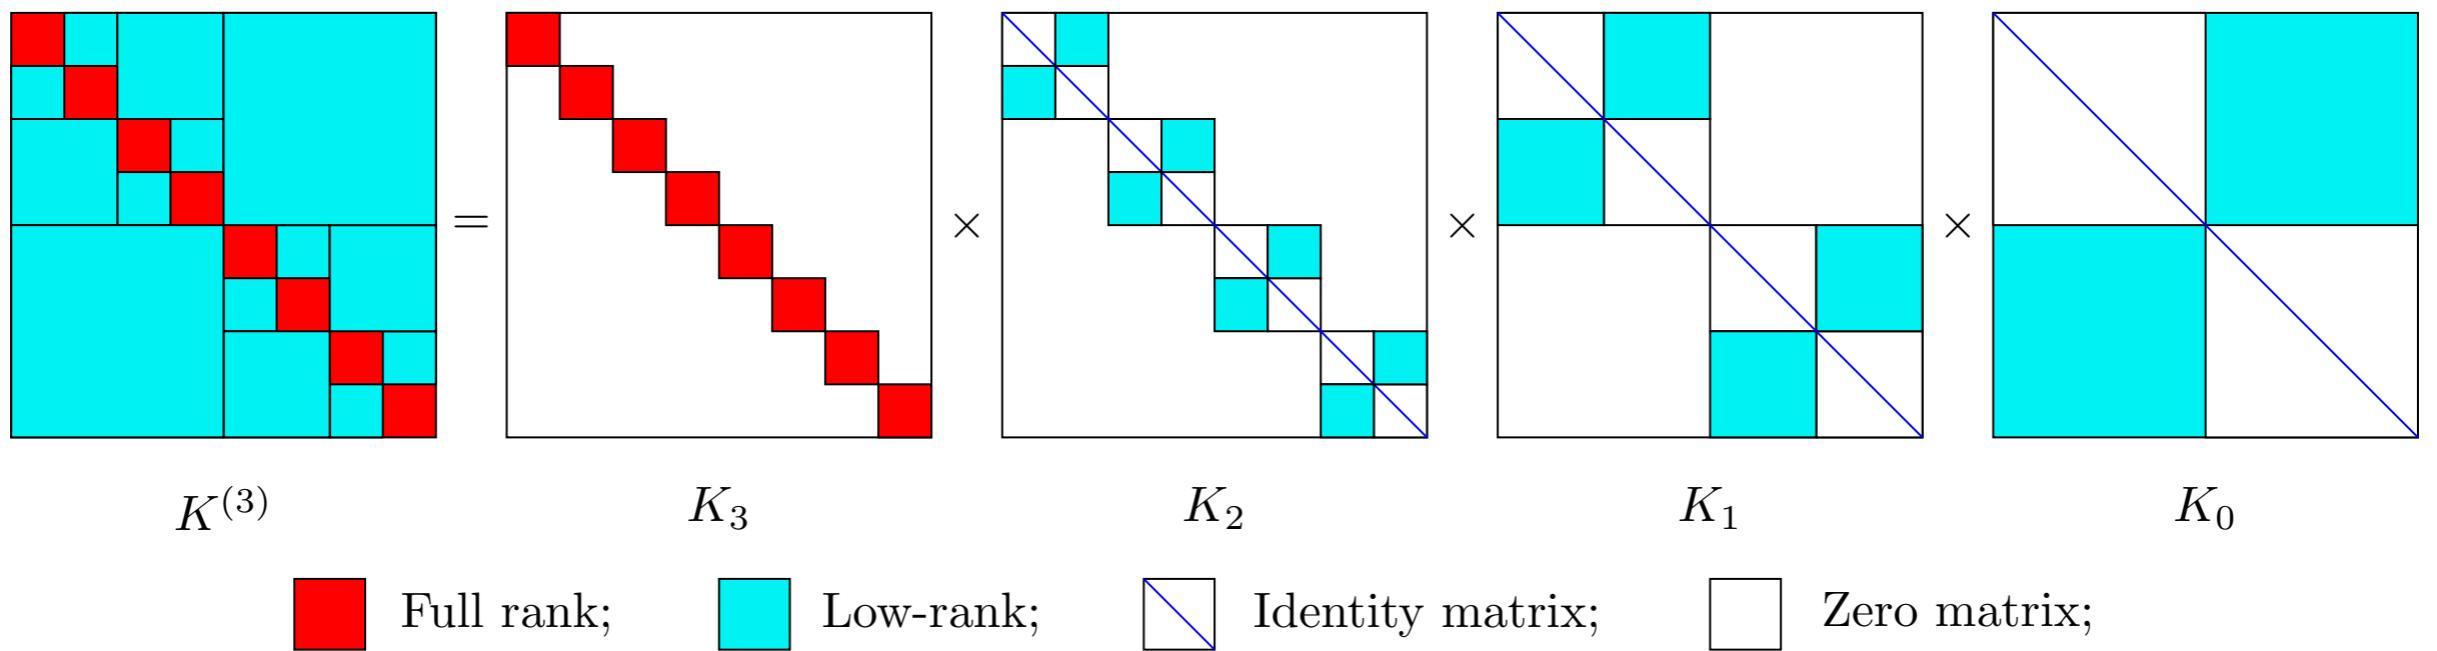




Aren't kernel matrices **Hierarchical Off-Diagonal Low-Rank?**

---

— no astronomer ever

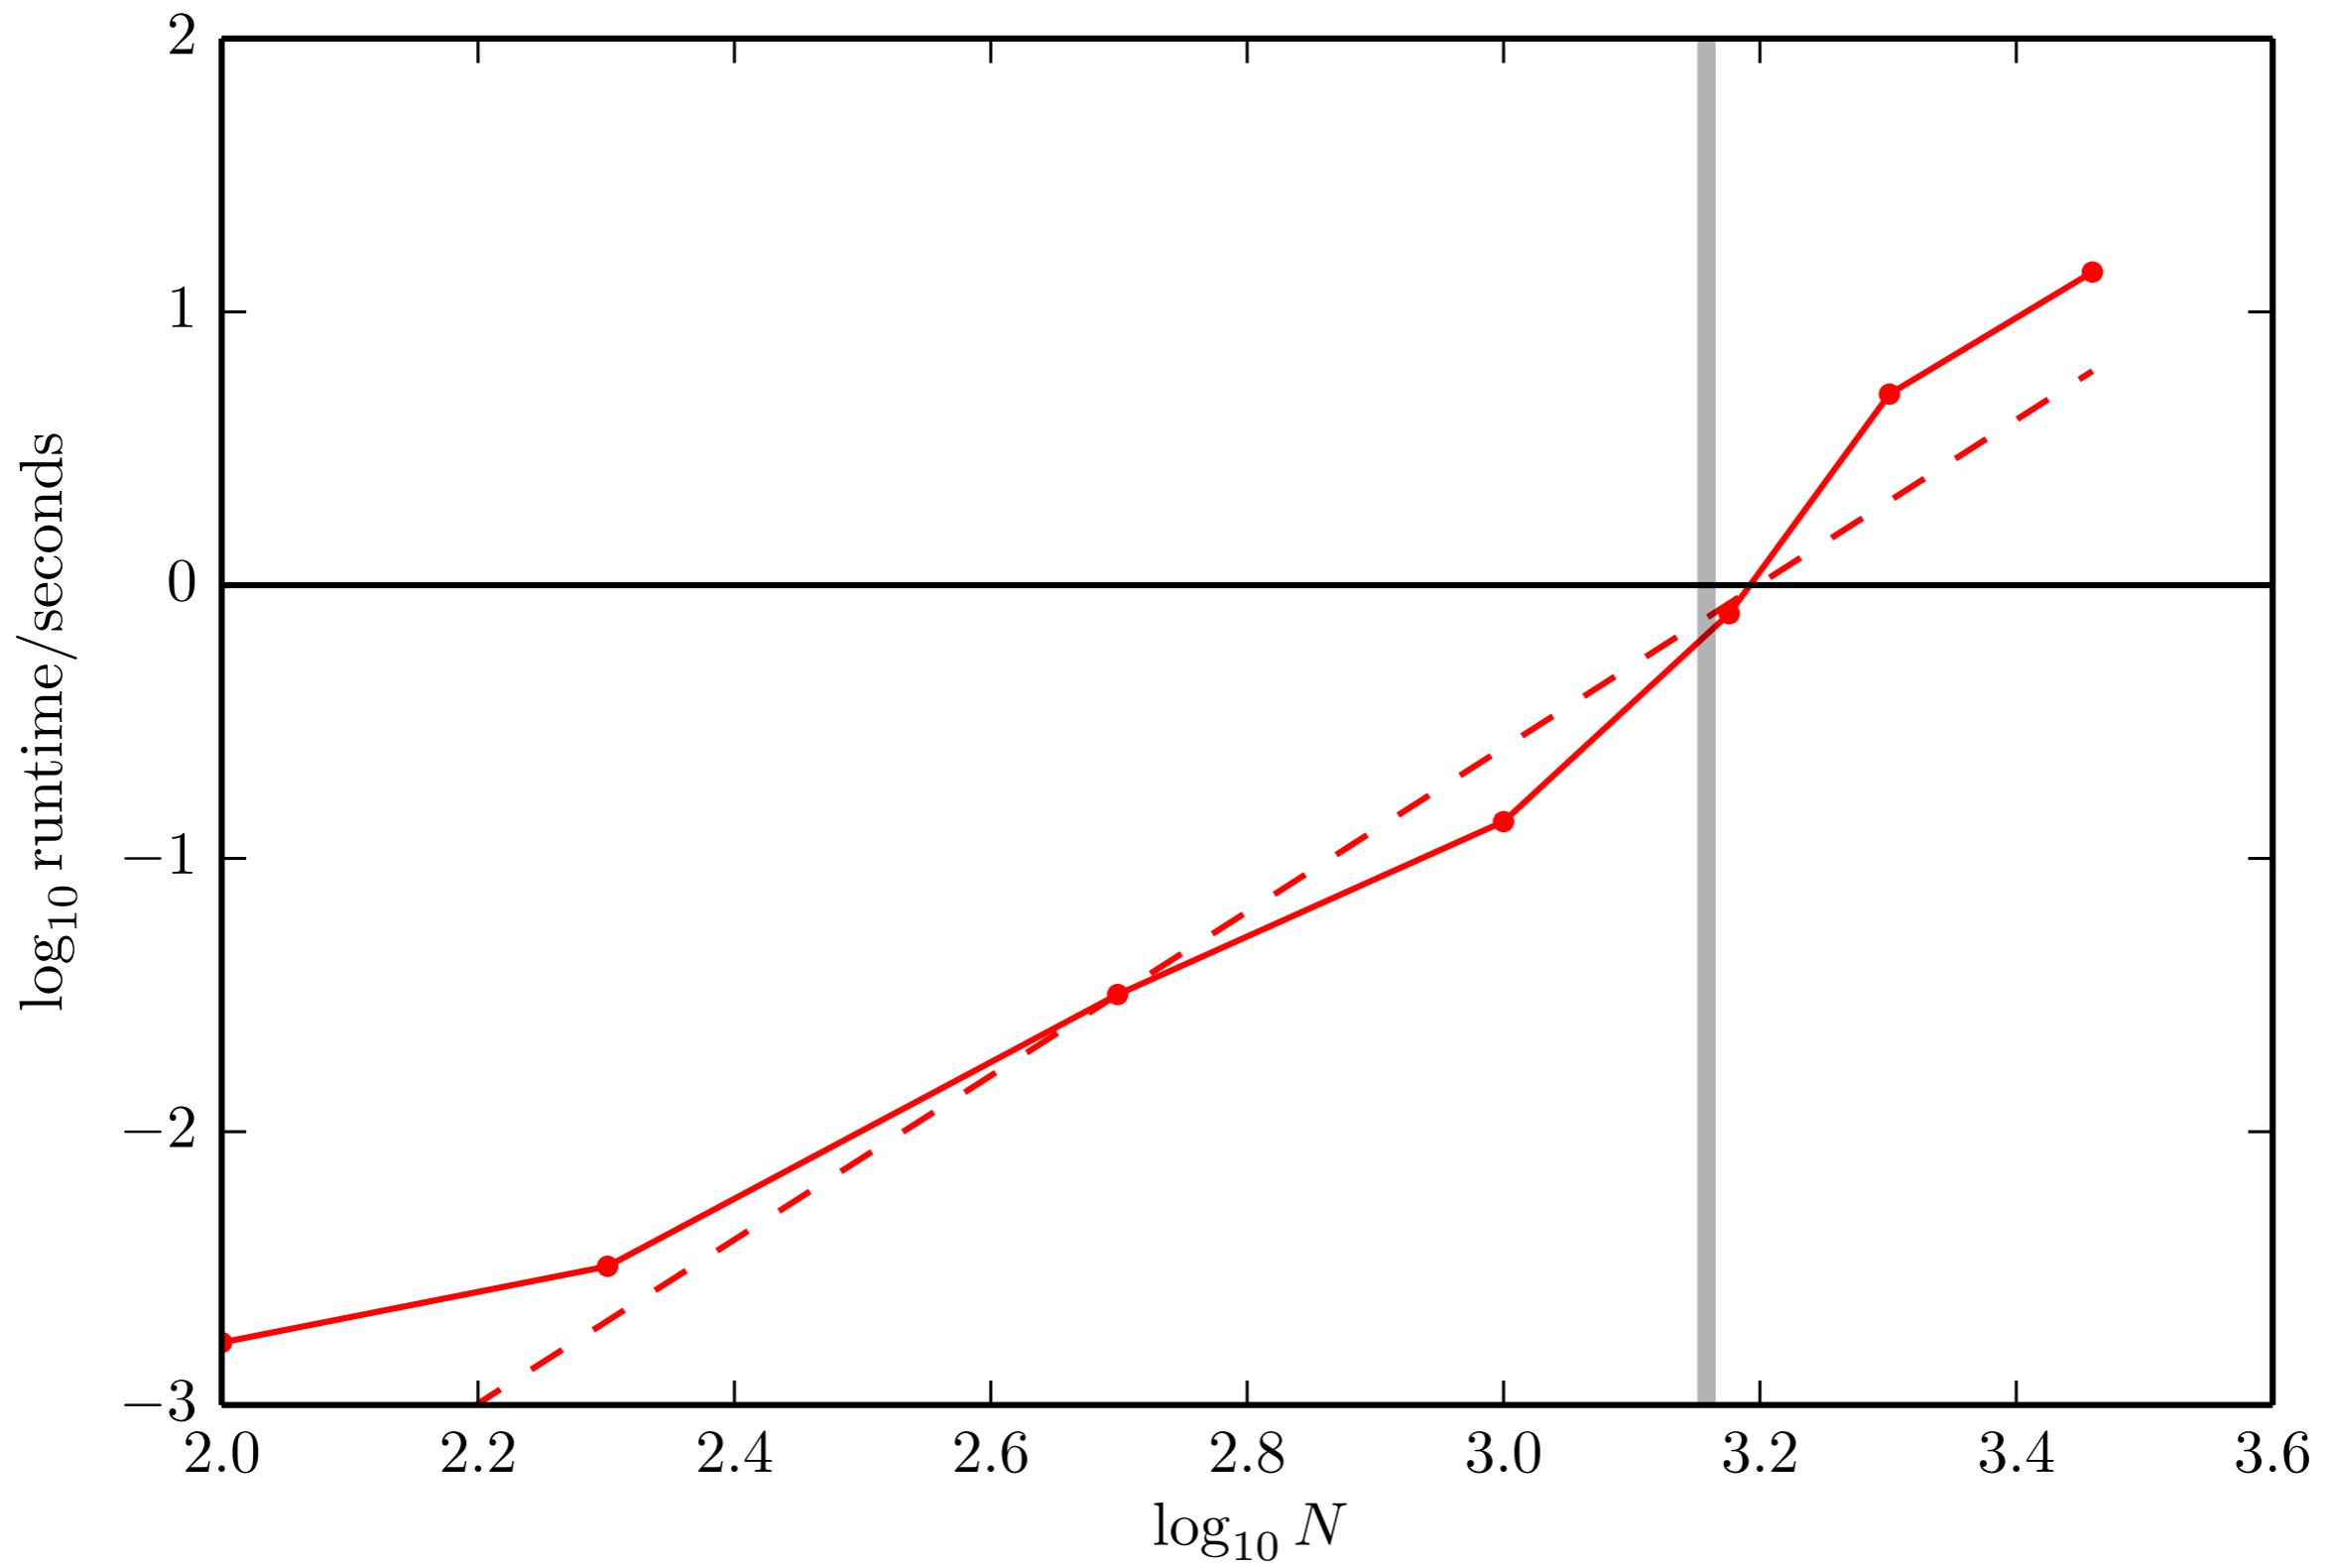


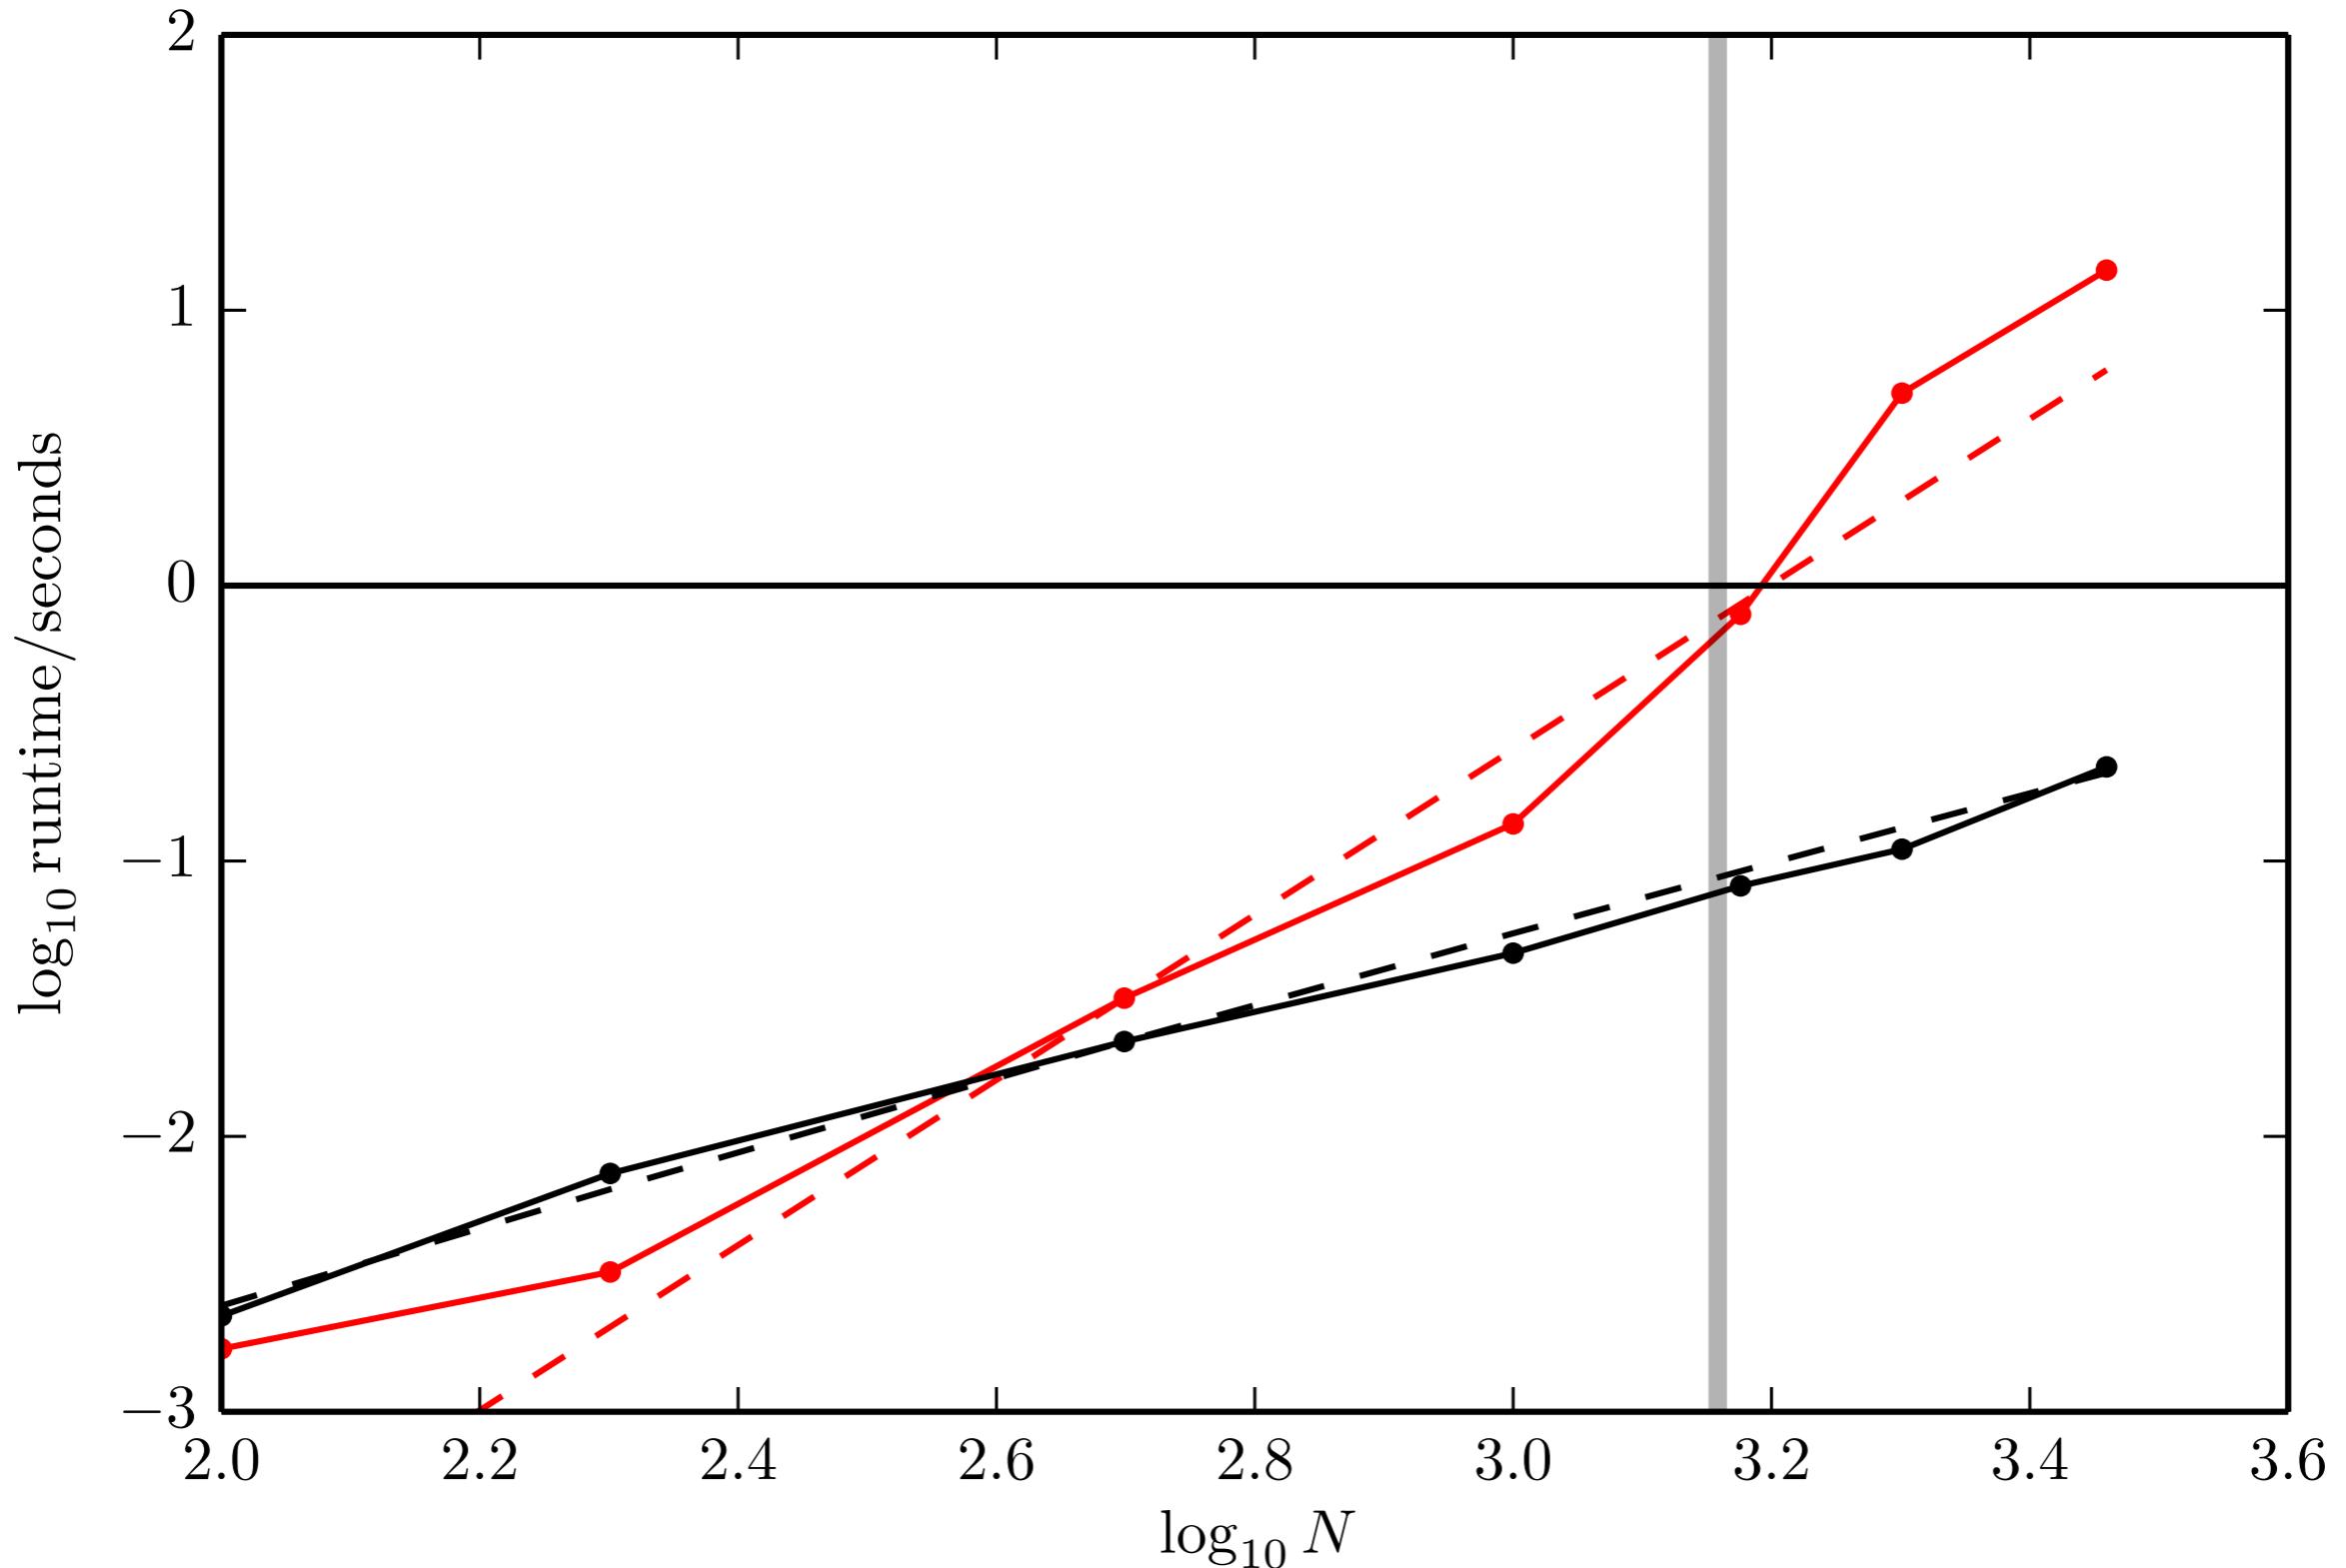
github.com/dfm/george

---

```
import numpy as np
from george import GaussianProcess, kernels

def george_lnlike(x, y, yerr, a, s):
    kernel = a * kernels.RBFKernel(s)
    gp = GaussianProcess(kernel)
    gp.compute(x, yerr)
    return gp.lnlikelihood(y)
```





and **short cadence** data?

---

**one month** of data in **4 seconds**



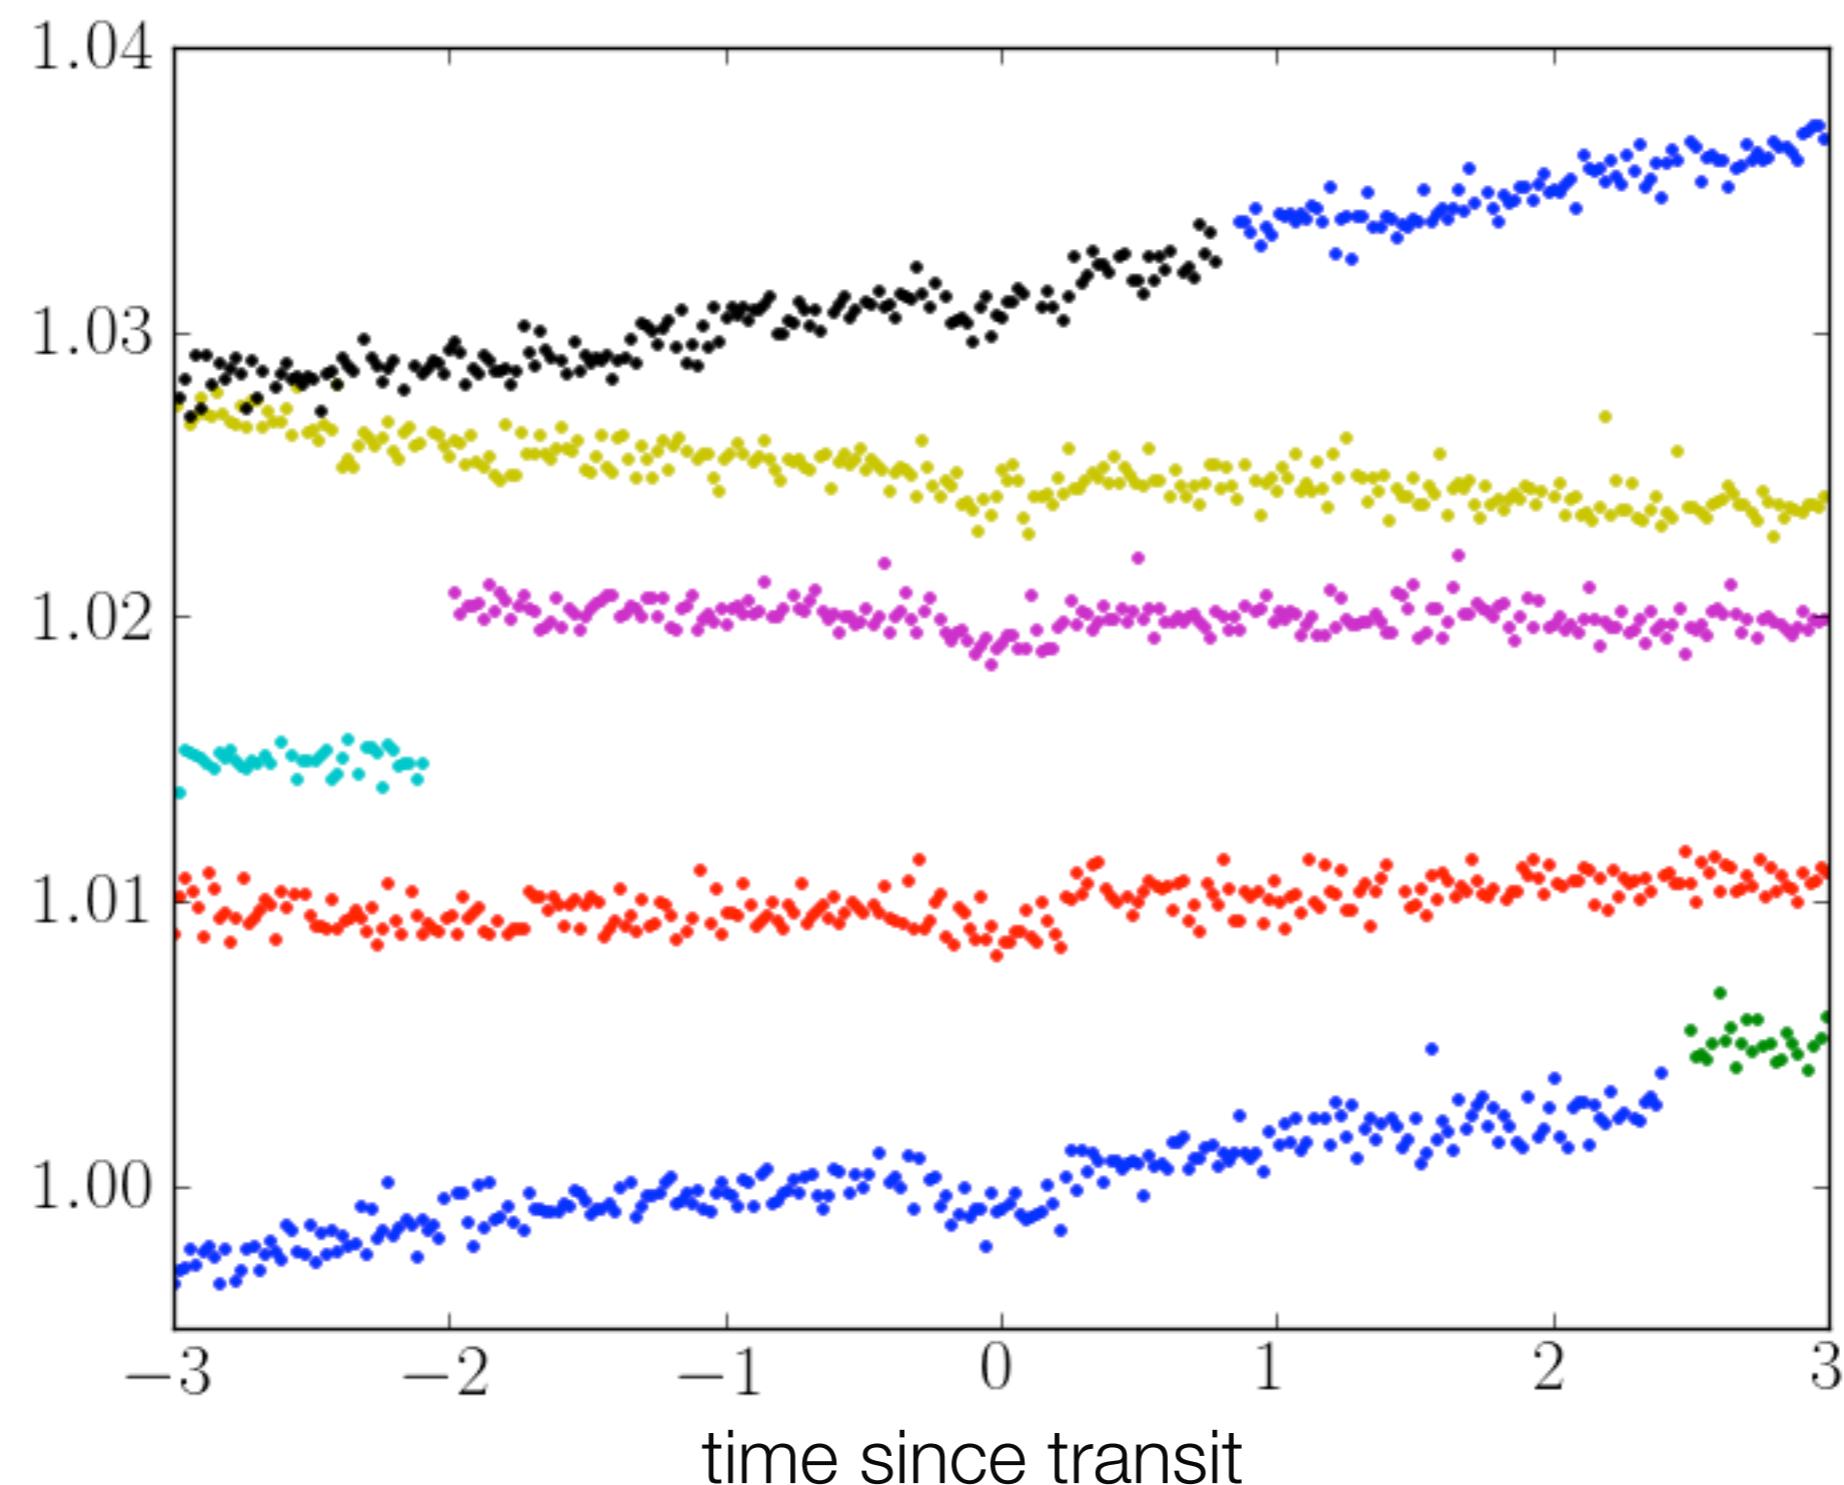
3

## Applications to Kepler data.

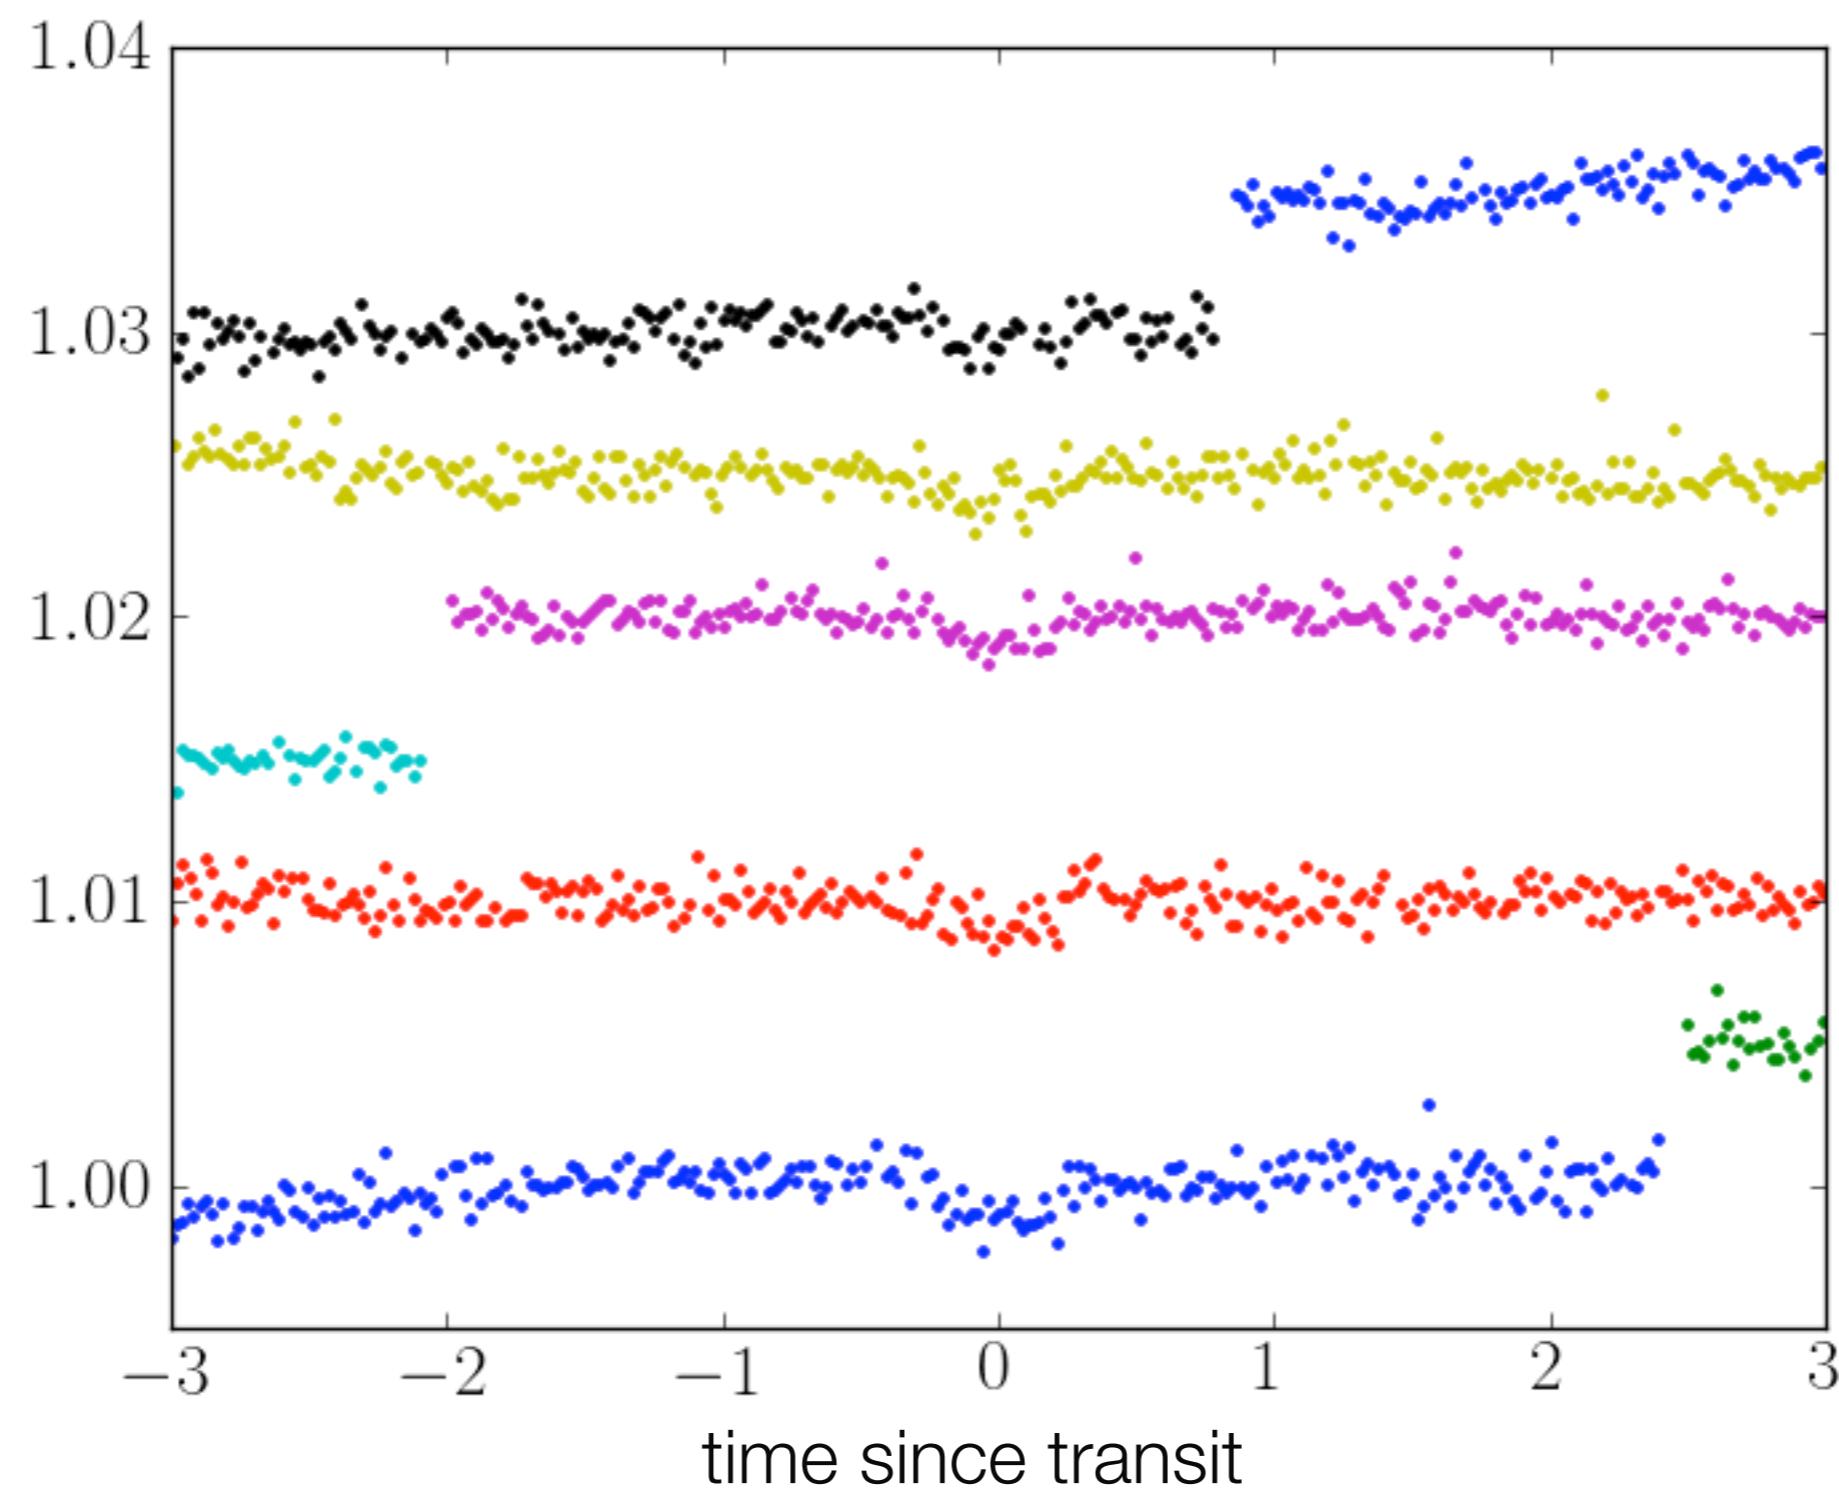
---

# Parameter Recovery

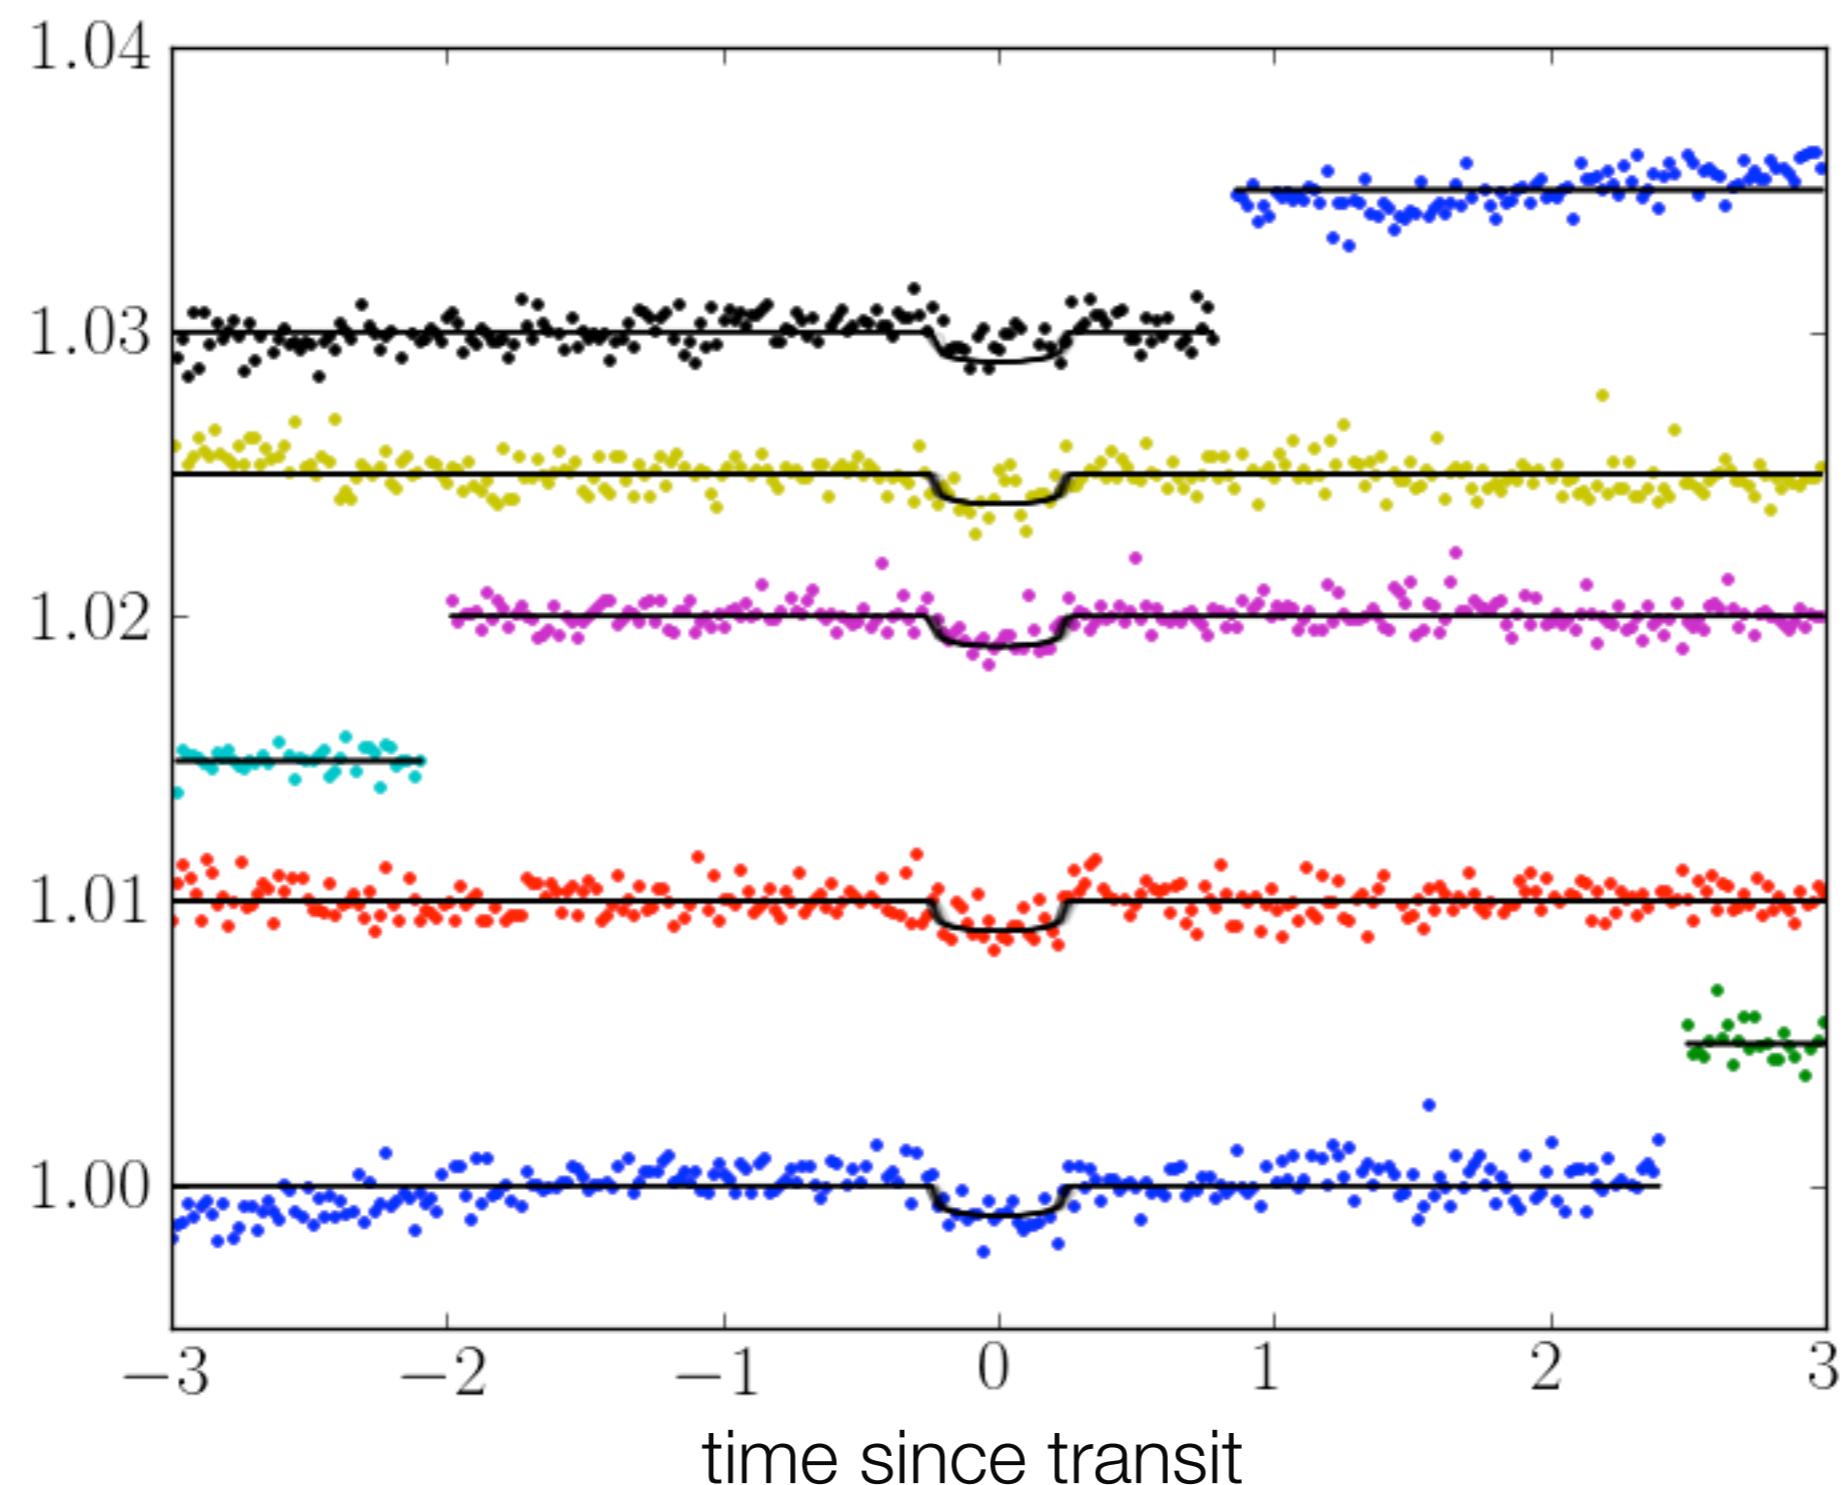
---



KIC 2301306 + injection

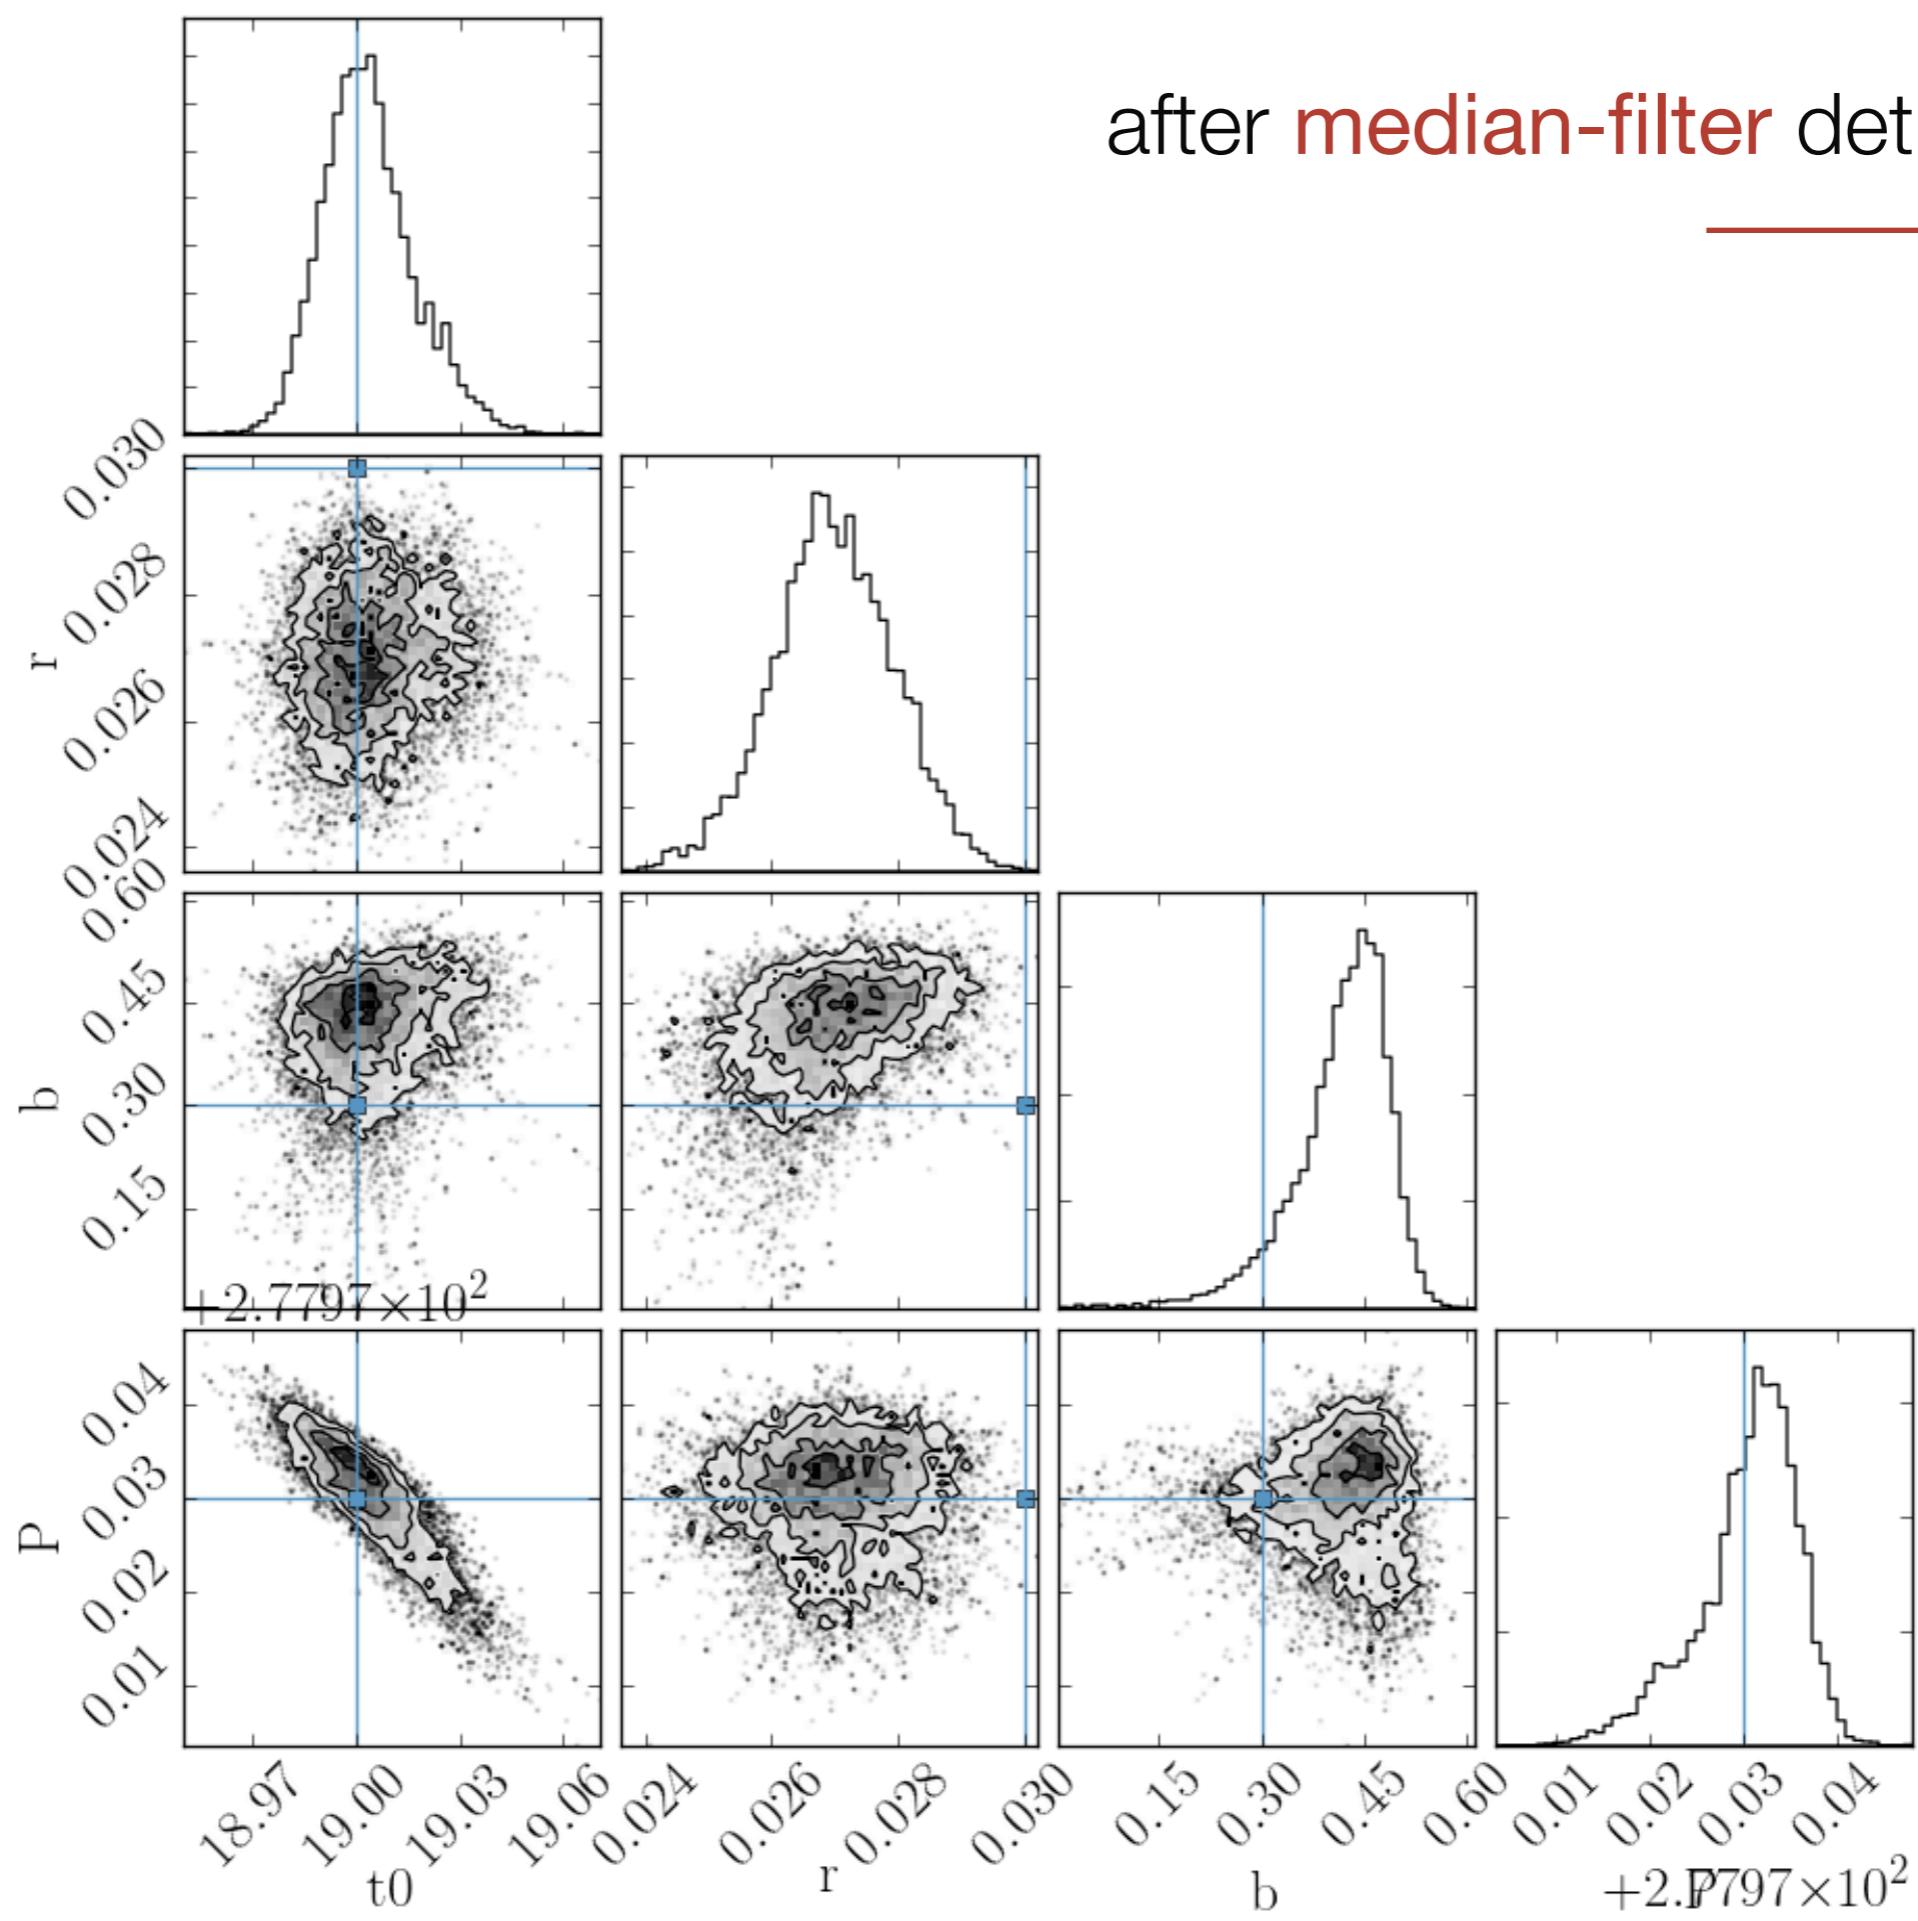


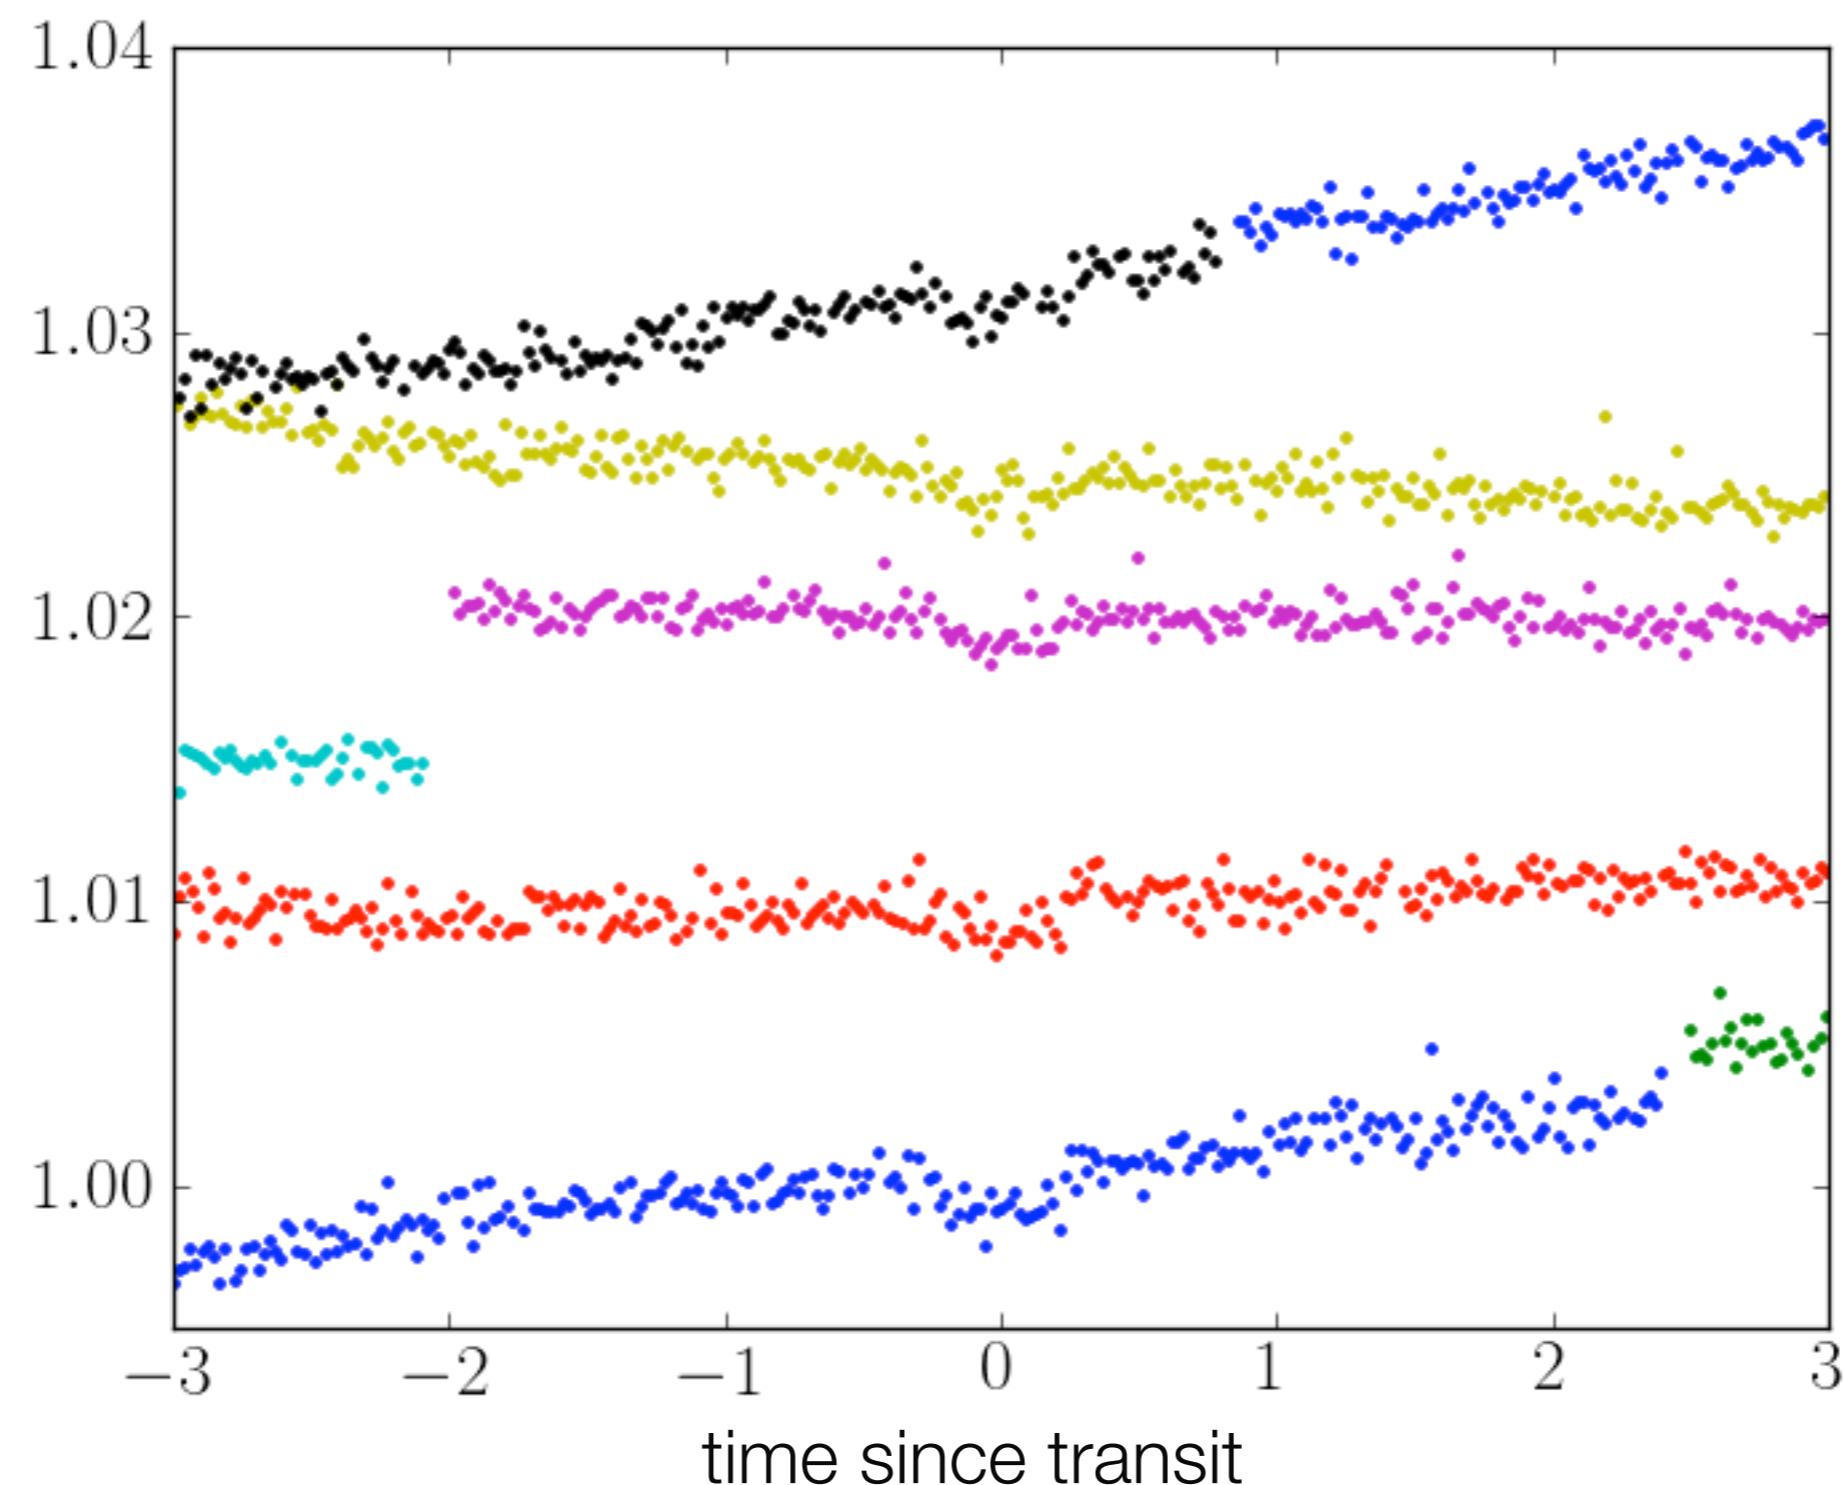
KIC 2301306 + injection



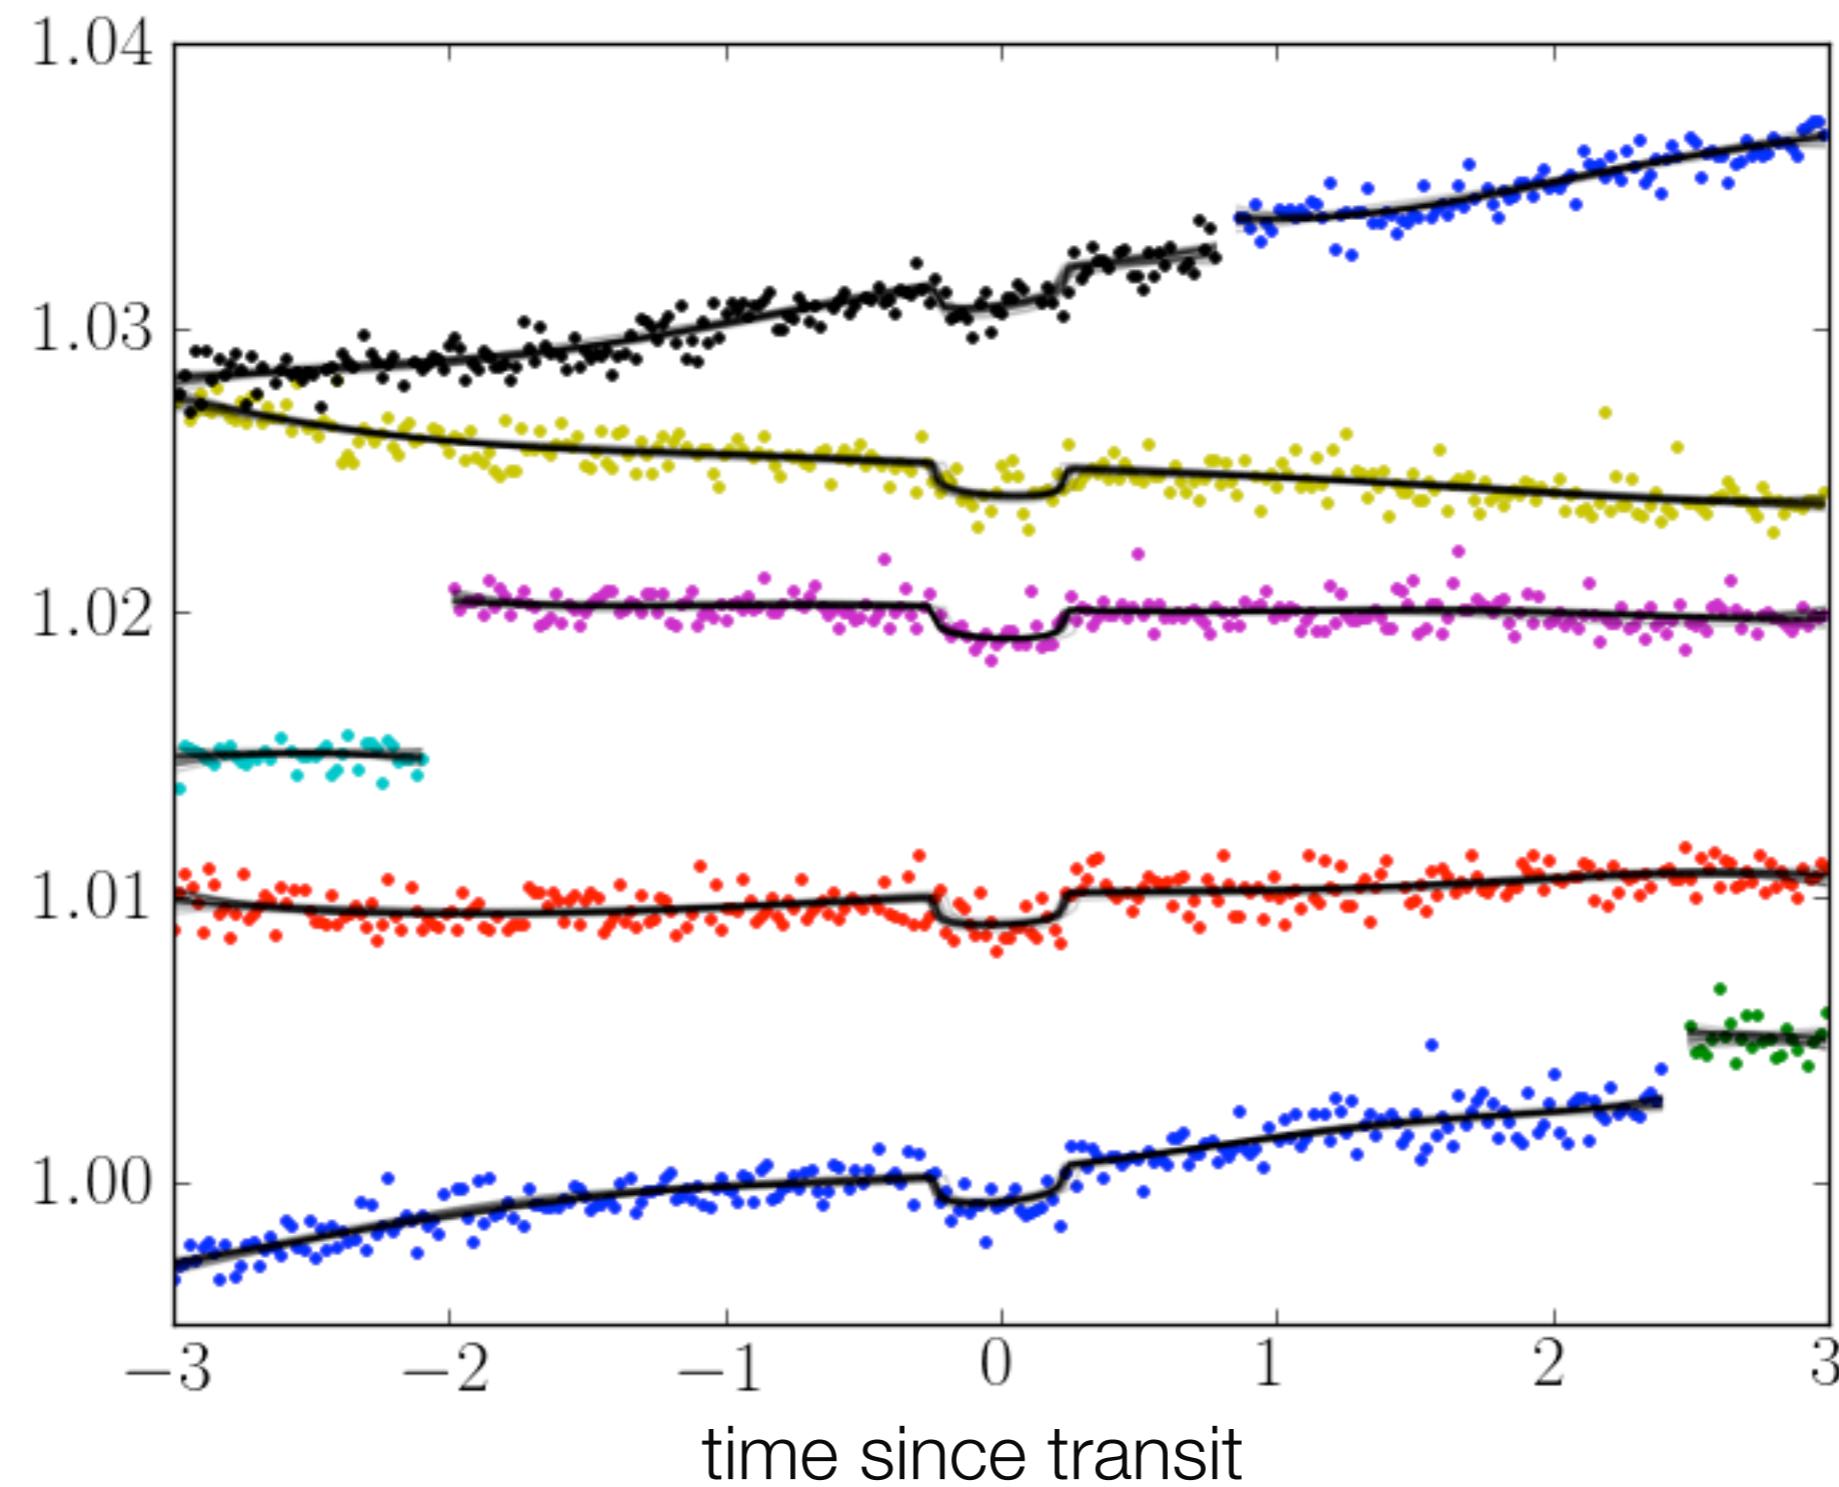
KIC 2301306 + injection

after median-filter detrending





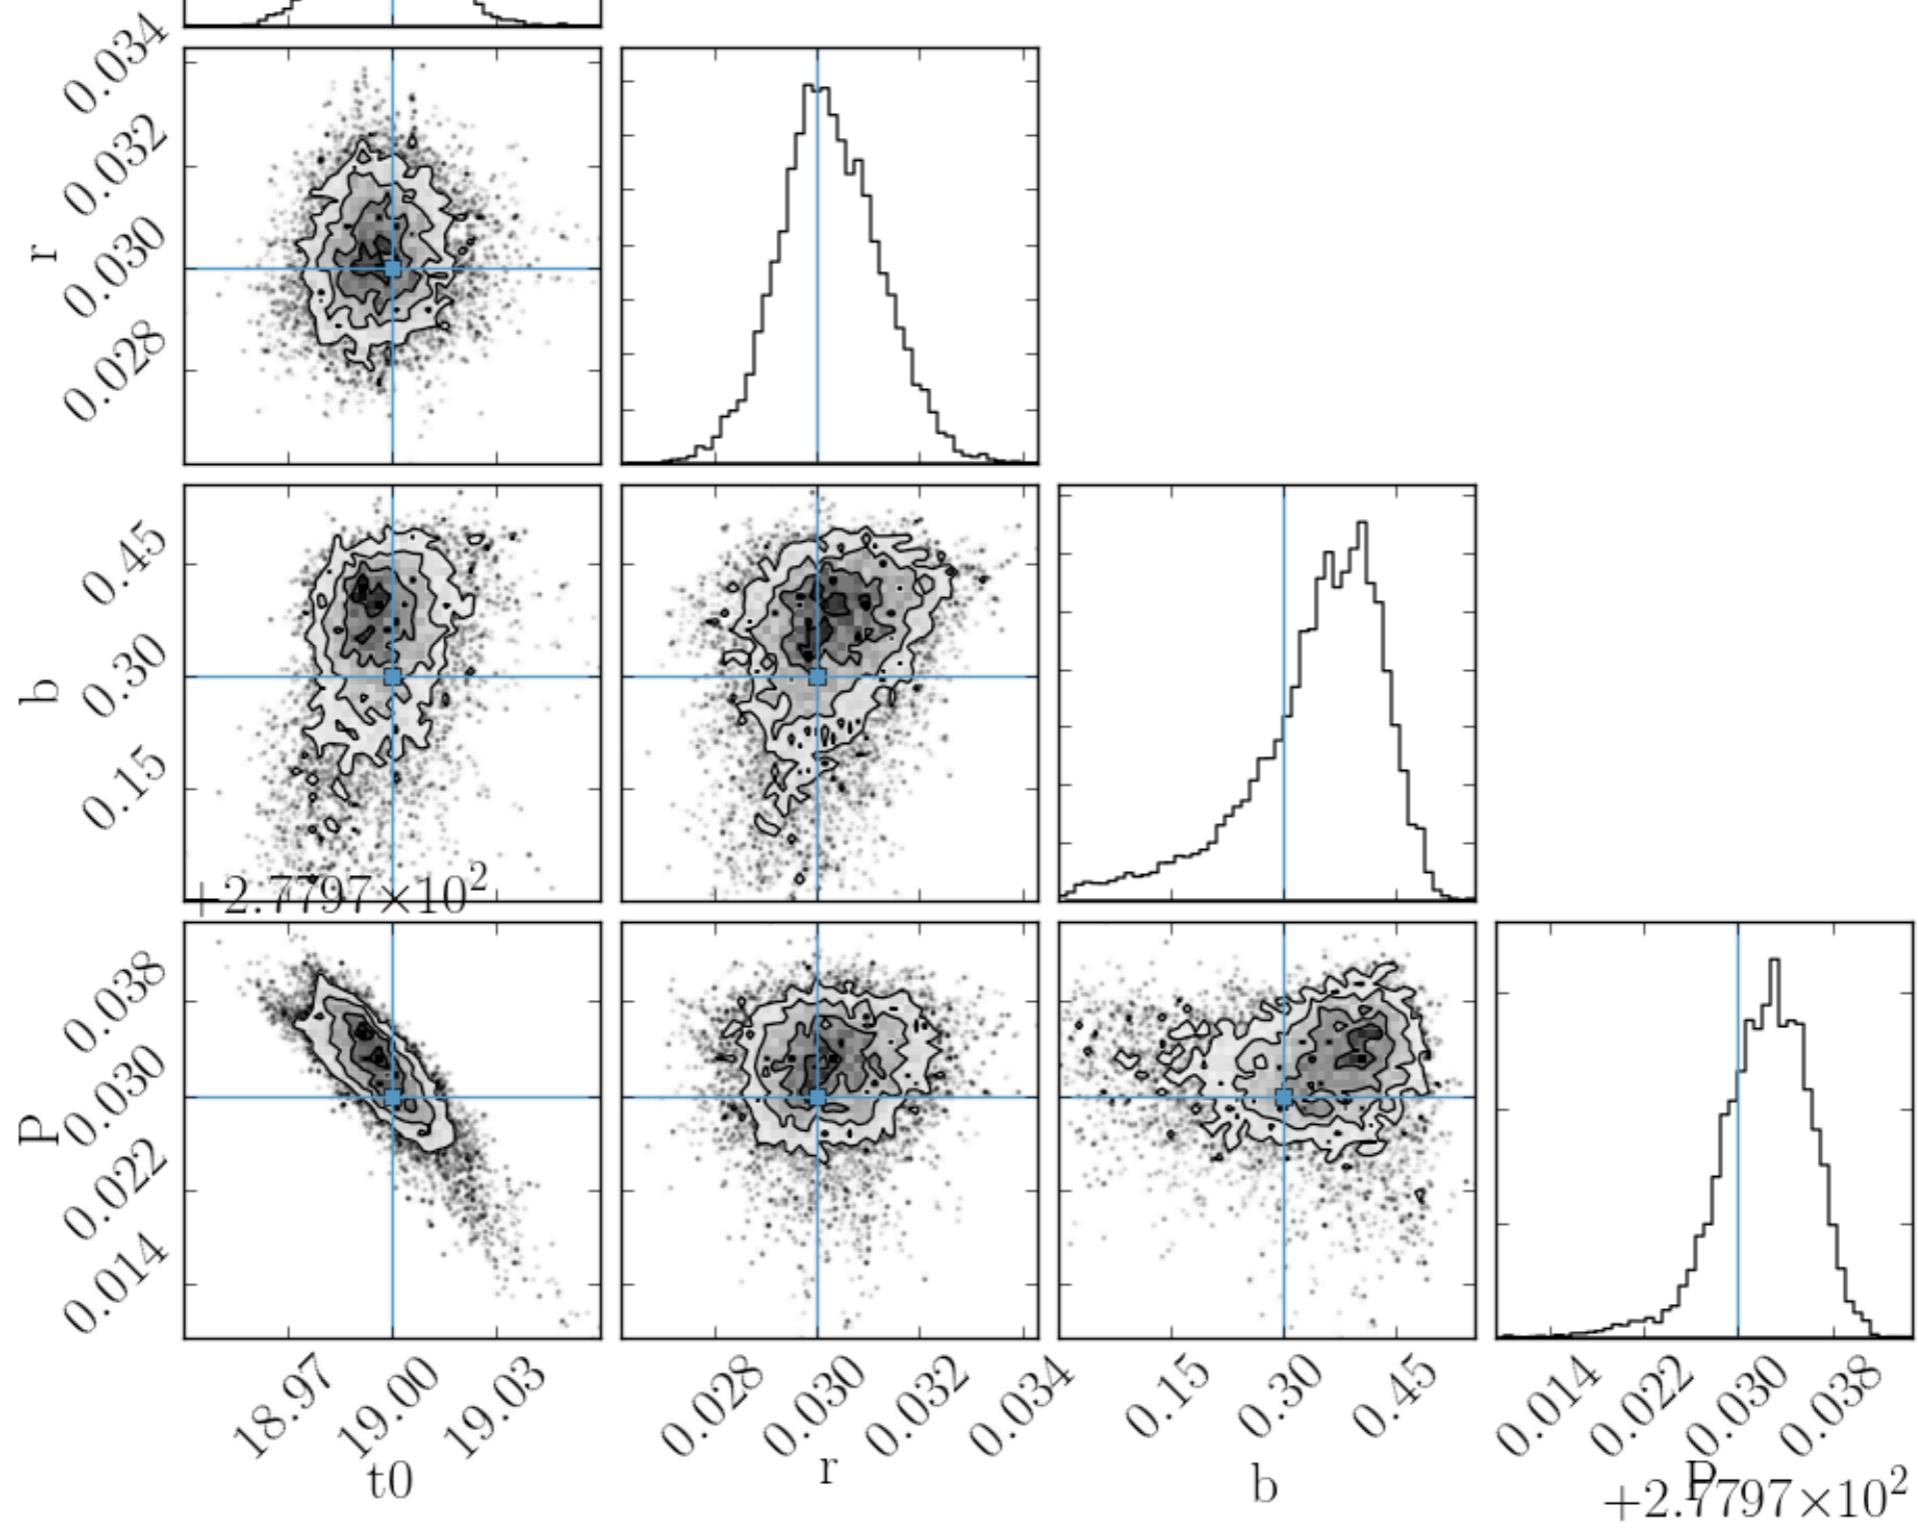
KIC 2301306 + injection

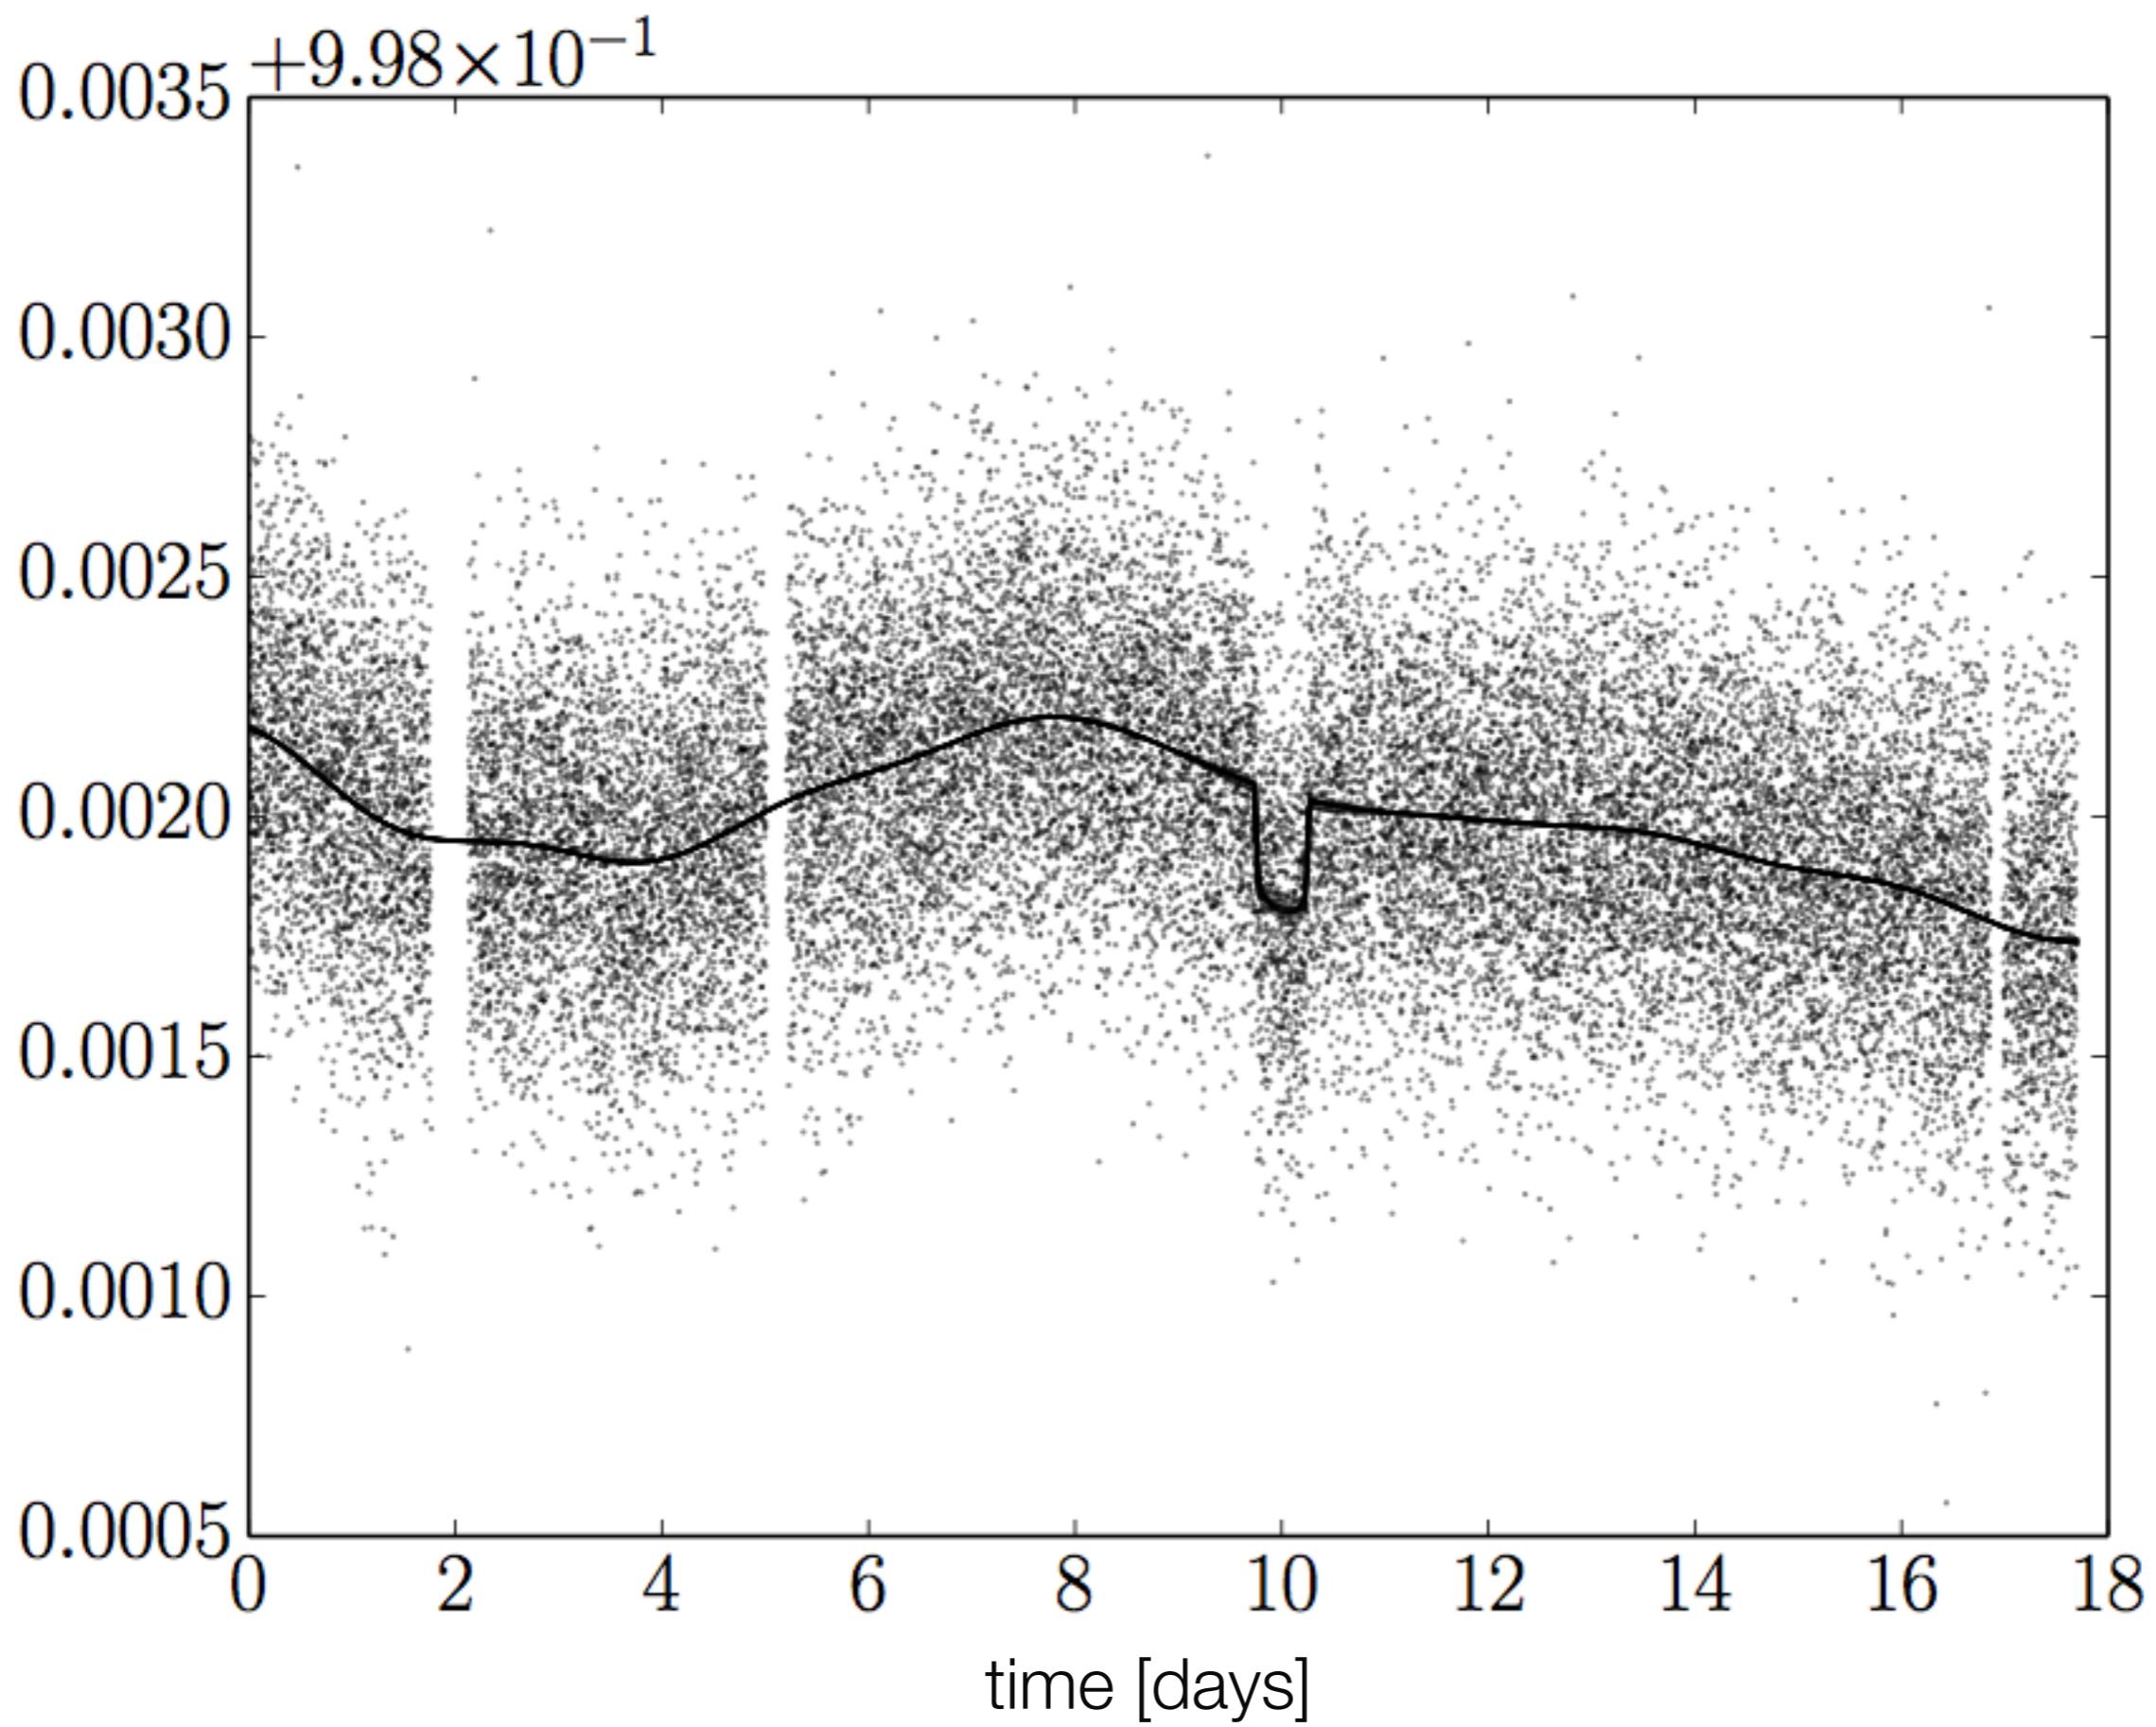


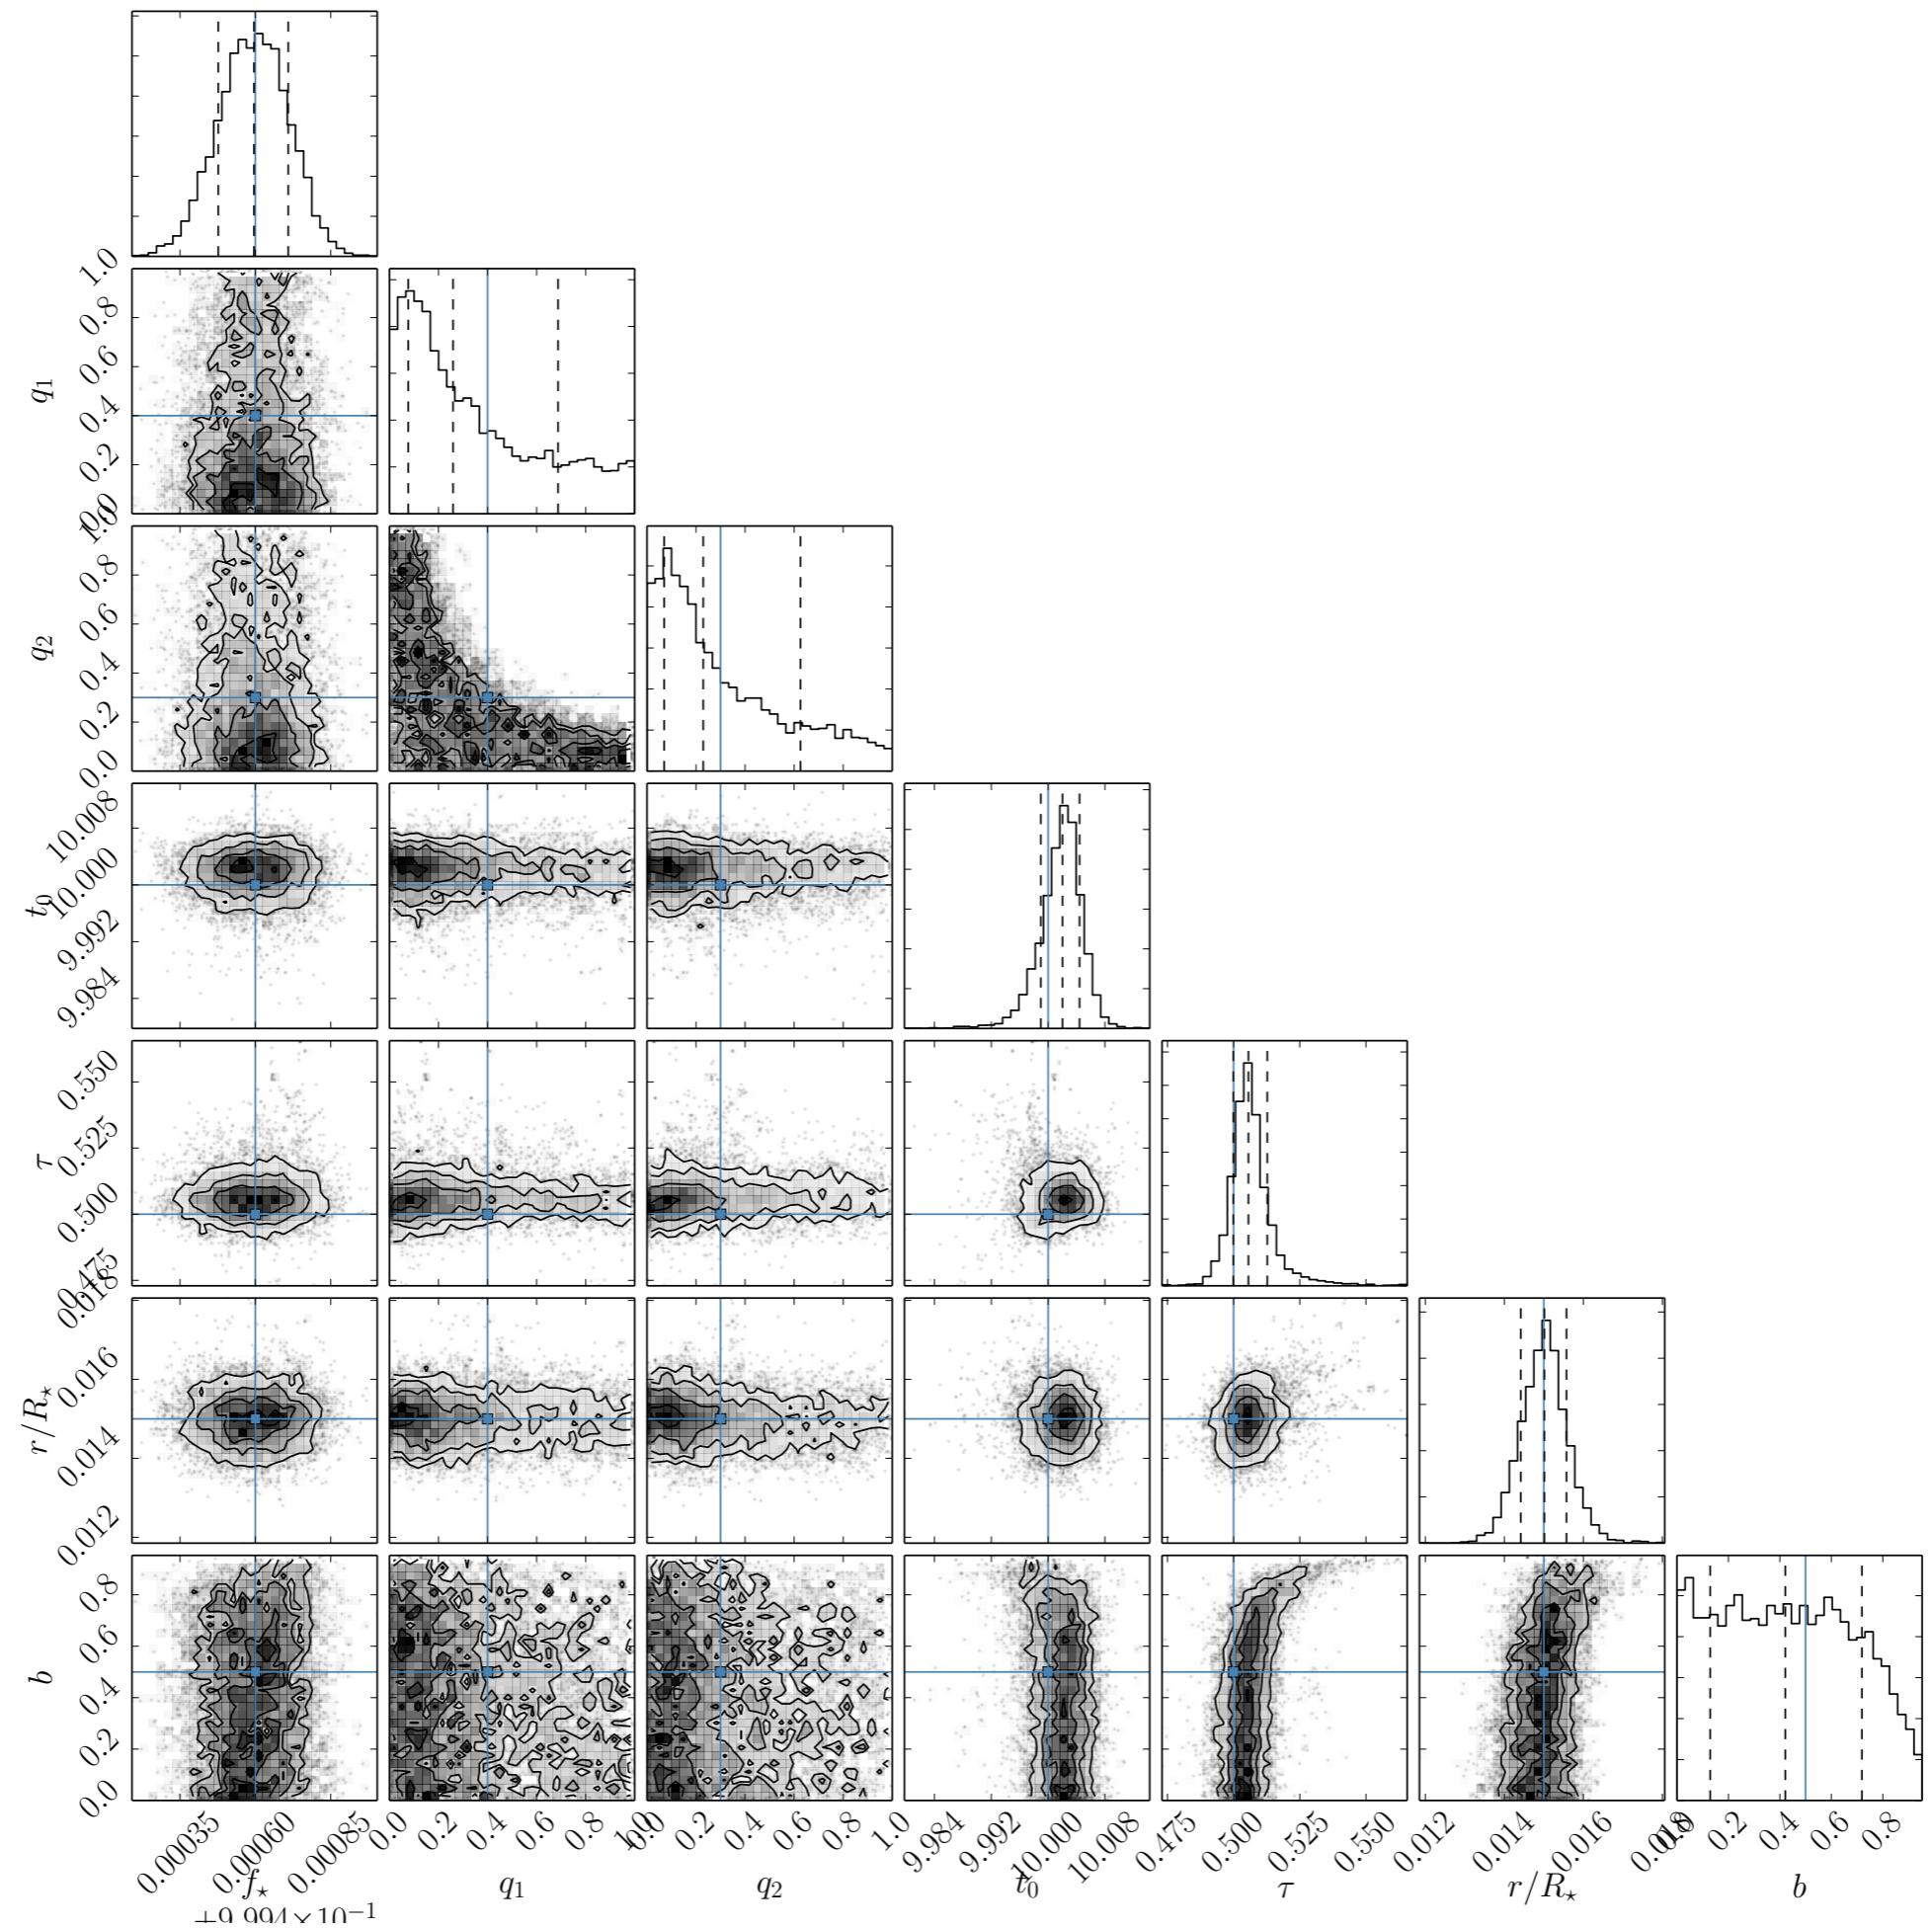
KIC 2301306 + injection

using Gaussian process **noise model**

---



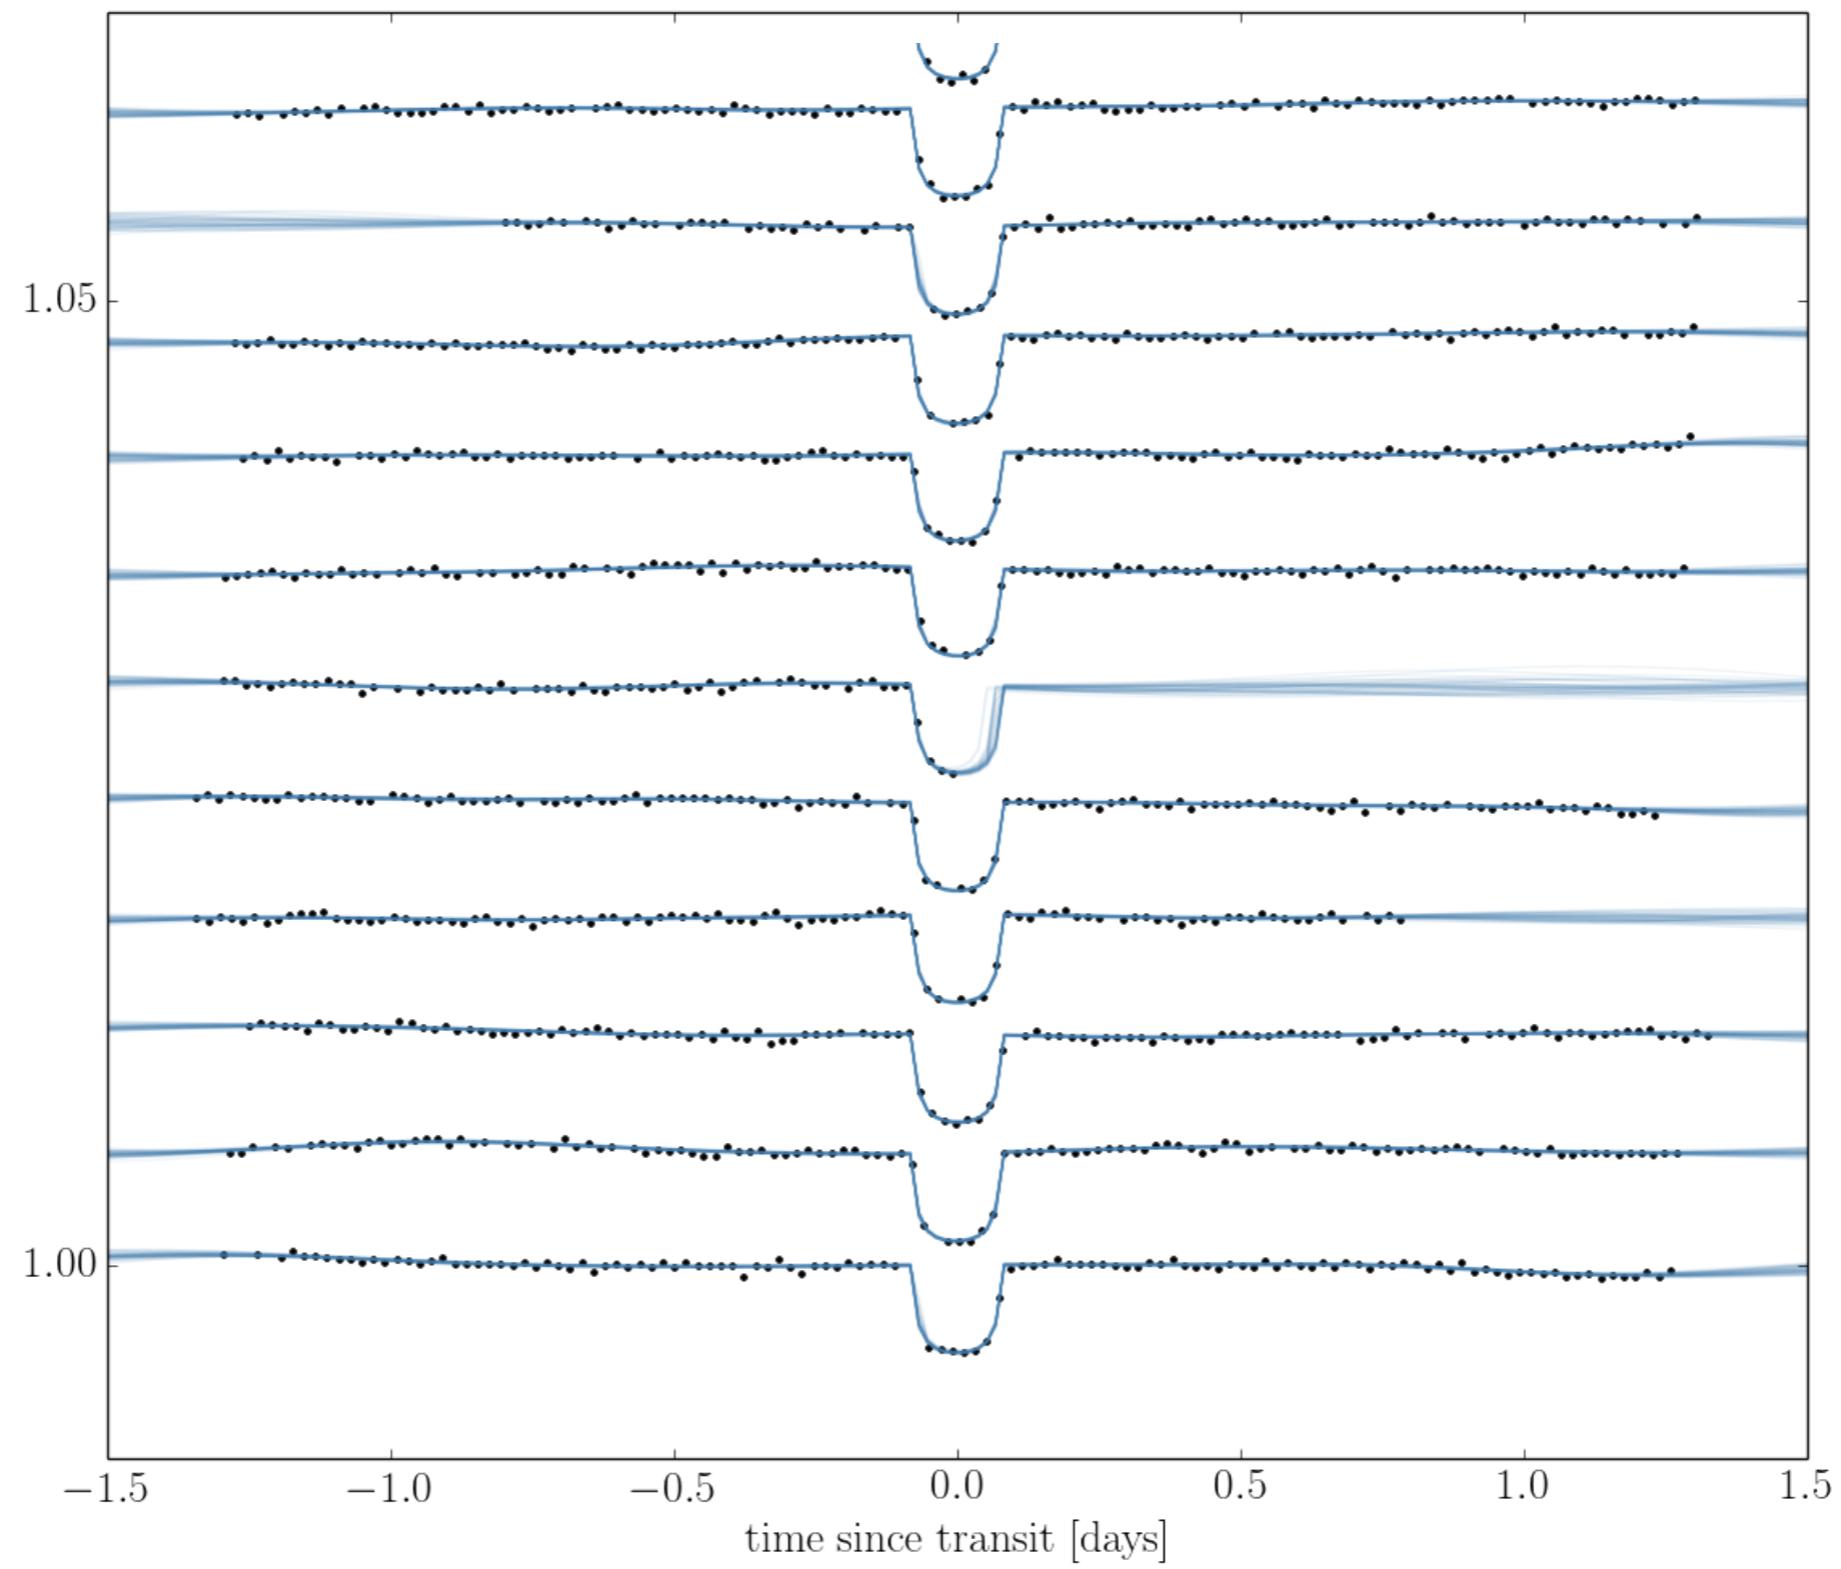


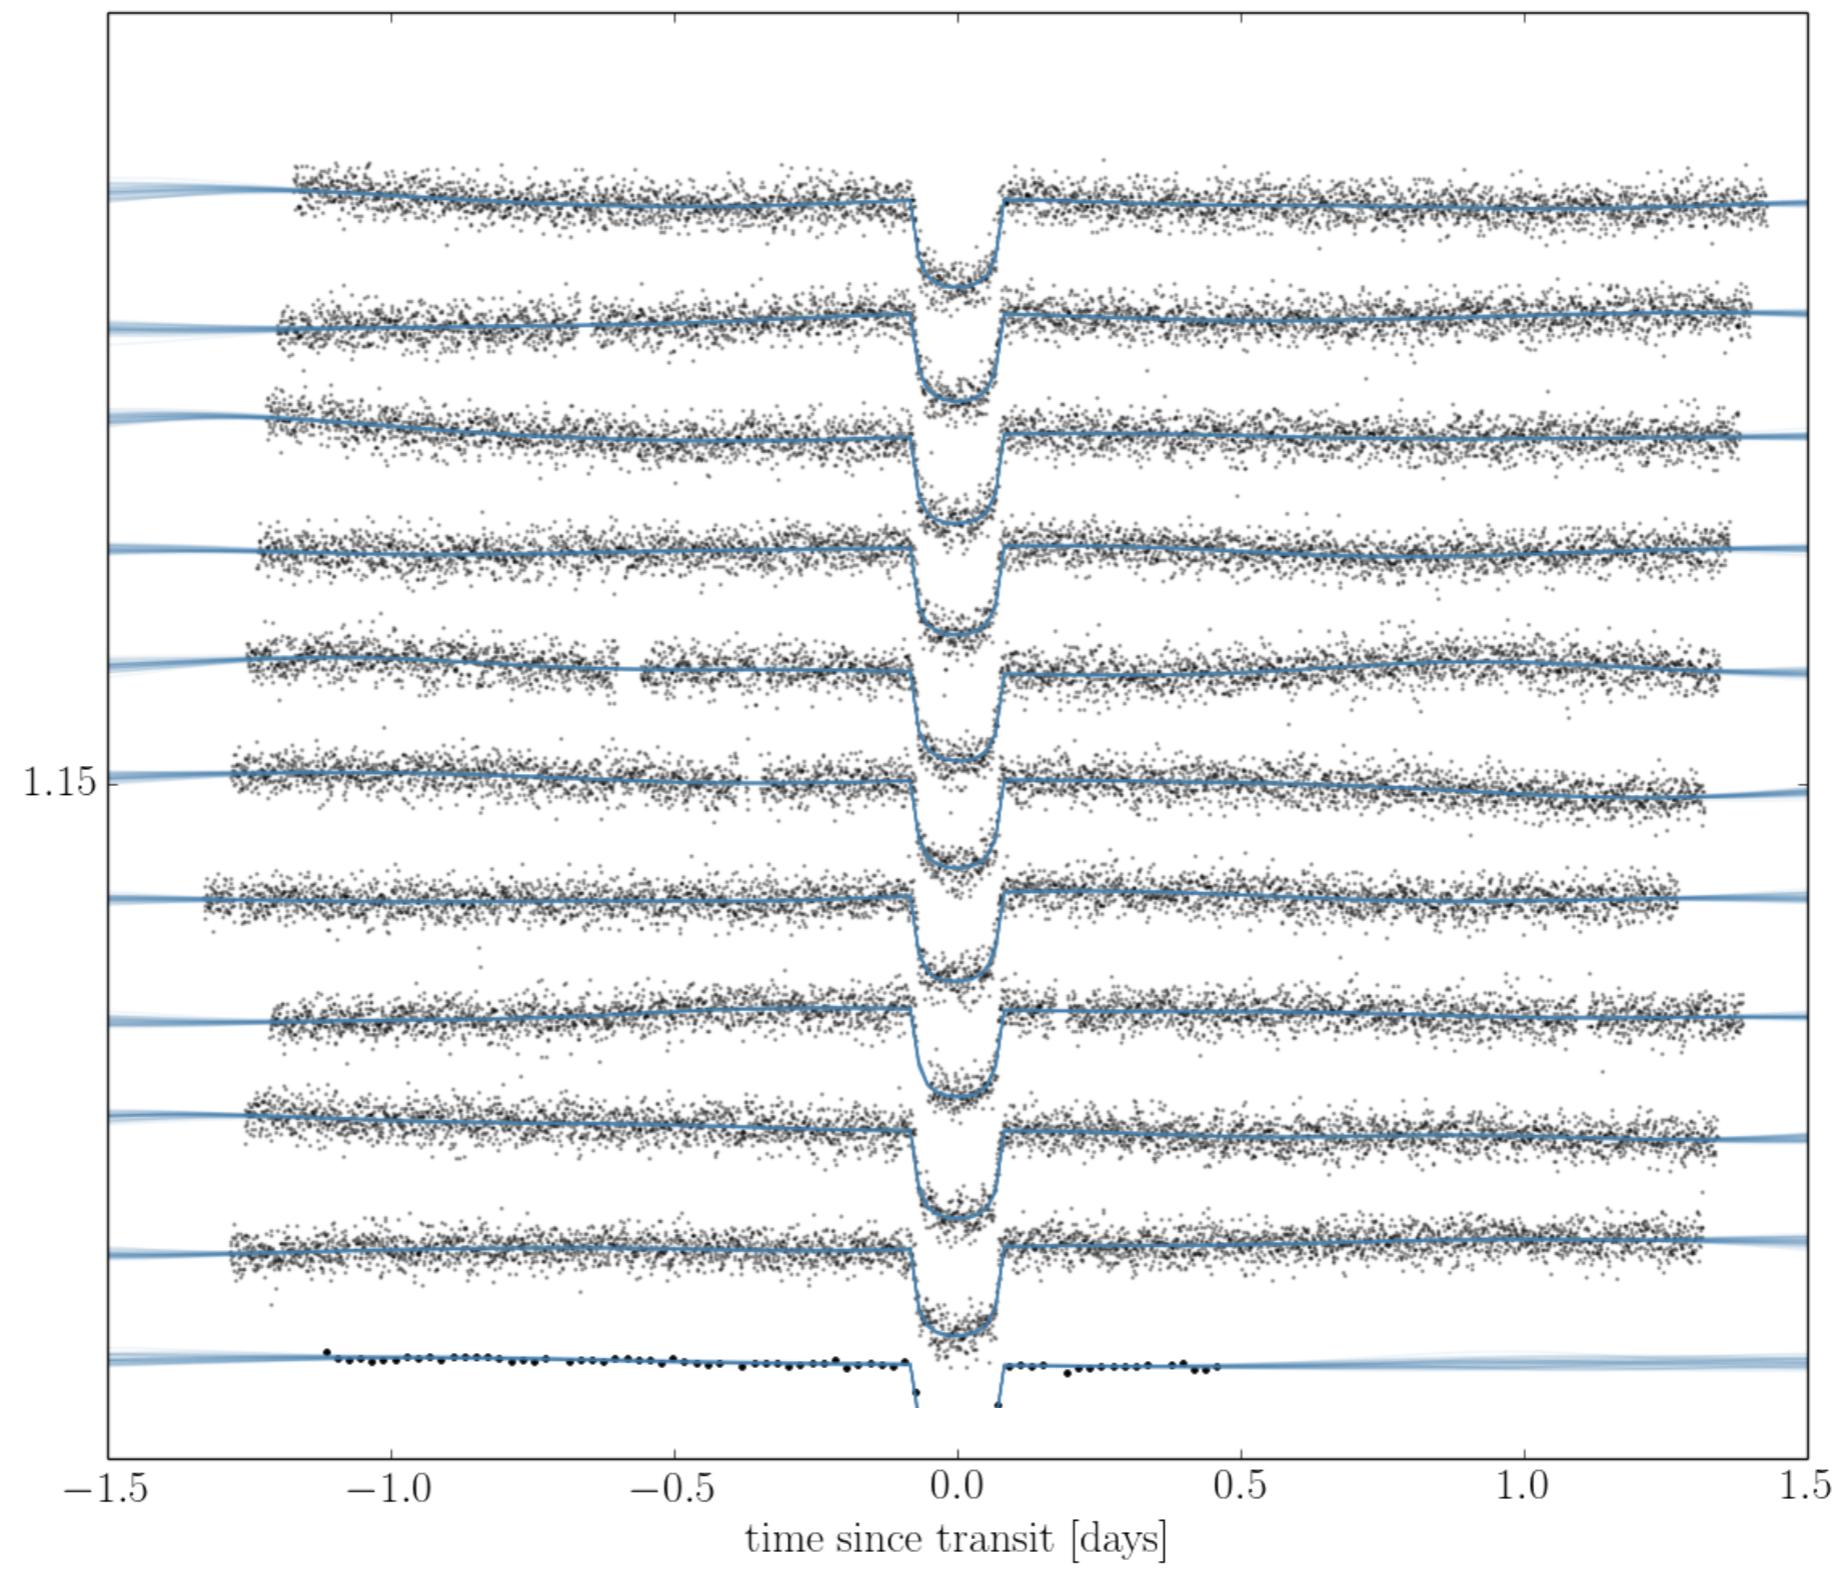


# KOI 1474.01

---

with Bekki Dawson, *et al.*

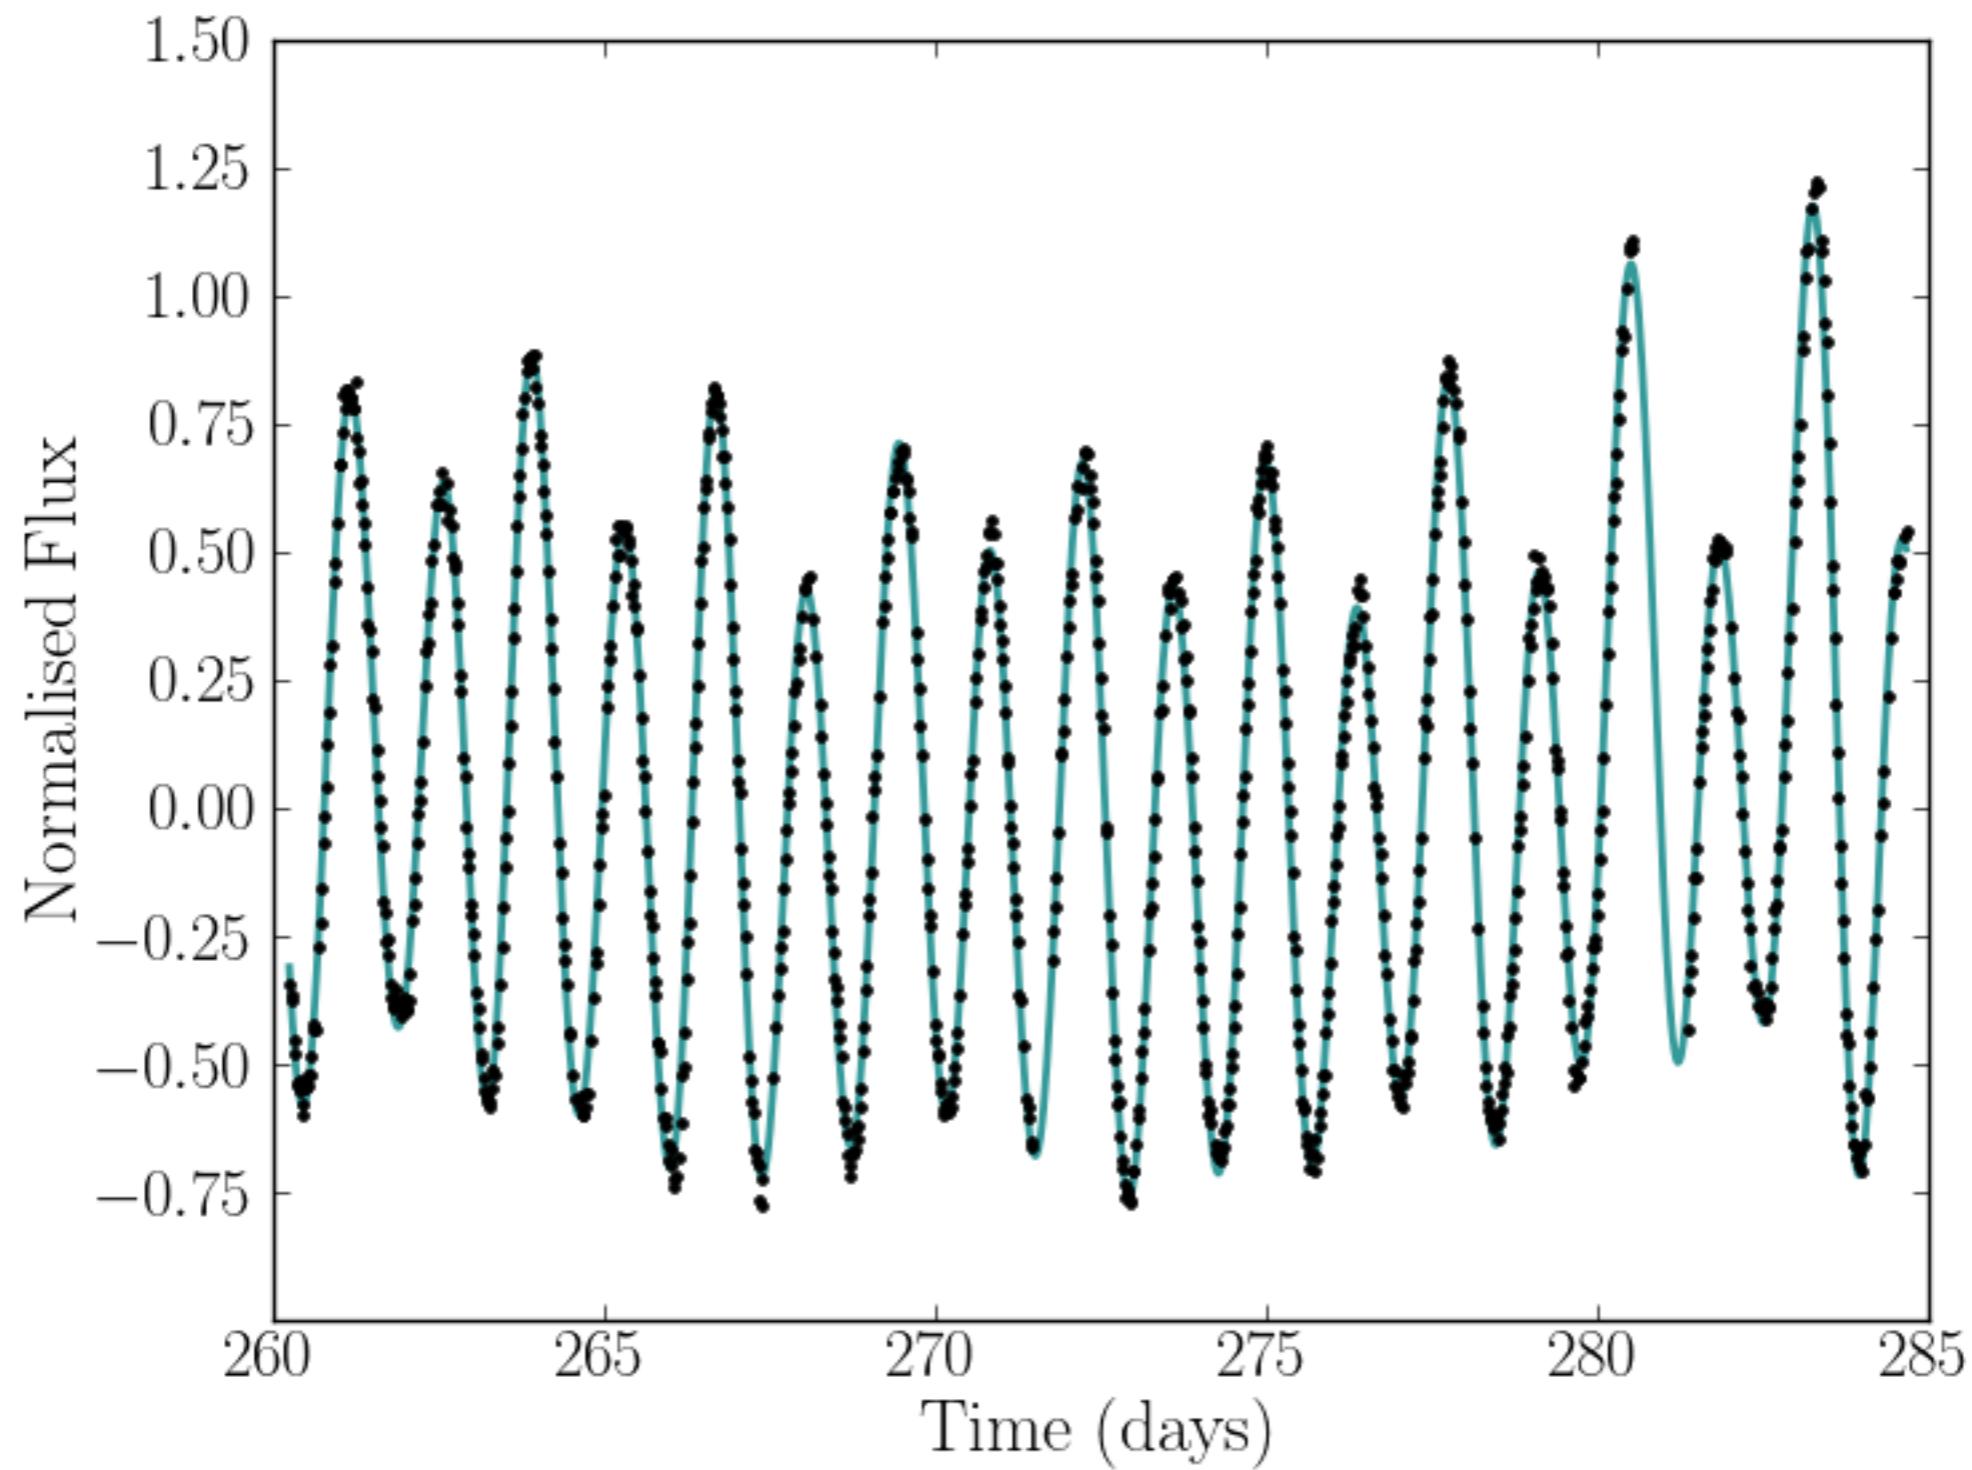




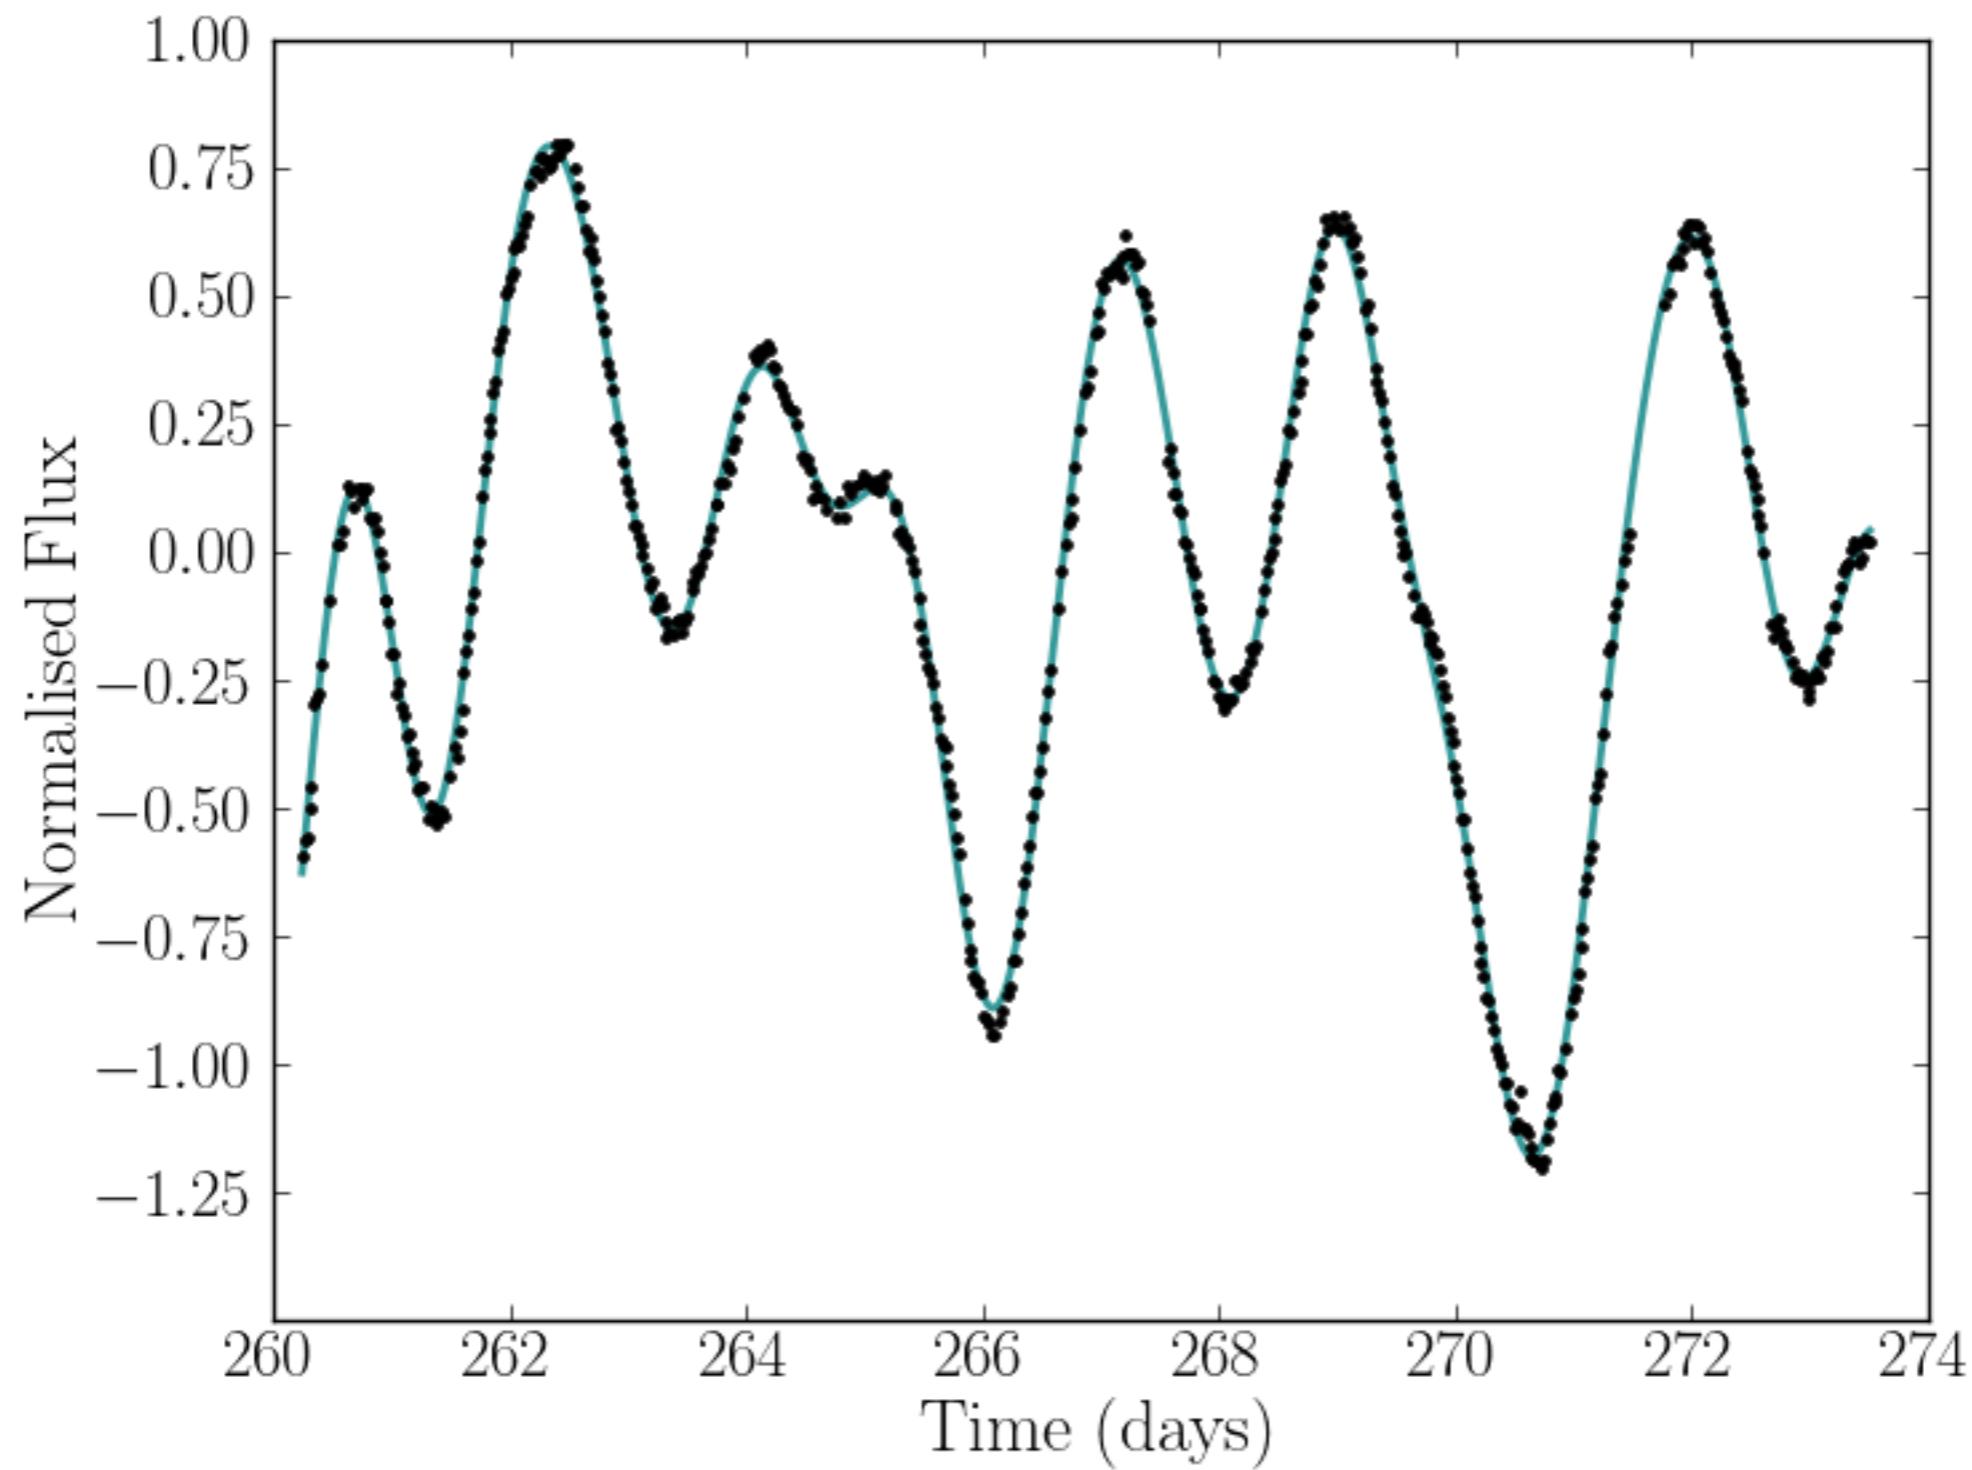
# Stellar Rotation

---

with Ruth Angus



figures from **Ruth Angus** (Oxford)



figures from Ruth Angus (Oxford)

# 4

## Conclusions & Summary.

---

correlated noise matters.

a Gaussian process provides a drop-in replacement likelihood function

if you can compute it

## Resources

---

[gaussianprocess.org/gpml](http://gaussianprocess.org/gpml)

[github.com/dfm/](https://github.com/dfm/) [gp  
george

[danfm@nyu.edu](mailto:danfm@nyu.edu)